

# **Database Development and Design (CPT201)**

## **Lecture 3b: B+ Tree Index**


Dr. Wei Wang  
Department of Computing

# Learning Outcomes

- B+-Tree Index
  - Queries
  - update


# B<sup>+</sup>-Tree Index

- B<sup>+</sup>-Tree is "short" and "Fat"
  - Disk-based: usually one node per block; large fan-out
  - Balanced (more or less): good performance guarantee.
- In a B<sup>+</sup>-Tree,
  - $n$  (or sometimes  $M$ ) is the number of pointers in a node; pointers:  $P_1, P_2, \dots, P_n$
  - Search keys:  $K_1 < K_2 < K_3 < \dots < K_{n-1}$
  - All paths (from root to leaf) have same length
  - Root must have at least two children
  - In each non-leaf node (inner node), more than 'half' ( $\geq \lceil n/2 \rceil$ ) pointers must be used
  - Each leaf node must contain at least  $\lceil (n-1)/2 \rceil$  keys



# Example

- An Example B+-Tree with  $n = 3$ 
  - All paths have same length. 🗨️
  - Root has (at least) two children
  - In each non-leaf node (inter node), more than half ( $\geq \lceil 3/2 \rceil = 2$ ) pointers are used
  - Each leaf node contains at least  $\lceil (3-1)/2 \rceil = 1$  key



# Queries on B<sup>+</sup>-Trees

- Find record with search-key value **V**.
  - 1.  $C = \text{root}$
  - 2. **While**  $C$  is not a leaf node
    - 2.1. Let  $i$  be least value such that  $V \leq K_i$ .
    - 2.2. If no such exists, set  $C =$  last non-null pointer in  $C$
    - 2.3. Else { if ( $V = K_i$ ) Set  $C = P_i + 1$  else set  $C = P_i$ }
  - 3. Let  $i$  be least value such that  $K_i = V$
  - 4. If there is such a value  $i$ , follow pointer  $P_i$  to the desired record.
  - 5. Else no record with search-key value  $V$  exists.

# Observations about B<sup>+</sup>-trees

- Since the inter-node connections are done by pointers, “logically” close blocks need not be “physically” close.
- The non-leaf levels of the B<sup>+</sup>-tree form a **hierarchy of sparse indices**.
- If there are K search-key values in the file
  - The B<sup>+</sup>-tree height is no more than  $\lceil \log_{\lceil n/2 \rceil}(K) \rceil$ .
  - Level below root has at least  $2 * \lceil n/2 \rceil$  values
  - Next level has at least  $2 * \lceil n/2 \rceil * \lceil n/2 \rceil$  values
  - .. etc.

# Observations about B<sup>+</sup>-trees cont'd

- Searching can be conducted efficiently.
  - a node is generally the same size as a disk block, typically 4 kilobytes
  - $n$  is typically around 100 (40 bytes per index entry).
  - with 1 million search key values and  $n = 100$
  - at most  $\log_{50}(1,000,000) = 4$  nodes are accessed in a lookup.
- Insertion and deletion to the main file can be handled efficiently, as the index can be restructured in **logarithmic** time.

# Updates on B<sup>+</sup>-Trees: Insertion

- 1. Find the leaf node in which the search-key value would appear
- 2. If the search-key value is already present in the leaf node
  - 2.1. Add record to the file
  - 2.2. If necessary add a pointer to the bucket.
- 3. If the search-key value is not present, then
  - 3.1. add the record to the main file (and create a bucket if necessary)
  - 3.2. If there is room in the leaf node, insert (key-value, pointer) pair in the leaf node
  - 3.3. Otherwise, split the node (along with the new (key-value, pointer) entry) as discussed in the next slide.




# Updates on B<sup>+</sup>-Trees: Insertion cont'd

- **Splitting a leaf node:**
  - take the (search-key value, pointer) pairs and the one being inserted) in an in-memory area  $M$  in sorted order. Assume there are  $n$  search key values in total.
  - Place the first  $\lceil n/2 \rceil$  in the original node, and the rest in a new node.
  - let the new node be  $p$ , and let  $k$  be the least key value in  $p$ . Insert  $(k,p)$  in the parent of the node being split.
  - If the parent is full, split it and propagate the split further up.
- Splitting of nodes proceeds upwards till a node that is not full is found.
  - In the worst case the root node may be split, increasing the height of the tree by 1.


# Updates on B<sup>+</sup>-Trees: Insertion cont'd

- **Splitting a non-leaf node**: when inserting  $(k,p)$  into an already full internal node  $N$ 
  - Copy  $N$  to an in-memory area  $M$  with space for  $n+1$  pointers and  $n$  keys
  - Insert  $(k,p)$  into  $M$  in sorted order
  - Copy  $P_1, K_1, \dots, K_{\lceil n/2 \rceil - 1}, P_{\lceil n/2 \rceil}$  from  $M$  back into node  $N$
  - Copy  $P_{\lceil n/2 \rceil + 1}, K_{\lceil n/2 \rceil + 1}, \dots, K_n, P_{n+1}$  from  $M$  into newly allocated node  $N'$
  - Insert  $(K_{\lceil n/2 \rceil}, N')$  into parent  $N$


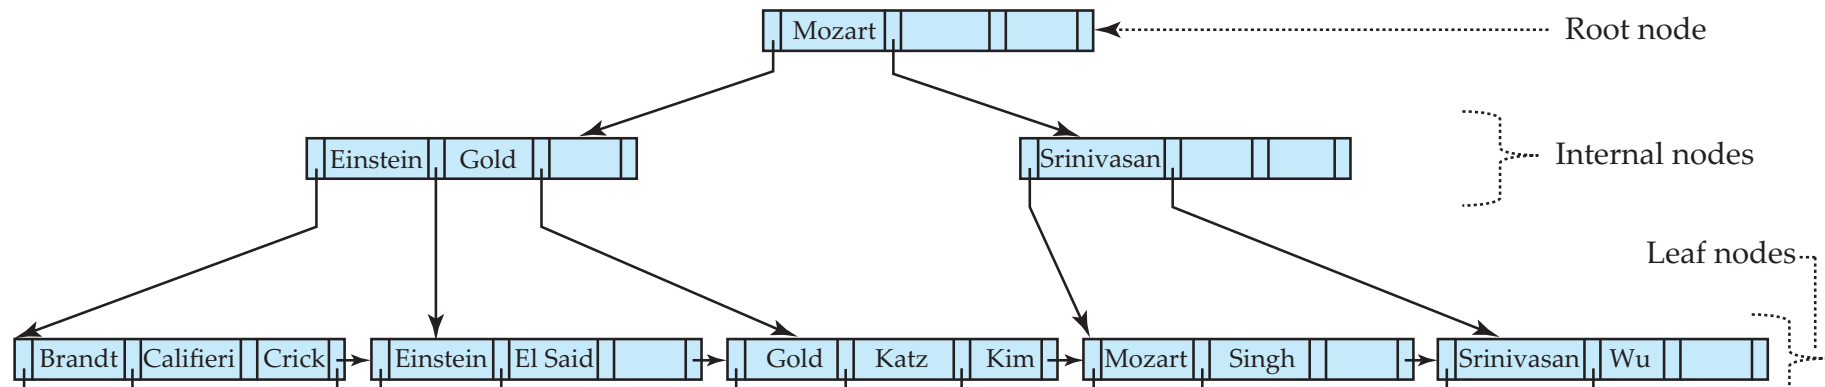
# Splitting a Leaf Node



# Splitting a Non-leaf Node

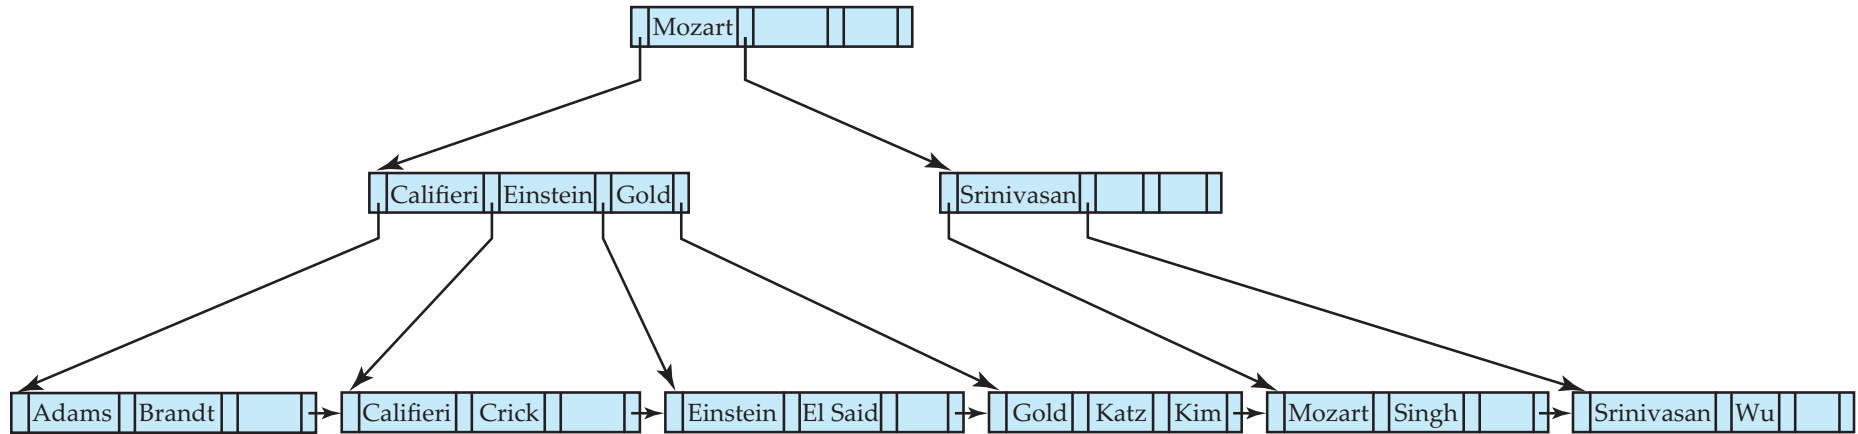


# Insertion Example



B+-Tree before and after insertion of "Adams"

# Insertion Example cont'd



**Question:**

What will happen after insertion of “Lamport”?

*Read pseudocode in textbook!*

# Exercise

- Construct a B+ tree for the following set of key values for  $n=3$ .
  - ( 2, 3, 5, 7, 11, 13, 17)

# Updates on B<sup>+</sup>-Trees: Deletion

- Find the record to be deleted, and remove it from the main file and from the bucket (if present)
- Remove (search-key value, pointer) from the leaf node if there is no bucket or if the bucket has become empty
- If the node has too few entries due to the removal, and the entries in the node and a sibling fit into a single node, then **merge siblings**:
  - Insert all the search-key values in the two nodes into a single node, and delete the other node.
  - If it is a non-leaf node, copy the value from the parent (between the two nodes) into the merged node
  - Delete the the value from the parent (between the two nodes). (Change may propagate to upper levels.)




# Updates on B<sup>+</sup>-Trees: Deletion

## cont'd

- Otherwise, if the node has too few entries due to the removal, but the entries in the node and a sibling do not fit into a single node, then **redistribute pointers**:
  - Redistribute the pointers between the node and a sibling such that both have more than the minimum number of entries; update the corresponding search-key value in the parent of the node.
  - If leaf node: take a proper value from sibling (value removed from sibling) and insert it to the underfull node; update the value in parent.
  - If non-leaf node: insert the value at (and remove from) parent to the underfull node, remove the value from sibling and update the parent.
  - **Read pseudocode in textbook!**
- The node deletions may cascade upwards till a node which has  $\lceil n/2 \rceil$  or more pointers is found.
- If the root node has only one pointer after deletion, it is **deleted** and the sole child becomes the **root**.


# Merge Siblings – at Leaf Node

- Merge siblings  $n1$  and  $n2$
- Delete  $K$  (and the appropriate pointer) from parent  $P$




# Merge Siblings – at non-Leaf Node

- Merge siblings  $n1$  and  $n2$  and  $K$
- Delete  $K$  (and the appropriate pointer) from parent  $P$




# Redistribute Pointers – at Leaf Node

- Copy  $K_i$  from  $n1$  and add it to  $n2$
- Delete  $K_i$  from  $n1$
- Replace the old value  $K$  in parent  $P$  with  $K_i$

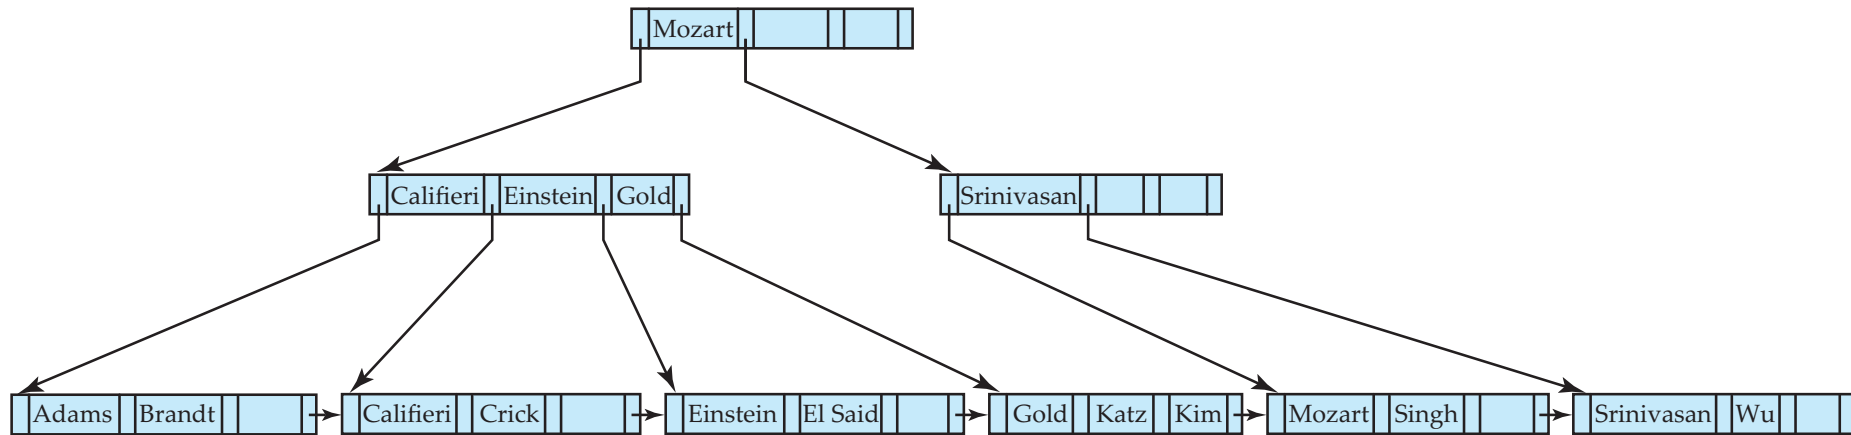


# Redistribute Pointers – at non-Leaf Node


- Copy  $K$  from parent  $P$  and add it to  $n2$
- Replace the old value  $K$  in parent  $P$  with  $K_i$  from  $n1$
- Delete  $K_i$  from  $n1$



# Deletion Example




Before and after deleting “Srinivasan”




- Deleting “Srinivasan” causes merging of under-full leaves

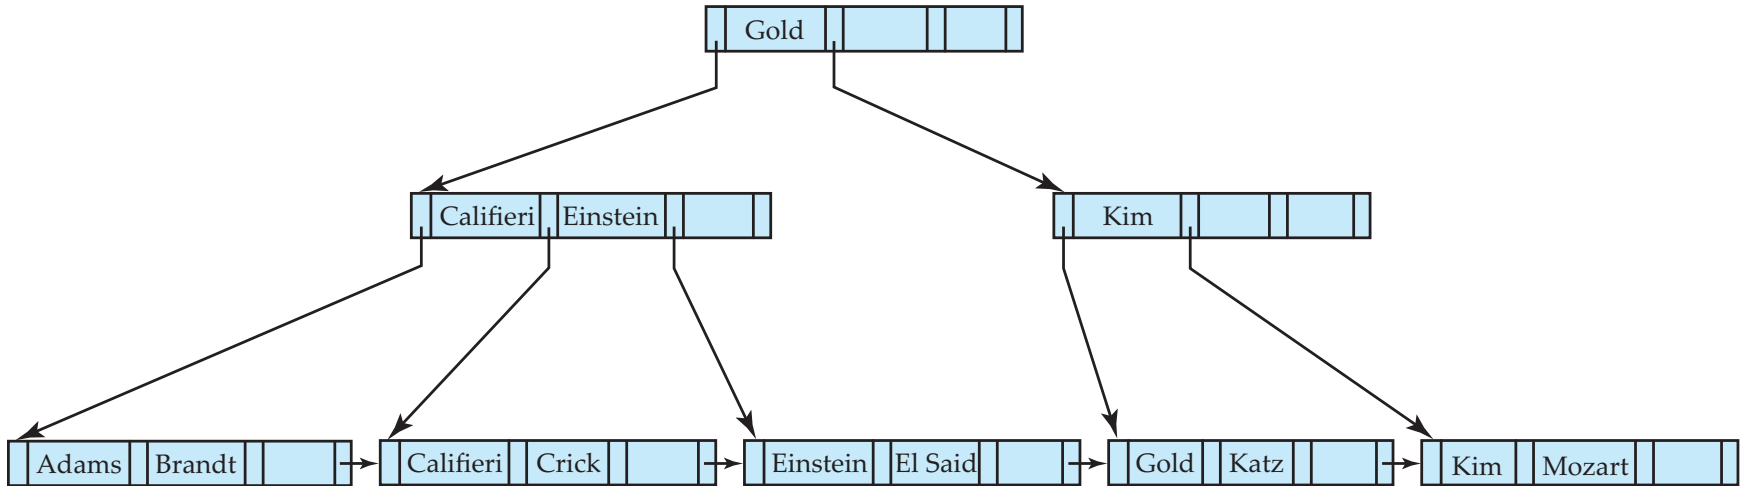
# Deletion Example cont'd




Before and after deleting “Singh and Wu”



# Deletion Example cont'd




Before and after deleting “Gold”






# More Example



# Another Example



# End of Lecture

- Summary
  - B+-Tree Index Files
    - lookup
    - Insertion
    - Deletion
- Reading
  - Database System Concepts, 6<sup>th</sup> edition, chapter 11.1, 11.2, 11.3
  - Database System Concepts, 7<sup>th</sup> edition, chapter 14.1, 14.2, 14.3