

# **Database Development and Design (CPT201)**

## **Lecture 4c: Query Evaluation - Join**

Dr. Wei Wang  
Department of Computing

# Learning Outcomes

- Algorithms for evaluating join operators
- Algorithms for evaluating other expressions

# Natural-Join Operation

**Notation:**  $r \bowtie s$

- Let  $r$  and  $s$  be relations on schemas  $R$  and  $S$  respectively.  
Then,  $r \bowtie s$  is a relation on schema  $R \cup S$  obtained as follows:
  - Consider each pair of tuples  $t_r$  from  $r$  and  $t_s$  from  $s$ .
  - If  $t_r$  and  $t_s$  have the same value on each of the attributes in  $R \cap S$ , add a tuple  $t$  to the result, where
    - $t$  has the same value as  $t_r$  on  $r$
    - $t$  has the same value as  $t_s$  on  $s$
- Example:

$R = (A, B, C, D)$

$S = (E, B, D)$

- Result schema =  $(A, B, C, D, E)$
- $r \bowtie s$  is defined as:


$$\Pi_{r.A, r.B, r.C, r.D, s.E} (\sigma_{r.B = s.B \wedge r.D = s.D} (r \times s))$$



# Join Operation

- Several different algorithms to implement joins exist (not counting with the ones involving parallelism)
  - Nested-loop join
  - Block nested-loop join
  - Indexed nested-loop join
  - Merge-join
  - Hash-join
- As for selection, the choice is based on cost estimate.

# Banking Example



- Examples in next slides use the following information:
  - Number of records of *customer*: 10,000
  - Number of blocks of *customer*: 400
  - Number of records of *depositor*: 5,000
  - Number of blocks of *depositor*: 100

# Nested-Loop Join

- The simplest join algorithms, that can be used **independently** of everything (like the linear search for selection)
- To compute the theta join:  $r \bowtie_{\theta} s$   
for each tuple  $t_r$  in  $r$  do begin  
  for each tuple  $t_s$  in  $s$  do begin  
    test pair  $(t_r, t_s)$  to see if they satisfy the join condition  $\theta$   
    if they do, add  $t_r \cdot t_s$  to the result.  
  end  
end
- $r$  is called the **outer relation** and  $s$  the **inner relation** of the join.
- Quite expensive in general, since it requires to examine every pair of tuples in the two relations.

# Nested-Loop Join cont'd

- In the worst case, if there is enough memory **only** to hold one block of each relation,  $n_r$  is the number of tuples in relation  $r$ , the estimated cost is  
 $n_r * b_s + b_r$  block transfers, plus  
 $n_r + b_r$  seeks
- If the smaller relation fits entirely in memory, use that as the inner relation.
  - Reduces cost to  $b_r + b_s$  block transfers and 2 seeks
- But in general, it is much better to have the **smaller** relation as the **outer** relation
- The choice of the inner and outer relation strongly depends on the estimate of the size of each relation.

# Nested-Loop Join Cost in Example

- Assuming **worst case** memory availability cost estimate is
  - with *depositor* as outer relation:
    - $5,000 * 400 + 100 = 2,000,100$  block transfers,
    - $5,000 + 100 = 5,100$  seeks
  - with *customer* as the outer relation
    - $10,000 * 100 + 400 = 1,000,400$  block transfers and 10,400 seeks
- If smaller relation (*depositor*) fits entirely in memory, the cost estimate will be 500 block transfers and 2 seeks
- Instead of iterating over records, one could iterate over blocks. This way, instead of  $n_r * b_s + b_r$  we would have  $b_r * b_s + b_r$  block transfers
- This is the basis of the block nested-loops algorithm.



# Block Nested-Loop Join

- Variant of nested-loop join in which every block of inner relation is paired with every block of outer relation.

```
for each block  $B_r$  of  $r$  do begin
  for each block  $B_s$  of  $s$  do begin
    for each tuple  $t_r$  in  $B_r$  do begin
      for each tuple  $t_s$  in  $B_s$  do begin
        Check if  $(t_r, t_s)$  satisfy the join condition
        if they do, add  $t_r \cdot t_s$  to the result.
```

end

end

end

end

# Block Nested-Loop Join Cost

- Worst case estimate:  $b_r * b_s + b_r$  block transfers and  $2 * b_r$  seeks
  - Each block in the inner relation  $s$  is **read once** for each *block* in the outer relation (instead of once for each tuple in the outer relation).
- Best case (when smaller relation fits into memory):  $b_r + b_s$  block transfers plus 2 seeks.
- Some improvements to nested loop and block nested loop algorithms can be made:
  - Scan inner loop forward and backward alternately, to make use of the blocks remaining in buffer, reduce the number of disk access
  - Use index on inner relation (if available) to quickly get the tuples which match the tuple of the outer relation.

# Indexed Nested-Loop Join

- Index lookups can replace file scans if
  - join is an equi-join or natural join and
  - an index is available on the inner relation's join attribute
    - In some cases, it pays to construct an index just to compute a join.
- For each tuple  $t_r$  in the outer relation  $r$ , use the index on  $s$  to look up tuples in  $s$  that satisfy the join condition with tuple  $t_r$ .
- Worst case: buffer has space for only one page of  $r$ , and, for each tuple in  $r$ , we perform an index lookup on  $s$ .
- Cost of the join:  $b_r + n_r * c$  block transfers and seeks
  - Where  $c$  is the cost of traversing index and fetching all matching  $s$  tuples for one tuple in  $r$
  - $c$  can be estimated as cost of a single selection on  $s$  using the join condition (usually quite low, when compared to the join)
- If indices are available on join attributes of both  $r$  and  $s$ , use the relation with **fewer** tuples as the **outer** relation.



# Example of Indexed Nested-Loop Join Costs

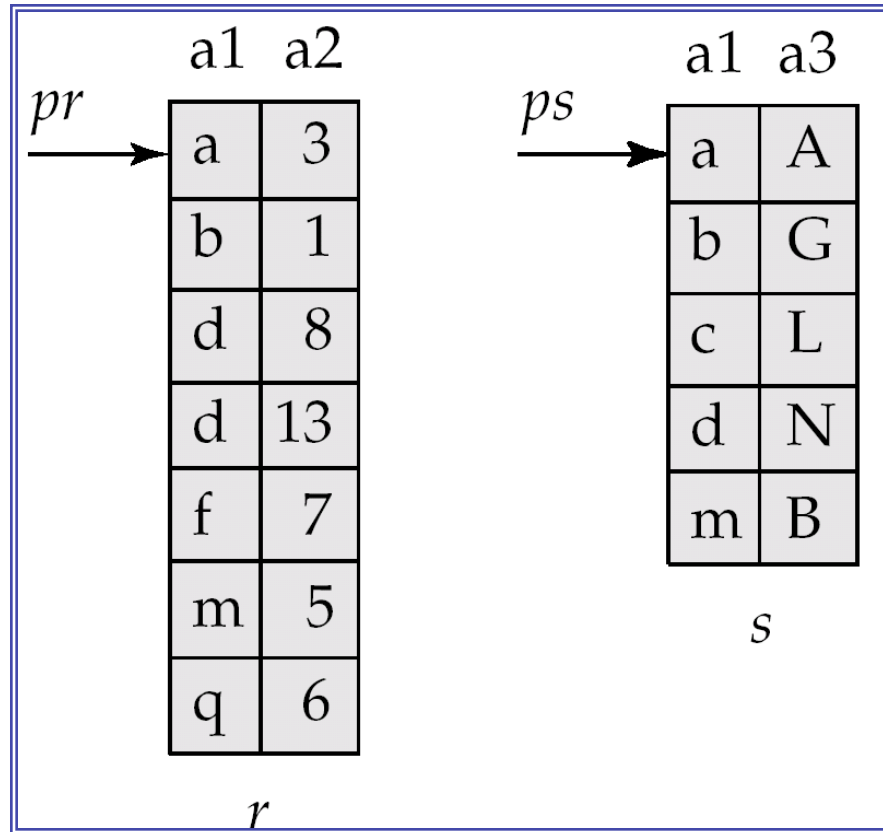
- Compute *depositor* ⋈ *customer*, with *depositor* as the **outer** relation.
- Let *customer* have a primary B<sup>+</sup>-tree index on the join attribute *customer-name*, with n=20.
- Since *customer* has 10,000 tuples, the height of the tree is 4, and one more access is needed to find the actual data
- *depositor* has 5,000 tuples
- **Nested loop join**: 2,000,100 block transfers and 5,100 seeks
- Cost of **block nested loops join**
  - $400 \times 100 + 100 = 40,100$  block transfers +  $2 \times 100 = 200$  seeks
- Cost of **indexed nested loops join**
  - $100 + 5,000 \times (4+1) = 25,100$  block transfers and seeks.
  - The number of block transfers is less than that for block nested loops join
  - But number of seeks is much larger
  - In this case using the index **doesn't pay** (this is specially so because the relations are small)

# Merge-Join

1. Sort both relations on their join attribute (if not already sorted on the join attributes).  
Join step is similar to the merge stage of the sort-merge algorithm.
2. Merge-join algorithm
  1. Initialise two pointers point to  $r$  and  $s$
  2. While not done
    1. the pointers to  $r$  and  $s$  move through the relation.
    2. A group of tuples of inner relation  $s$  with the same value on the join attributes is read into  $S_s$ .
    3. Do join on tuple pointed by  $p_r$  and tuples in  $S_s$ ;
  3. End while

# Merge-Join cont'd

Read pseudocode in the textbook!



# Merge-Join cont'd

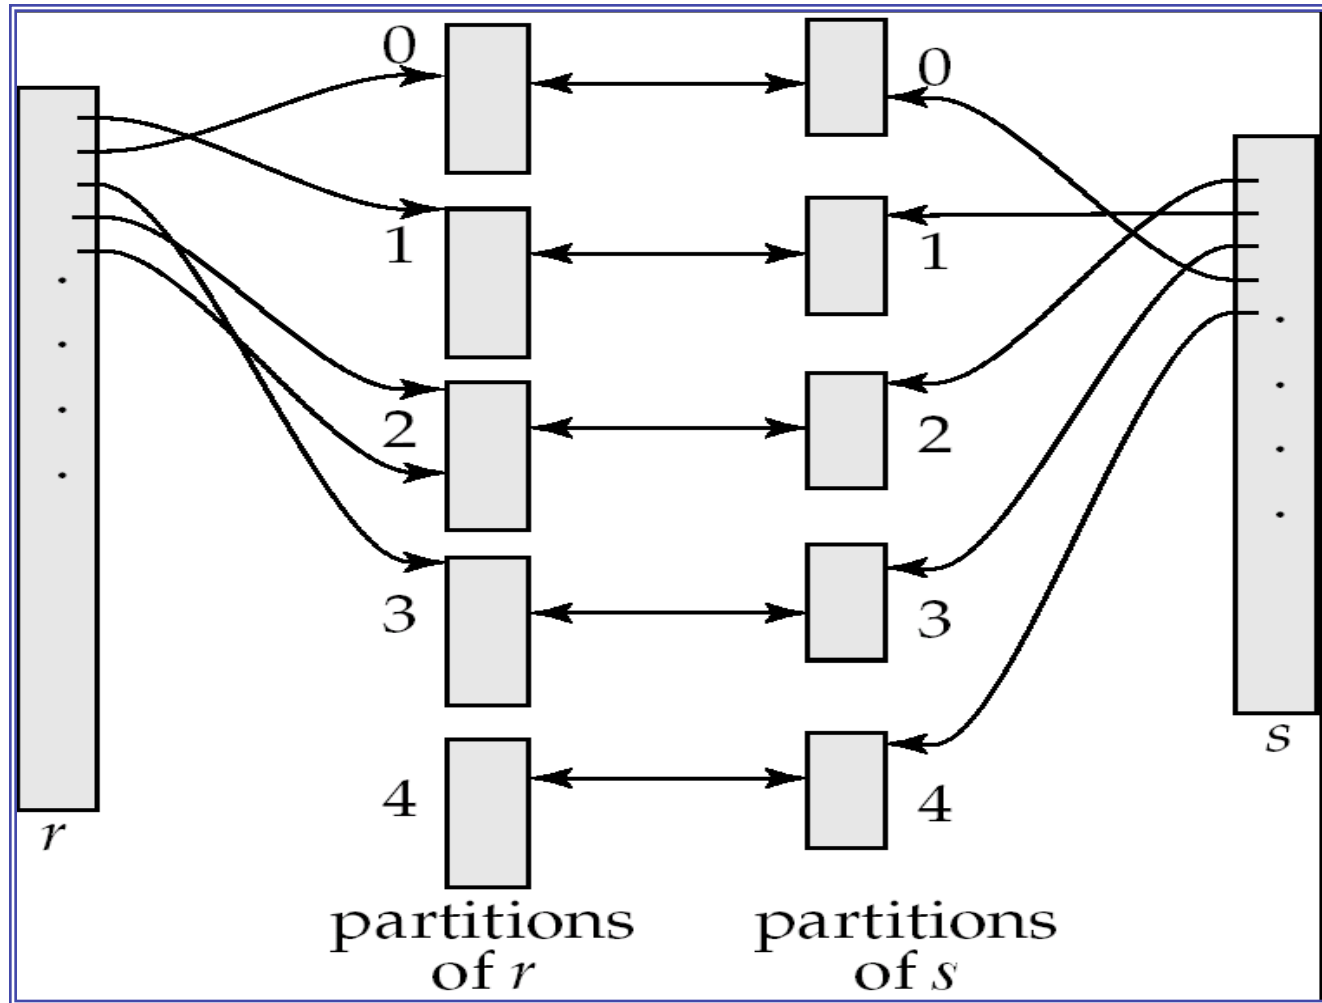
- Can be used only for equi-joins and natural joins
- Each block needs to be read only once (assuming that all tuples for any given value of the join attributes fit in memory)
- Thus the cost of merge join is (where  $b_b$  is the number of blocks in allocated in memory for each relation):  
$$b_r + b_s \text{ block transfers} +$$
$$\lceil b_r / b_b \rceil + \lceil b_s / b_b \rceil \text{ seeks}$$
  - **Plus** the cost of sorting if relations are unsorted.
  - Since seeks are much more expensive than data transfer, it makes sense to allocate multiple buffer blocks to each relation, provided extra memory is available.

# Hash-Join

- Also only applicable for equi-joins and natural joins.
- A hash function  $h$  is used to partition tuples of both relations
- $h$  maps  $JoinAttrs$  values to  $\{0, 1, \dots, n\}$ , where  $JoinAttrs$  denotes the common attributes of  $r$  and  $s$  used in the natural join.
  - $r_0, r_1, \dots, r_n$  denote partitions of  $r$  tuples
    - Each tuple  $t_r \in r$  is put in partition  $r_i$  where  $i = h(t_r[JoinAttrs])$ .
  - $s_0, s_1, \dots, s_n$  denotes partitions of  $s$  tuples
    - Each tuple  $t_s \in s$  is put in partition  $s_i$ , where  $i = h(t_s[JoinAttrs])$ .
- General idea:
  - Partition the relations according to this
  - Then perform the join on each partition  $r_i$  and  $s_i$ 
    - There is no need to compute the join between different partitions since an  $r$  tuple and an  $s$  tuple that satisfy the join condition will have the same value for the join attributes. If that value is hashed to some value  $i$ , the  $r$  tuple has to be in  $r_i$  and the  $s$  tuple in  $s_i$ .



# Hash-Join cont'd



# Hash-Join Algorithm

1. Partition the relation  $s$  using hashing function  $h$ .  
When partitioning a relation, some blocks of memory ( $b_b$ ) are reserved as the output buffer for each partition.
2. Partition  $r$  similarly.
3. For each  $i$ :
  - (a) Load  $s_i$  into memory and build an in-memory hash index on it using the join attribute. This hash index uses a **different hash** function than the earlier  $h$  for partitioning.
  - (b) Read the tuples in  $r_i$  from the disk one by one. For each tuple  $t_r$  locate each matching tuple  $t_s$  in  $s_i$  using the in-memory hash index. Output the concatenation of their attributes.

---

Relation  $s$  is called the **build input** and  $r$  is called the **probe input**.



# Hash-Join algorithm cont'd

- The number of partitions  $n$  for the hash function  $h$  is chosen such that each  $s_i$  should fit in memory.
  - Typically  $n$  is chosen as  $\lceil b_s/M \rceil * f$  where  $f$  is a "fudge factor", typically around 1.2, to avoid overflows
  - The probe relation partitions  $r_i$  need not fit in memory

# Cost of Hash-Join

- The cost of hash join is  $3(b_r + b_s) + 4 * n_h$  block transfers, and  $2(\lceil b_r / b_b \rceil + \lceil b_s / b_b \rceil) + 2 * n_h$  seeks
  - each of the  $n_h$  partitions could have a partially filled block that has to be written and read back
  - The build and probe phases require only one seek for each of the  $n_h$  partitions of each relation, since each partition can be read sequentially.
- If the entire build input can be kept in main memory (then no partitioning is required), Cost estimate goes down to  $b_r + b_s$  and 2 seeks.

# Cost of Hash-Join in Example

- For the running example, assume that memory size is 20 blocks  $b_{\text{ depositor}} = 100$  and  $b_{\text{ customer}} = 400$ .
- *depositor* is to be used as build input. Partition it into five partitions, each of size 20 blocks. This partitioning can be done in one pass. Similarly, partition *customer* into five partitions, each of size 80. This is also done in one pass.
- Assuming 3 blocks are allocated for the input buffer and each output buffer
- Therefore total cost, ignoring cost of writing partially filled blocks:
$$3(100 + 400) = 1,500 \text{ block transfers} +$$
$$2(\lceil 100/3 \rceil + \lceil 400/3 \rceil) + 2*5 = 344 \text{ seeks}$$
- We had up to here:
  - 40,100 block transfers plus 200 seeks (for block nested loop)
  - 25,100 block transfers and seeks (for index nested loop).

# Other Operations: Duplicate Elimination

- **Duplicate elimination** can be implemented via hashing or sorting.
  - On sorting duplicates will come adjacent to each other, and all but one set of duplicates can be deleted.
  - *Optimisation:* duplicates can be deleted during run generation as well as at intermediate merge steps in external sort-merge.
  - Hashing is similar - duplicates will come into the same bucket.
- **Projection:**
  - perform projection on each tuple;
  - followed by duplicate elimination.

# Other Operations: Aggregation

- **Aggregation** can be implemented similarly to duplicate elimination.
  - Sorting or hashing can be used to bring tuples in the same group together, and then the aggregate functions can be applied on each group.
  - *Optimisation*: combine tuples in the same group during run generation and intermediate merges, by computing partial aggregate values
    - For count, min, max, sum: keep aggregate values on tuples found so far in the group.
      - When combining partial aggregate for count, add up the aggregates
    - For avg, keep sum and count, and divide sum by count at the end

# Other Operations: Set Operations

- **Set operations** ( $\cup$ ,  $\cap$  and  $-$ ): can either use variant of merge-join after sorting, or variant of hash-join.
- Set operations using **hashing**:
  1. Partition both relations using the same hash function
  2. Process each partition  $i$  as follows.
    1. Using a **different hashing function**, build an **in-memory hash index** on  $r_i$ .
    2. Process  $s_i$  as follows
      - $r \cup s$ :
        1. Add tuples in  $s_i$  to the hash index if they are not in it.
        2. At the end, add the tuples in the hash index to the result.
      - $r \cap s$ :
        1. output tuples in  $s_i$  to the result if they are already in the hash index
      - $r - s$ :
        1. for each tuple in  $s_i$ , if it is in the hash index, delete it from the index.
        2. At the end, add remaining tuples in the hash index to the result.



# End of Lecture

## ■ Summary

- Join
  - Nested-Loop Join
  - Block-Nested-Loop Join
  - Indexed-Nested-Loop Join
  - Sorted-Merge-Join
  - Hash Join
- Other Operations

## ■ Reading

- 6<sup>th</sup> edition, Chapters 12.5 and 12.6
- 7<sup>th</sup> edition, Chapters 15.5 and 15.6