

# **Database Development and Design (CPT201)**

## **Lecture 5a: Introduction to Query Optimisation 1**

Dr. Wei Wang

Department of Computing

# Learning Outcomes

- Introduction to Query Optimisation
  - Transformation of Relational Expressions




# Evaluation of Expressions

- So far we have seen algorithms for individual operations
  - These have then to be combined to evaluate complex expressions, with multiple operations.
- Alternatives for evaluating an entire expression tree
  - **Materialisation:** generate results of an expression whose inputs are relations or relations that are already computed. Temporary relations must be **materialised** (stored) on disk.
  - **Pipelining:** pass on tuples to parent operations even as the operation is being executed.

# Materialisation

- **Materialised evaluation:** evaluate one operation at a time, starting at the lowest-level. Use intermediate results materialised into temporary relations to evaluate next-level operations.
  - e.g., in figure below, compute and store the selection, then compute its join with *customer* and store the result, and finally compute the projections on *customer-name*.



# Materialisation cont'd

- Materialised evaluation is always applicable
- It may require considerable storage space.  
Moreover, cost of writing results to disk and reading them back can be quite high
  - Our cost formulas for operations ignore cost of writing final results to disk, so:
  - Overall cost = Sum of costs of individual operations +  
cost of writing intermediate results to disk
- Double buffering: use two output buffers for each operation, when one is full, write it to disk while the other is getting filled
  - Allows overlap of disk writes with computation and reduces execution time

# Pipelining

- **Pipelined evaluation:** evaluate several operations simultaneously, passing the results of one operation on to the next.
  - e.g., in previous expression tree, don't store result of the selection
    - instead, pass tuples directly to the join.
    - Similarly, don't store result of join, pass tuples directly to projection.
- It is much cheaper than materialisation: there is no need to store a temporary relation to disk.
- Pipelining may not always be possible - e.g., sort and hash-join where a preliminary phase is required over the whole relations.
- Pipelines can be executed in two ways: **demand driven** and **producer driven**.

# Producer-Driven Pipelining

- In producer-driven (or eager or push) pipelining
  - Operators produce tuples eagerly and pass them up to their parents
    - buffer maintained between operators, child puts tuples in buffer, parent removes tuples from buffer
    - if buffer is full, child waits till there is space in the buffer, and then generates more tuples
  - System schedules operations that have space in output buffer and can process more input tuples.



# Demand-Driven Pipelining

- In demand driven (or lazy, or pull) evaluation
  - system repeatedly requests next tuple from top level operation
  - Each operation requests next tuple from child operations as required, in order to output its next tuple
  - In between calls, operation has to maintain "state" so it knows what to return next.



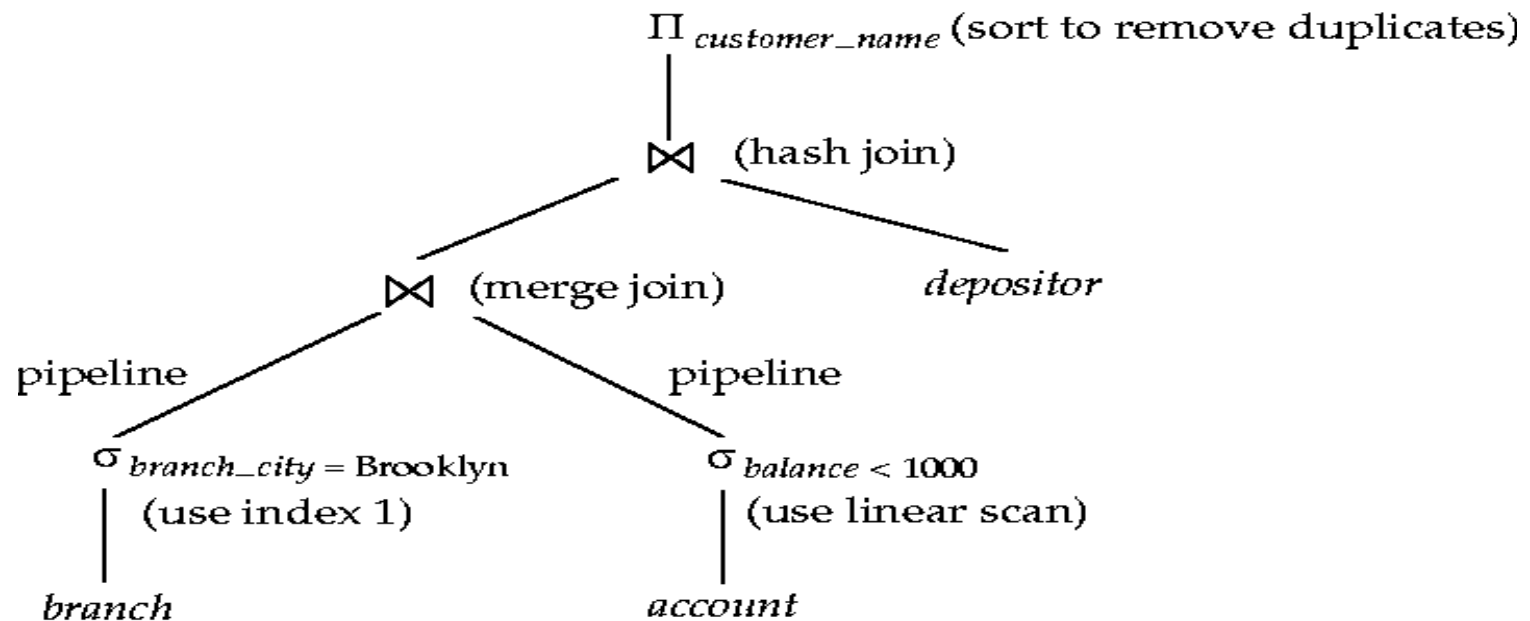
# Equivalent expressions

- Alternative ways of evaluating a given query
  - Equivalent expressions
  - Different algorithms for each operation



# Evaluation Plan

- An **evaluation plan** defines exactly what algorithm is used for each operation, and how the execution of the operations is coordinated.



# Cost-based query Optimisation

- Cost difference between evaluation plans for a query can be enormous
  - E.g. seconds vs. days in some cases
- **Cost-based query optimisation**
  - Find logically equivalent expressions of the given expression (but more efficient to execute)
  - Select a detailed strategy for processing the query, such as choosing the algorithm to use for executing an operation or choosing the specific indices to use
- Estimation of plan cost based on:
  - **Statistical information** about relations, e.g., number of tuples, number of distinct values for an attribute
  - **Statistical estimation for intermediate results** to compute cost of complex expressions
  - **Cost formulae** for algorithms, computed using statistics
- It should be noted that since the cost is an estimate, the selected plan is **not** necessarily the least-costly plan; however, as long as the estimates are good, the plan will not be much more costly than it.

# Transformation of Relational Expressions

- Two relational algebra expressions are said to be **equivalent** if the two expressions generate the same set of tuples on every legal database instance
  - Note: order of tuples is irrelevant
  - In SQL, inputs and outputs are multisets of tuples
  - Two expressions in the multiset version of the relational algebra are said to be equivalent if the two expressions generate the same multiset of tuples on every legal database instance.
- An **equivalence rule** says that expressions of two forms are equivalent if
  - Can replace expression of first form by second, or vice versa

# Equivalence Rules

- Rule 1: Conjunctive selection operations can be deconstructed into a sequence of individual selections.

$$\sigma_{\theta_1 \wedge \theta_2}(E) = \sigma_{\theta_1}(\sigma_{\theta_2}(E))$$

- Rule 2: Selection operations are commutative.

$$\sigma_{\theta_1}(\sigma_{\theta_2}(E)) = \sigma_{\theta_2}(\sigma_{\theta_1}(E))$$

- Rule 3: Only the last one in a sequence of projection operations is needed, the others can be omitted.

$$\Pi_{L_1}(\Pi_{L_2}(\dots(\Pi_{L_n}(E))\dots)) = \Pi_{L_1}(E)$$

- Rule 4: Selections can be combined with Cartesian products and theta joins.

- (a).  $\sigma_{\theta}(E_1 \times E_2) = E_1 \bowtie_{\theta} E_2$
- (b).  $\sigma_{\theta_1}(E_1 \bowtie_{\theta_2} E_2) = E_1 \bowtie_{\theta_1 \wedge \theta_2} E_2$

# Equivalence Rules cont'd

- Rule 5: Theta-join operations (and natural joins) are commutative.

$$E_1 \bowtie_{\theta} E_2 = E_2 \bowtie_{\theta} E_1$$

- Rule 6.

- (a) Natural join operations are associative:

$$(E_1 \bowtie E_2) \bowtie E_3 = E_1 \bowtie (E_2 \bowtie E_3)$$

- (b) Theta joins are associative in the following manner:

$$(E_1 \bowtie_{\theta_1} E_2) \bowtie_{\theta_2 \wedge \theta_3} E_3 = E_1 \bowtie_{\theta_1 \wedge \theta_3} (E_2 \bowtie_{\theta_2} E_3)$$

where  $\theta_2$  involves attributes from only  $E_2$  and  $E_3$ .

# Equivalence Rules cont'd

- Rule 7. The selection operation distributes over the theta join operation under the following two conditions:

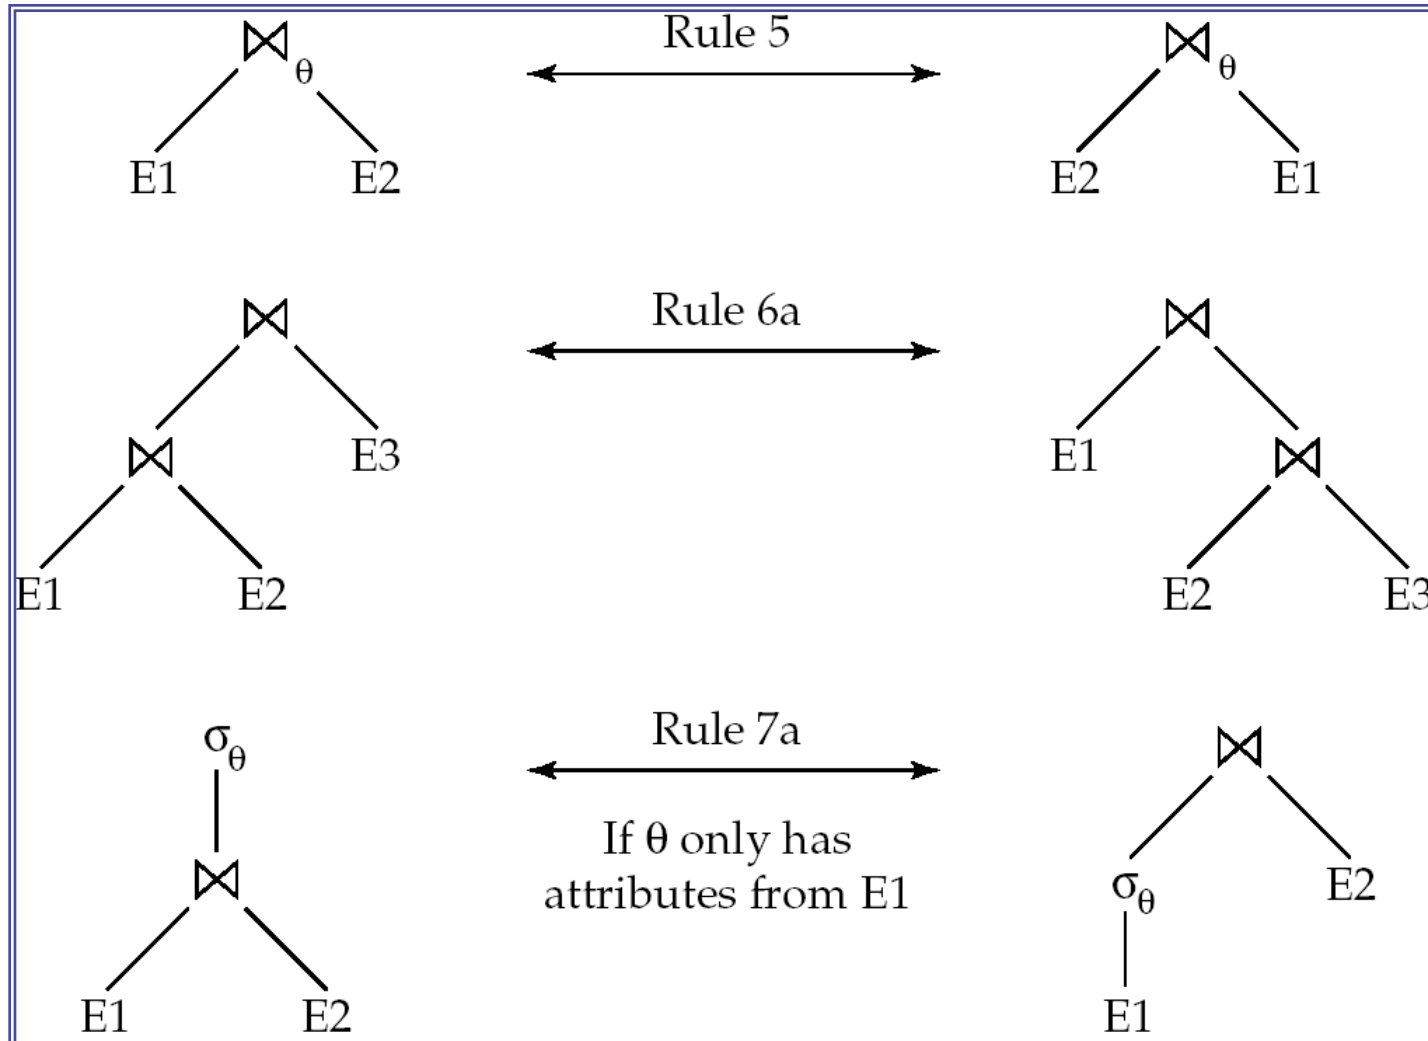
- (a) When  $\theta_0$  involves only the attributes of one of the expressions ( $E_1$ ) being joined.

$$\sigma_{\theta_0}(E_1 \bowtie_{\theta} E_2) = (\sigma_{\theta_0}(E_1)) \bowtie_{\theta} E_2$$

- (b) When  $\theta_1$  involves only the attributes of  $E_1$  and  $\theta_2$  involves only the attributes of  $E_2$ .

$$\sigma_{\theta_1 \wedge \theta_2}(E_1 \bowtie_{\theta} E_2) = (\sigma_{\theta_1}(E_1)) \bowtie_{\theta} (\sigma_{\theta_2}(E_2))$$

# Pictorial Depiction of Equivalence Rules





# Equivalence Rules cont'd

## ■ Rule 8. The projection operation distributes over the theta join operation as follows:

- (a) Let  $L_1$  and  $L_2$  be attributes from  $E_1$  and  $E_2$ , if  $\theta$  involves only attributes from  $L_1 \cup L_2$ :

$$\Pi_{L_1 \cup L_2} (E_1 \bowtie_{\theta} E_2) = (\Pi_{L_1} (E_1)) \bowtie_{\theta} (\Pi_{L_2} (E_2))$$

- (b) Consider a join  $E_1 \bowtie_{\theta} E_2$ .
  - let  $L_1$  and  $L_2$  be sets of attributes from  $E_1$  and  $E_2$ , respectively.
  - let  $L_3$  be attributes of  $E_1$  that are involved in join condition  $\theta$ , but are not in  $L_1 \cup L_2$ , and
  - let  $L_4$  be attributes of  $E_2$  that are involved in join condition  $\theta$ , but are not in  $L_1 \cup L_2$ .

$$\Pi_{L_1 \cup L_2} (E_1 \bowtie_{\theta} E_2) = \Pi_{L_1 \cup L_2} ((\Pi_{L_1 \cup L_3} (E_1)) \bowtie_{\theta} (\Pi_{L_2 \cup L_4} (E_2)))$$



# Equivalence Rules cont'd

- Rule 9. The set operations union and intersection are commutative (set difference is not commutative)

$$\begin{aligned}E_1 \cup E_2 &= E_2 \cup E_1 \\E_1 \cap E_2 &= E_2 \cap E_1\end{aligned}$$

- Rule 10. Set union and intersection are associative.

$$\begin{aligned}(E_1 \cup E_2) \cup E_3 &= E_1 \cup (E_2 \cup E_3) \\(E_1 \cap E_2) \cap E_3 &= E_1 \cap (E_2 \cap E_3)\end{aligned}$$

- Rule 11. The selection operation distributes over  $\cup$ ,  $\cap$  and  $-$ .

$$\sigma_{\theta}(E_1 - E_2) = \sigma_{\theta}(E_1) - \sigma_{\theta}(E_2)$$

Also: 
$$\sigma_{\theta}(E_1 - E_2) = \sigma_{\theta}(E_1) - E_2$$

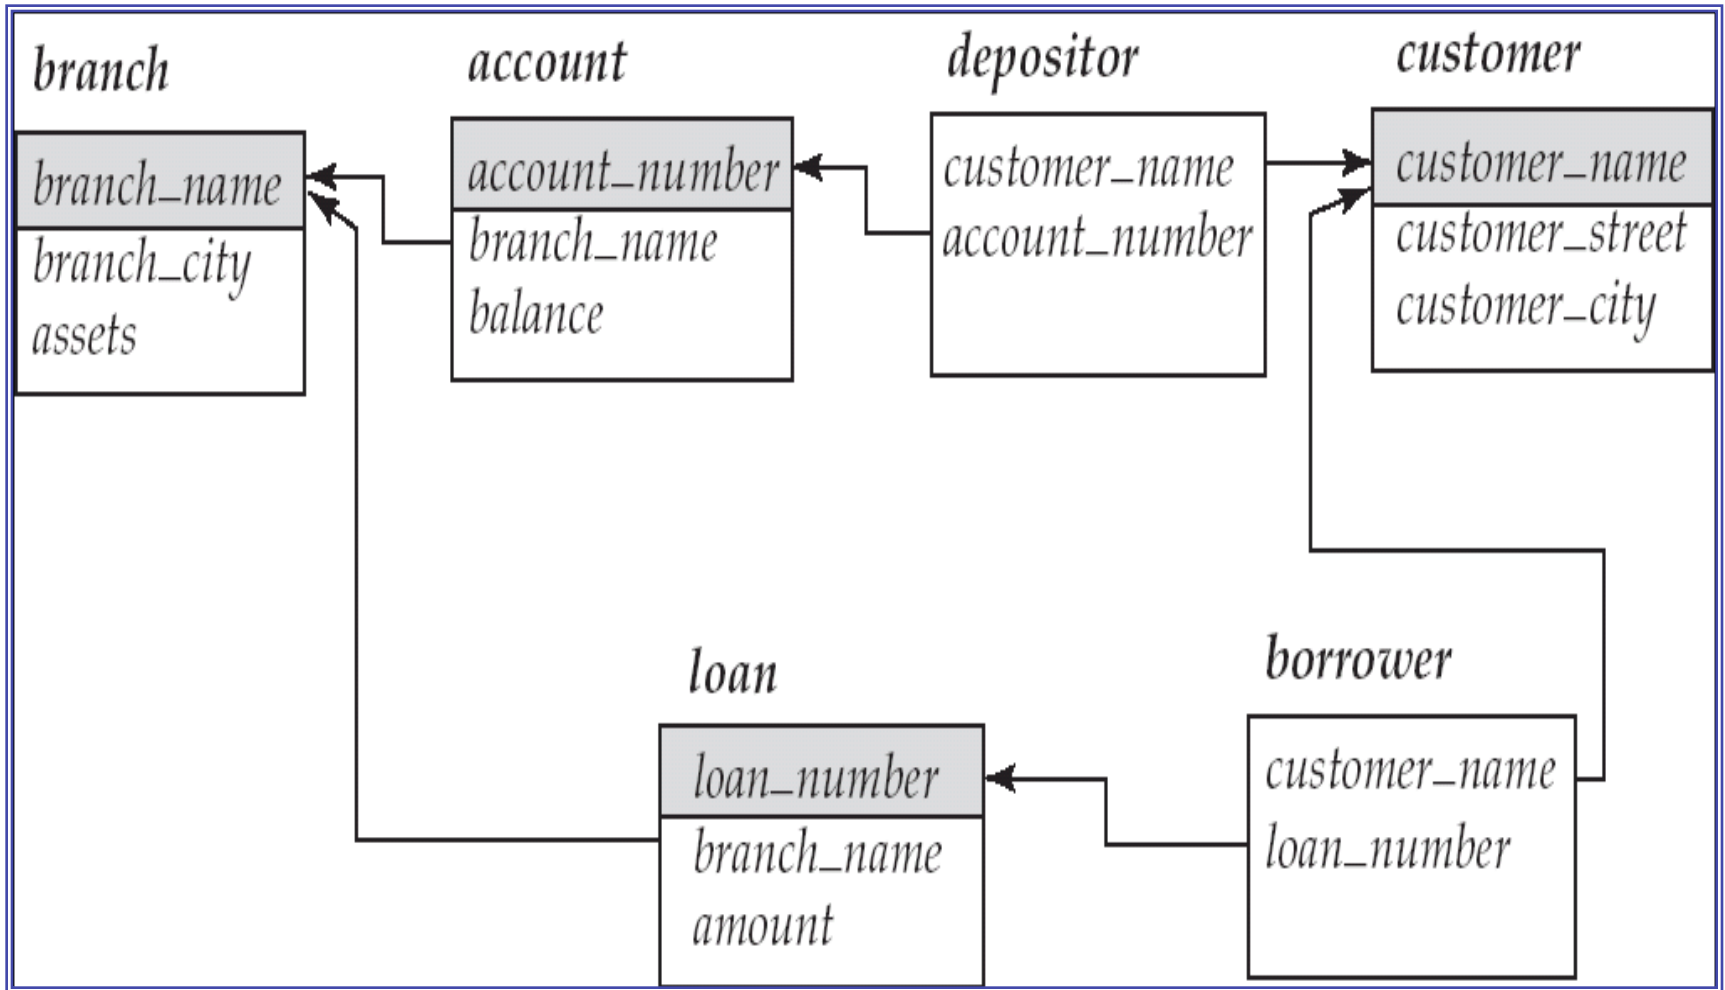
and similarly for  $\cap$  in place of  $-$ , but not for  $\cup$

- Rule 12. The projection operation distributes over union

$$\Pi_L(E_1 \cup E_2) = (\Pi_L(E_1)) \cup (\Pi_L(E_2))$$



# Banking Example



# Example: Pushing Selections

- Query: Find the names of all customers who have an account at some branch located in Brooklyn.

$\Pi_{customer\_name}(\sigma_{branch\_city = "Brooklyn"}(branch \bowtie (account \bowtie depositor)))$

- Transformation using rule 7a (distribute the selection).

$\Pi_{customer\_name}((\sigma_{branch\_city = "Brooklyn"}(branch)) \bowtie (account \bowtie depositor))$

- Performing the selection as early as possible reduces the size of the relation to be joined.

# Example: Multiple Transformations

- Query: Find the names of all customers with an account at a Brooklyn branch whose account balance is over \$1000.

$\Pi_{customer\_name}(\sigma_{branch\_city = "Brooklyn" \wedge balance > 1000} (branch \bowtie (account \bowtie depositor)))$

- Transformation using join associatively (Rule 6a and 7a):


$\Pi_{customer\_name}((\sigma_{branch\_city = "Brooklyn" \wedge balance > 1000} (branch \bowtie account)) \bowtie depositor)$

- Second form provides an opportunity to apply the “**perform selections early**” rule, resulting in the subexpression 7b

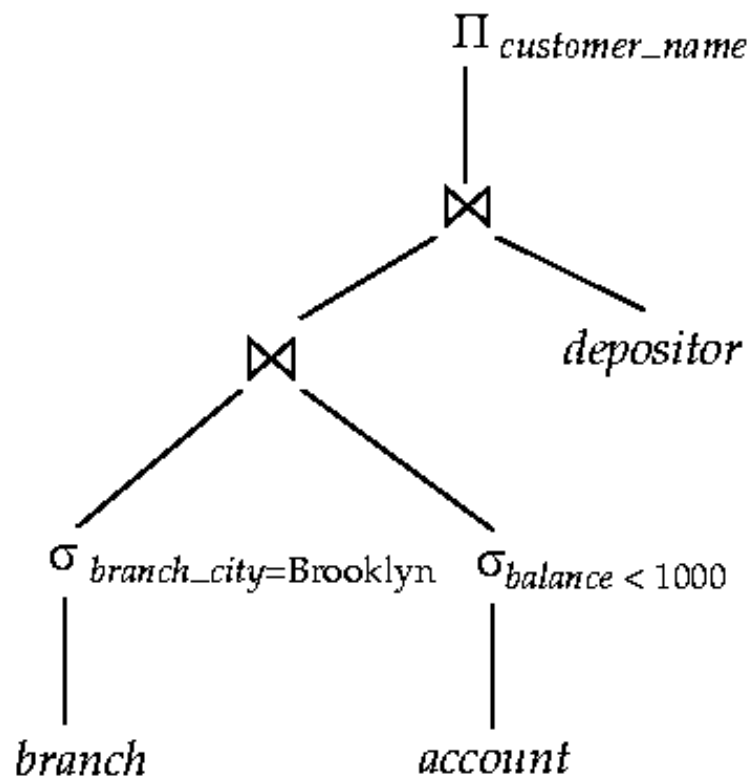
$\sigma_{branch\_city = "Brooklyn"} (branch) \bowtie \sigma_{balance > 1000} (account)$

- Thus a sequence of transformations can be useful

# Multiple Transformations cont'd



(a) Initial expression tree



(b) Tree after multiple transformations

# Transformation Example: Pushing Projections

$\Pi_{customer\_name}((\sigma_{branch\_city = \text{"Brooklyn"}} (branch) \bowtie account) \bowtie depositor)$

- When we compute

$(\sigma_{branch\_city = \text{"Brooklyn"}} (branch) \bowtie account)$

we obtain a relation whose schema is:

$(branch\_name, branch\_city, assets, \text{account\_number}, balance)$

- Push projections using equivalence rules 8b; eliminate unneeded attributes from intermediate results to get:

$\Pi_{customer\_name}((\Pi_{account\_number}(\sigma_{branch\_city = \text{"Brooklyn"}} (branch) \bowtie account)) \bowtie depositor))$

(HINT: L1 is null, L2 is customer\_name; L3=L4=account\_number)

- Performing projection as early as possible reduces the size of the tuples to be joined.

# Join Ordering Example

- For all relations  $r_1, r_2$ , and  $r_3$ ,

$$(r_1 \bowtie r_2) \bowtie r_3 = r_1 \bowtie (r_2 \bowtie r_3)$$

(Join Associativity)

- If  $r_2 \bowtie r_3$  is quite large and  $r_1 \bowtie r_2$  is small, we choose

$$(r_1 \bowtie r_2) \bowtie r_3$$

so that we compute and store a smaller temporary relation.



# Join Ordering Example cont'd

- Consider the expression

$\Pi_{customer\_name} ((\sigma_{branch\_city = \text{"Brooklyn"}}(branch)) \bowtie (account \bowtie depositor))$

- Could compute  $account \bowtie depositor$  first, and join result with

$\sigma_{branch\_city = \text{"Brooklyn"}}(branch)$   
but  $account \bowtie depositor$  is likely to be a large relation.

- Only a small fraction of the bank's customers are likely to have accounts in branches located in Brooklyn
  - it is better to compute first

$\sigma_{branch\_city = \text{"Brooklyn"}}(branch) \bowtie account$

# Enumeration of Equivalent Expressions

- Query optimisers use equivalence rules to systematically generate expressions equivalent to the given expression
- The approach is very expensive in space and time

**procedure** genAllEquivalent( $E$ )

$EQ = \{E\}$

**repeat**

    Match each expression  $E_i$  in  $EQ$  with each equivalence rule  $R_j$

**if** any subexpression  $e_i$  of  $E_i$  matches one side of  $R_j$

        Create a new expression  $E'$  which is identical to  $E_i$ , except that  
         $e_i$  is transformed to match the other side of  $R_j$

        Add  $E'$  to  $EQ$  if it is not already present in  $EQ$

**until** no new expression can be added to  $EQ$

# End of Lecture

- Summary
  - Transformation of Relational Expressions
- Reading
  - Textbook chapter 13.1, 13.2, 13.3, and 13.4