

# **Database Development and Design (CPT201)**

## **Lecture 5b: Introduction to Query Optimisation 2**

Dr. Wei Wang

Department of Computing

# Learning Outcomes

- Introduction to Query Optimisation
  - Catalog Information for Cost Estimation
  - Cost-based optimisation




# Catalog Information for Cost Estimation

- $n_r$ : number of tuples in a relation  $r$ .
- $b_r$ : number of blocks containing tuples of  $r$ .
- $l_r$ : size of a tuple of  $r$ .
- $f_r$ : blocking factor of  $r$ . i.e., the number of tuples of  $r$  that fit into one block.
- $V(A, r)$ : number of **distinct** values that appear in  $r$  for attribute  $A$ ; same as the size of  $\Pi_A(r)$ .
- If tuples of  $r$  are stored together physically in a file, then: 
$$b_r = \left\lceil \frac{n_r}{f_r} \right\rceil$$

# Histograms

- Histogram on attribute *age* of relation *person*
- Equi-width histograms
- Equi-depth histograms



# Estimation of the Size of Selection

- $\sigma_{A=v}(r)$ 
  - $n_r / V(A,r)$ : number of records that will satisfy the selection
  - Equality condition on a **key** attribute (primary key): *size estimate* = 1
- $\sigma_{A \leq v}(r)$  (case of  $\sigma_{A \geq v}(r)$  is symmetric)
  - Let  $c$  denote the estimated number of tuples satisfying the condition. Let  $\min(A,r)$  and  $\max(A,r)$  denote the lowest and highest values for attribute  $A$ .
  - If  $\min(A,r)$  and  $\max(A,r)$  are available in catalog
    - $c = 0$  if  $v < \min(A,r)$
    - $c = n_r \cdot \frac{v - \min(A,r)}{\max(A,r) - \min(A,r)}$
  - If histograms available, can refine above estimate
  - In absence of statistical information  $c$  is assumed to be  $n_r / 2$ .

# Estimation of the Size of Joins

- The Cartesian product  $r \times s$  contains  $n_r \cdot n_s$  tuples; each tuple occupies  $s_r + s_s$  bytes.
- If  $R \cap S = \emptyset$ , then  $r \bowtie s$  is the same as  $r \times s$ .
- If  $R \cap S$  is a key for  $R$ , then a tuple of  $s$  will join with at most one tuple from  $r$ 
  - therefore, the number of tuples in  $r \bowtie s$  is no greater than the number of tuples in  $s$ .
- If  $R \cap S$  is a foreign key in  $S$  referencing  $R$ , then the number of tuples in  $r \bowtie s$  is exactly the same as the number of tuples in  $s$ .
  - The case for  $R \cap S$  being a foreign key referencing  $S$  is symmetric.
- In the example query *depositor*  $\bowtie$  *customer*, *customer\_name* in *depositor* is a foreign key (of *customer*)
  - hence, the result has exactly  $n_{\text{depositor}}$  tuples, which is 5000



# Estimation of the Size of Joins

## cont'd

- If  $R \cap S = \{A\}$  is not a key for  $R$  or  $S$ .

If we assume that every tuple  $t$  in  $R$  produces tuples in  $R \bowtie S$ , the number of tuples in  $R \bowtie S$  is estimated to be:

$$\frac{n_r * n_s}{V(A, s)}$$

If the reverse is true, the estimate obtained will be:

$$\frac{n_r * n_s}{V(A, r)}$$

The **lower of these two estimates** is probably the more accurate one.

- Can improve on above if histograms are available
  - Use formula similar to above, for each cell of histograms on the two relations

# Join Operation: Running Example

- Running example: *depositor* ⋈ *customer*
- Catalog information for join examples:
  - $n_{customer} = 10,000$ .
  - $f_{customer} = 25$ , which implies that  $b_{customer} = 10,000/25 = 400$ .
  - $n_{depositor} = 5000$ .
  - $f_{depositor} = 50$ , which implies that  $b_{depositor} = 5,000/50 = 100$ .
  - $V(customer\_name, depositor) = 2,500$ , which implies that, on average, each customer has two accounts.
    - Also assume that *customer\_name* in *depositor* is a foreign key on *customer*.
    - $V(customer\_name, customer) = 10,000$  (primary key)



# Join Operation: Running Example cont'd

- Compute the size estimates for *depositor* ⋈ *customer* without using information about foreign keys:
  - $V(\text{customer\_name}, \text{depositor}) = 2,500$ , and  
 $V(\text{customer\_name}, \text{customer}) = 10,000$
  - The two estimates are  $5,000 * 10,000 / 2,500 = 20,000$   
and  $5,000 * 10,000 / 10,000 = 5,000$
  - We choose the lower estimate, which in this case, is the same as our earlier computation using foreign keys.

# Size Estimation for Other Operations

- Projection: estimated size of  $\Pi_A(r) = V(A,r)$
- Set operations
  - For unions/intersections of selections on the **same** relation: rewrite and use size estimate for selections
    - e.g.,  $\sigma_{\theta_1}(r) \cup \sigma_{\theta_2}(r)$  can be rewritten as  $\sigma_{\theta_1 \vee \theta_2}(r)$
  - For operations on **different** relations:
    - estimated size of  $r \cup s$  = size of  $r$  + size of  $s$ .
    - estimated size of  $r \cap s$  = minimum size of  $r$  and size of  $s$ .
    - estimated size of  $r - s$  =  $r$ .
    - All the three estimates may be quite **inaccurate**, but provide **upper bounds** on the sizes.

# Estimation of Number of Distinct Values in Selection

- If  $\theta$  forces  $A$  to take a specified value:  $V(A, \sigma_{\theta}(r)) = 1$ .
  - e.g.,  $A = 3$
- If  $\theta$  forces  $A$  to take on one of a specified set of values:
$$V(A, \sigma_{\theta}(r)) = \text{number of specified values.}$$
  - (e.g.,  $(A = 1 \vee A = 3 \vee A = 4)$ ),
- If the selection condition  $\theta$  is of the form  $A \text{ op } v$  ( $\text{op}$  is  $>$ ,  $<$ , etc),
$$V(A, \sigma_{\theta}(r)) = V(A, r) * s$$
  - where  $s$  is the selectivity of the selection.
- In all the other cases: use approximate estimate of  $\min(V(A, r), n_{\sigma_{\theta}(r)})$

# Estimation of Distinct Values cont'd

Joins:  $r \bowtie s$

- If all attributes in  $A$  are from  $r$ ,  
estimated  $V(A, r \bowtie s) = \min(V(A, r), n_{r \bowtie s})$
- If  $A$  contains attributes  $A1$  from  $r$  and  $A2$  from  $s$ , then  
estimated  
 $V(A, r \bowtie s) =$   
 $\min(V(A1, r) * V(A2 - A1, s), V(A1 - A2, r) * V(A2, s), n_{r \bowtie s})$ 
  - More accurate estimate can be got using probability theory, but  
this one works fine generally
- Projections: Estimation of distinct values are  
straightforward for projections.
  - They are the same in  $\Pi_A(r)$  as in  $r$ .

# Choice of Evaluation Plans

- Must consider the **interaction** of evaluation techniques when choosing evaluation plans
  - choosing the cheapest algorithm for each operation independently may not yield best overall algorithm, e.g.
    - merge-join may be costlier than hash-join, but may provide a sorted output which reduces the cost for an outer level aggregation.
    - nested-loop join may provide opportunity for **pipelining**
- Practical query optimisers incorporate elements of the following two broad approaches:
  - Search all the plans and choose the best plan in a cost-based fashion.
  - Uses heuristics to choose a plan.

# Cost-Based Join Order Optimisation

- Consider finding the best join-order for

$$r_1 \bowtie r_2 \bowtie \dots \bowtie R_n.$$

- There are  $(2(n-1))!/(n-1)!$  different join orders for above expression. With  $n = 7$ , the number is 665280, with  $n = 10$ , the number is greater than 176 billion!
- No need to generate all the join orders. Using **dynamic programming**, the least-cost join order for any subset of  $\{r_1, r_2, \dots, r_n\}$  is computed only once and stored for future use.

# Dynamic Programming in Optimisation

- To find best plan (join tree) for a set of  $n$  relations:
  - Consider all possible plans of the form:  $S_1 \bowtie (S - S_1)$ , where  $S_1$  is any non-empty subset of  $S$ .
  - Recursively compute cost for joining subsets of  $S$  to find the cost of each plan. Choose the cheapest of the alternatives.
  - Base case for recursion: single relation access plan
    - Find the best selection strategy for a particular relation  $R_i$
  - When plan for any subset is computed, store it and reuse it when it is required again, instead of re-computing it.

# Join Order Optimisation Algorithm

```
// initialise bestplan[S].cost to  $\infty$ 
procedure findbestplan(S)
  if (bestplan[S].cost  $\neq \infty$ )
    return bestplan[S]
  // else bestplan[S] has not been computed earlier, compute it now
  if (S contains only 1 relation)
    set bestplan[S].plan and bestplan[S].cost based on the best way
    of accessing S /* Using selections on S and indices on S */
  else for each non-empty subset S1 of S such that S1  $\neq$  S
    P1= findbestplan(S1)
    P2= findbestplan(S - S1)
    A = best algorithm for joining results of P1 and P2
    cost = P1.cost + P2.cost + cost of A
    if cost < bestplan[S].cost
      bestplan[S].cost = cost
      bestplan[S].plan = "execute P1.plan; execute P2.plan;
                        join results of P1 and P2 using A"
  return bestplan[S]
```





# Cost of Join Order Optimisation

- With dynamic programming time complexity of optimisation with bushy trees is  $O(3^n)$ .
  - With  $n = 10$ , this number is 59000 instead of 176 billion!
- Space complexity is  $O(2^n)$  as the number of subsets of the  $S$  is  $2^n$ .
- Although both numbers still increase rapidly with  $n$ , commonly occurring joins usually have less than 10 relations, and can be handled easily.

# Cost-Based Optimisation with Equivalence Rules

- Many optimisers follow an approach based on
  - Using heuristic transformations to handle constructs other than joins
  - applying the cost-based join order selection algorithm to subexpressions involving only joins and selections
- General-purpose cost-based optimiser based on equivalence rules
  - easy to extend the optimiser with new rules to handle different query constructs
  - but the procedure to enumerate all equivalent expressions is very expensive

# Cost-Based Optimisation with Equivalence Rules cont'd


- To make the approach work efficiently requires the following:
  - A space-efficient representation of expressions
  - Efficient techniques for detecting duplicate derivations of the same expression
  - dynamic programming based on memoisation
  - avoid generating all possible equivalent plans

# Heuristic Optimisation

- Cost-based optimisation is expensive, even with dynamic programming.
- Systems may use *heuristics* to reduce the number of choices that must be made in a cost-based fashion.
- Heuristic optimisation transforms the query-tree by using a set of rules that typically (but not in all cases) improve execution performance:
  - Perform selection early (reduces the number of tuples)
  - Perform projection early (reduces the number of attributes)
  - Perform the most restrictive selection and join operations (i.e., with smallest result size) before other similar operations.
- Some systems use only heuristics, others combine heuristics with partial cost-based optimisation.

# Other heuristics: Left Deep Join Trees

- In **left-deep join trees**, the right-hand-side input for each join is a relation, not the result of an intermediate join.



# Cost of left-deep join Optimisation

- To find best left-deep join tree for a set of  $n$  relations:
  - Consider  $n$  alternatives with one relation as right-hand side input and the other relations as left-hand side input.
  - Modify optimisation algorithm:
    - Replace "for each non-empty subset  $S_1$  of  $S$  such that  $S_1 \neq S$ "
    - By: for each relation  $r$  in  $S$ , let  $S_1 = S - r$ .
- If only left-deep trees are considered, time complexity of finding best join order is  $O(n!)$ , with dynamic programming this can be reduced to  $O(n 2^n)$ 
  - Space complexity remains at  $O(2^n)$
- Cost-based optimisation is expensive, but worthwhile for queries on large datasets (typical queries have small  $n$ , generally  $< 10$ )

# Structure of Query Optimisers

- Many optimisers considers only left-deep join orders.
  - Plus heuristics to push selections and projections down the query tree
  - Reduces optimisation complexity and generates plans amenable to pipelined evaluation.
- Heuristic optimisation used in some versions of Oracle:
  - Repeatedly pick “best” relation to join next
    - Starting from each of n starting points. Pick best among these

# Structure of Query Optimisers

## cont'd

- Some query optimisers integrate heuristic selection and the generation of alternative access plans.
  - Frequently used approach
    - heuristic rewriting of nested block structure and aggregation
    - followed by cost-based join-order optimisation for each block
  - Some optimisers (e.g. SQL Server) apply transformations to entire query and do not depend on block structure
- Even with the use of heuristics, cost-based query optimisation imposes a substantial overhead.
  - But is worth for expensive queries
  - Optimisers often use simple heuristics for very cheap queries, and perform exhaustive enumeration for more expensive queries



# End of Lecture

## ■ Summary

- Transformation of Relational Expressions
- Catalog Information for Cost Estimation
- Cost-based optimisation
- Dynamic Programming for Choosing Evaluation Plans

## ■ Reading

- Textbook chapter 13.1, 13.2, 13.3, and 13.4