# Database Development and Design (CPT201)

## Lecture 6a: Transaction ManagementConcepts of Transaction

Dr. Wei Wang
Department of Computing

#### Learning Outcomes

- Transaction Concept: Atomicity,
   Consistency, Isolation, Durability
- Concurrent Executions and Schedule
  - Serialisability
  - Recoverability
  - Testing for Serialisability



#### **Transaction Concept**

- A transaction is a unit of program execution that accesses and possibly updates various data items.
  - A transaction must see a consistent database.
  - During transaction execution the database may be temporarily inconsistent.
  - When the transaction completes successfully (is committed), the database must be consistent.
  - After a transaction commits, the changes it has made to the database persist, even if there are system failures.
- Two main issues to deal with:
  - Concurrent execution of multiple transactions
  - Recovery from failures of various kinds, such as hardware failures and system crashes



#### **ACID Properties**

- Atomicity: Either all operations of the transaction are properly reflected in the database or none are.
- Consistency: Execution of a transaction in isolation preserves the consistency of the database.
- Isolation: Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions.
  - That is, for every pair of transactions  $T_i$  and  $T_j$ , it appears to  $T_i$  that either  $T_j$ , finished execution before  $T_i$  started, or  $T_j$  started execution after  $T_i$  finished.
- Durability: After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.



## Example of Fund Transfer

The following is a transaction to transfer \$50 from account A to account B:

- 1. read(A)
- 2. A := A 50
- 3. write(A)
- 4. read(B)
- 5. B := B + 50
- 6. **write**(*B*)





#### Example of Fund Transfer cont'd

- Atomicity requirement if the transaction fails after step 3 and before step 6, the system should ensure that its updates are not reflected in the database, else inconsistency will result.
- Consistency requirement the sum of A and B is unchanged by the execution of the transaction.

- 1. read(A)
- 2. A := A 50
- 3. **write**(*A*)
- 4. **read**(*B*)
- 5. B := B + 50
- 6. **write**(*B*)





#### Example of Fund Transfer cont'd

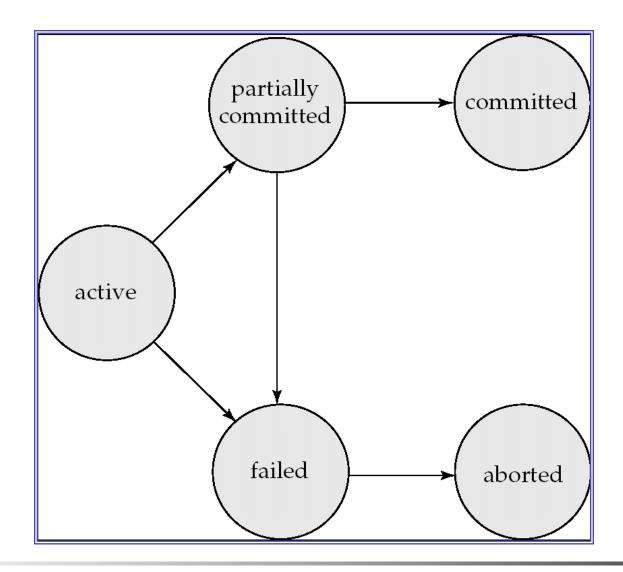
- Isolation requirement if between steps 3 and 6, another transaction is allowed to access the partially updated database, it will see an inconsistent database (the sum A + B will be less than it should be).
- Durability requirement once the user has been notified that the transaction has completed (i.e., the transfer of the \$50 has taken place), the updates to the database by the transaction must persist despite failures.

- 1. **read**(*A*)
- 2. A := A 50
- 3. **write**(*A*)
- 4. **read**(*B*)
- 5. B := B + 50
- 6. **write**(*B*)





#### **Transaction States**





#### **Transaction State**

- Active the initial state; the transaction stays in this state while it is executing
- Partially committed after the final statement has been executed.
- Failed –normal execution can no longer proceed.
- Aborted after the transaction has been rolled back and the database restored to its state prior to the start of the transaction.
- Committed after successful completion.



#### **Concurrent Executions**

- Multiple transactions are allowed to run concurrently in the system. Advantages are:
  - increased processor and disk utilisation, leading to better transaction throughput: one transaction can be using the CPU while another is reading from or writing to the disk
  - reduced average response time for transactions: short transactions need not wait behind long ones.
- Concurrency control schemes mechanisms to achieve isolation; that is, to control the interaction among the concurrent transactions in order to prevent them from destroying the consistency of the database.



#### Schedules

- Schedule a sequence of instructions that specifies the chronological order in which concurrent transactions are executed
  - a schedule for a set of transactions must consist of all instructions of those transactions
  - must preserve the order in which the instructions appear in each individual transaction.
- A transaction that successfully completes its execution will have a commit instruction as the last statement (will be omitted if it is obvious)
- A transaction that fails to successfully complete its execution will have an abort instruction as the last statement (will be omitted if it is obvious)

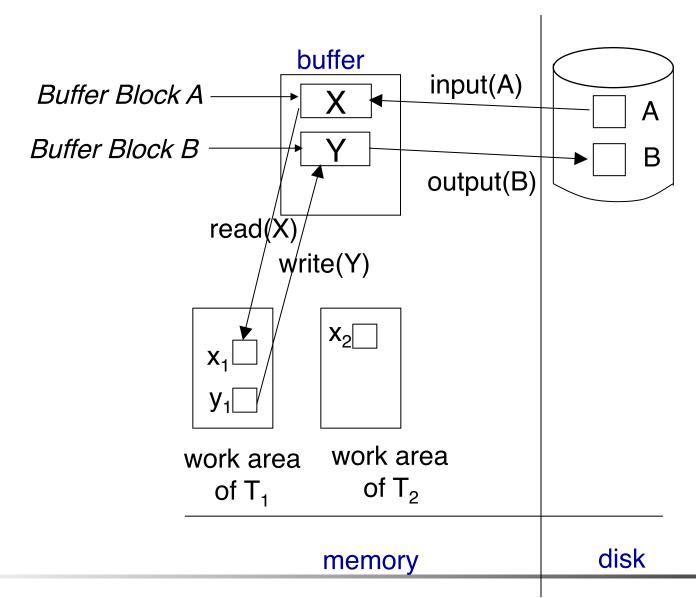


#### **Data Access**

- Physical blocks are those blocks residing on the disk.
- Buffer blocks are the blocks residing temporarily in main memory.
- Block movements between disk and main memory are initiated through the following two operations:
  - input(B) transfers the physical block B to main memory (buffer).
  - output(B) transfers the buffer block B to the disk, and replaces the appropriate physical block there.
- We assume, for simplicity, that each data item fits in, and is stored inside, a single block.



#### Example of Data Access





## Example of Schedules

Let A=1,000, B=2,000. Let  $T_1$  transfer \$50 from A to B, and  $T_2$  transfer 10% of the balance from A to B.

$T_1$	T2
read(A)	
A := A - 50	
write $(A)$	
read(B)	
B := B + 50	
write(B)	
	read(A)
	temp := A * 0.1
	A := A - temp
	write(A)
	read(B)
	B := B + temp
	write(B)

$T_1$	$T_2$
	read(A)
	temp := A * 0.1
	A := A - temp
	write(A)
	read(B)
	B := B + temp
	write(B)
read(A)	
A := A - 50	
write(A)	
read(B)	
B := B + 50	
write(B)	

1

1 and 2 are serial schedules;

How much is A, B, and A+B?





## Example of Schedules cont'd

Let  $T_1$  transfer \$50 from A to B, and  $T_2$  transfer 10% of the balance from A to B.

$T_1$	$T_2$
read(A)	
A := A - 50	
write(A)	
	read(A)
	temp := A * 0.1
	A := A - temp
	write(A)
read(B)	
B := B + 50	
write(B)	
	read(B)
	B := B + temp
	write(B)

$T_1$	$T_2$
read(A)	
A := A - 50	
	read(A)
	temp := A * 0.1
	A := A - temp
	write(A)
	read(B)
write(A)	
read(B)	
B := B + 50	
write(B)	
	B := B + temp
	write(B)

3

4

- 3 and 4 are not serial schedules;
- Does 3 preserve the value of (A+B)? How much is A and B?
- Does 4 preserve the value of (A+B)? How much is A and B?



## Serialisability

- Basic Assumption Each transaction preserves database consistency.
- A (possibly concurrent) schedule is serialisable if it is equivalent to a serial schedule. Different forms of schedule equivalence give rise to the notions of:
  - conflict serialisability
  - view serialisability
- We ignore operations other than read and write instructions for now for simplicity.
- We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes.



#### Conflicting Instructions

- Instructions  $l_i$  and  $l_j$  of transactions  $T_i$  and  $T_j$ , conflict if and only if there exists some data item Q accessed by both  $l_i$  and  $l_j$ , and at least one of these instructions is write Q.
- Three conflicts can be identified
  - write-read (WR) conflict: reading uncommitted data
  - read-write (RW) conflict: unrepeatable reads
  - write-write (WW) conflict: overwriting uncommitted data
- Intuitively, a conflict between  $l_i$  and  $l_j$  forces a (logical) temporal order between them.
  - If  $l_i$  and  $l_j$  are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule (the order of the two instructions can be swapped).



## Write-Read (WR) Conflicts

 Transaction T2 reads a database object that has been modified by transaction T1 which has not committed ("dirty read")

T1	T2
read(A) $A := A - 50$ $write(A)$	read( $A$ ) $A := A + A*10\%$ write( $A$ ) read( $B$ ) $B := B + B*10\%$ write( $B$ ) commit
read( $B$ ) $B := B + 50$ write( $B$ ) Rollback	Commit





## Read-Write (RW) Conflicts

- Transaction T2 could change the value of an object that has been read by a transaction T1, while T1 is still in progress. (unrepeatable read)
- "unrepeatable read" example

T1	T2
read(A)	read(A) $write(A)$ $commit$
read(A) write(A) commit	

T1	<b>T2</b>
read(A) A:=A-1 write(A) commit	read(A) A:=A-1 write(A) commit

Let A=1 and assume that there is an integrity constraint that prevents A from becoming negative. T1 will get an error.





#### Write-Write (WW) Conflicts

- Transaction T2 could overwrite the value of an object which has already been modified by T1, while T1 is still in progress. (Blind Write)
- "Blind Write" example
  - T1: set both Steven and Paul salaries at USD 1m.
  - T2: set both Steven and Paul salaries at RMB 1m.
- With this schedule, Steven has salary
   1m RMB and Paul has salary 1m USD.

T1	T2
write(Steven)	write(Paul) write(Steven) commit
write(Paul) commit	





## **Conflict Serialisability**

If a schedule S can be transformed into a schedule S' by a series of swaps of non-conflicting instructions, we say that S and S' are conflict equivalent. A schedule S is conflict serialisable if it is conflict equivalent to a serial schedule.

$T_1$	$T_2$
read(A)	
write(A)	
	read(A)
	write(A)
read(B)	
write(B)	
	read(B)
	write(B)

swaps of non-conflicting instructions.

$T_1$	$T_2$
read(A)	
write(A)	
read(B)	
write(B)	
	read(A)
	write(A)
	read(B)
	write(B)



6

3



## Conflict Serialisability cont'd

- Example of a schedule that is not conflict serialisable
- We are unable to swap instructions in the schedule to obtain either the serial schedule  $\langle T_3, T_4 \rangle$ , or the serial schedule  $\langle T_4, T_3 \rangle$ .

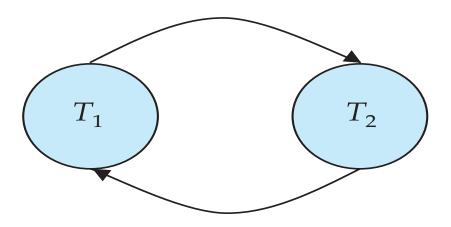
$T_3$	$T_4$
read(Q)	
	write(Q)
write(Q)	





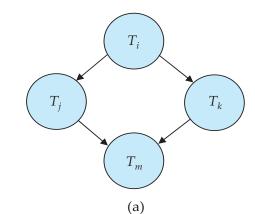
## Testing for Serialisability

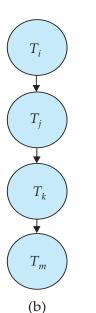
- Consider a schedule with a set of transactions  $T_1$ ,  $T_2$ , ...,  $T_n$
- Precedence graph a directed graph where
  - vertices are the transactions (names).
  - an arc (edge) is drawn from  $T_i$  to  $T_j$  if the two transactions conflict, and  $T_i$  accesses the same data item before  $T_j$  does.

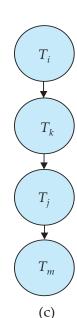


#### Test for Conflict Serialisability

- A schedule is conflict serialisable if and only if its precedence graph is acyclic.
- Cycle-detection algorithms exist which take order  $n^2$  time, where n is the number of vertices in the graph.
  - (Better algorithms take order n + e where e is the number of edges.)
- If precedence graph is acyclic, the serialisability order can be obtained by a topological sorting of the graph.
  - This is a linear order consistent with the partial order of the graph.
  - Could have more than one serialisability orders
- It is possible for two schedules to produce the same outcome, but are not conflict serialisable.









## View Serialisability

- Let S and S' be two schedules with the same set of transactions.
   S and S' are view equivalent if the following three conditions are met, for each data item Q,
  - If in schedule S, transaction  $T_i$  reads the initial value of Q, then in schedule S'also transaction  $T_i$  must read the initial value of Q.
  - If in schedule S transaction  $T_i$  executes read(Q), and that value was produced by transaction  $T_j$  (if any), then in schedule S' also transaction  $T_i$  must read the value of Q that was produced by the same write(Q) operation of transaction  $T_j$ .
  - The transaction (if any) that performs the final write(Q) operation in schedule S must also perform the final write(Q) operation in schedule S'.
- As can be seen, view equivalence is also based purely on reads and writes alone.



#### View Serialisability cont'd

- A schedule 5 is view serialisable if it is view equivalent to a serial schedule.
- Every conflict serialisable schedule is also view serialisable but not vice versa.
  - Below is a schedule which is view-serialisable but not conflict serialisable.

$T_{27}$	$T_{28}$	$T_{29}$
read (Q)		
write (Q)	write (Q)	
		write (Q)

- What serial schedule is above equivalent to?
  - Answer: <T27, T28, T29>
- Every view serialisable schedule that is not conflict serialisable has blind writes.



## Test for View Serialisability

- The precedence graph test for conflict serialisability cannot be used directly to test for view serialisability.
  - Extension to test for view serialisability has cost exponential in the size of the precedence graph.
- The problem of checking if a schedule is view serialisable falls in the class of NP-complete problems.
  - Thus existence of an efficient algorithm is extremely unlikely.
- However practical algorithms that just check some sufficient conditions for view serialisability can still be used.
- It is not used in practice due to its high degree of computational complexity



#### Recoverable Schedules

- Recoverable schedule if a transaction  $T_j$  reads a data item previously written by a transaction  $T_i$ , then the commit operation of  $T_i$  appears before the commit operation of  $T_j$ .
- $\, \blacksquare \,$  The following schedule (Schedule 11) is not recoverable as  $T_7$  commits immediately after the read

$T_6$	$T_7$
read(A) write(A)	
1 5979 10	read(A)
read(B)	Commit

- If  $T_6$  should abort,  $T_7$  would have read (and possibly shown to the user) an inconsistent database state.
- Hence, database must ensure that schedules are recoverable.





## Cascading Rollbacks

- Cascading rollback a single transaction failure leads to a series of transaction rollbacks. Consider the following schedule where none of the transactions has yet committed.
- If  $T_{10}$  fails (aborted later and needs to roll back),  $T_{11}$  and  $T_{12}$  must also be rolled back.
- Can lead to the undoing of a significant amount of work

$T_{10}$	$T_{11}$	$T_{12}$
read(A)		
read(B)		
write(A)		
	read(A)	
	write(A)	
		read(A)





#### Cascadeless Schedules

- Cascadeless schedules cascading rollbacks cannot occur; for each pair of transactions  $T_i$  and  $T_j$  such that  $T_j$  reads a data item previously written by  $T_i$ , the commit operation of  $T_i$  appears before the read operation of  $T_i$ .
- Every cascadeless schedule is also recoverable.
- It is desirable to restrict the schedules to those that are cascadeless.



#### End of Lecture

#### Summary

- Transactions and the ACID properties
- Schedules and conflict serialisable schedules
- Potential anomalies with interleaving

#### Reading

- Textbook 6<sup>th</sup> edition, chapter 14.1, 14.2, 14.3, 14.4, 14.5, 14.6 and 14.7
- Textbook 7<sup>th</sup> edition, chapter 17.1, 17.2, 17.3, 17.4, 17.5, 17.6 and 17.7

