

Database Development and Design (CPT201)

Lecture 9: Database Connectivity

Dr. Wei Wang
Department of Computing

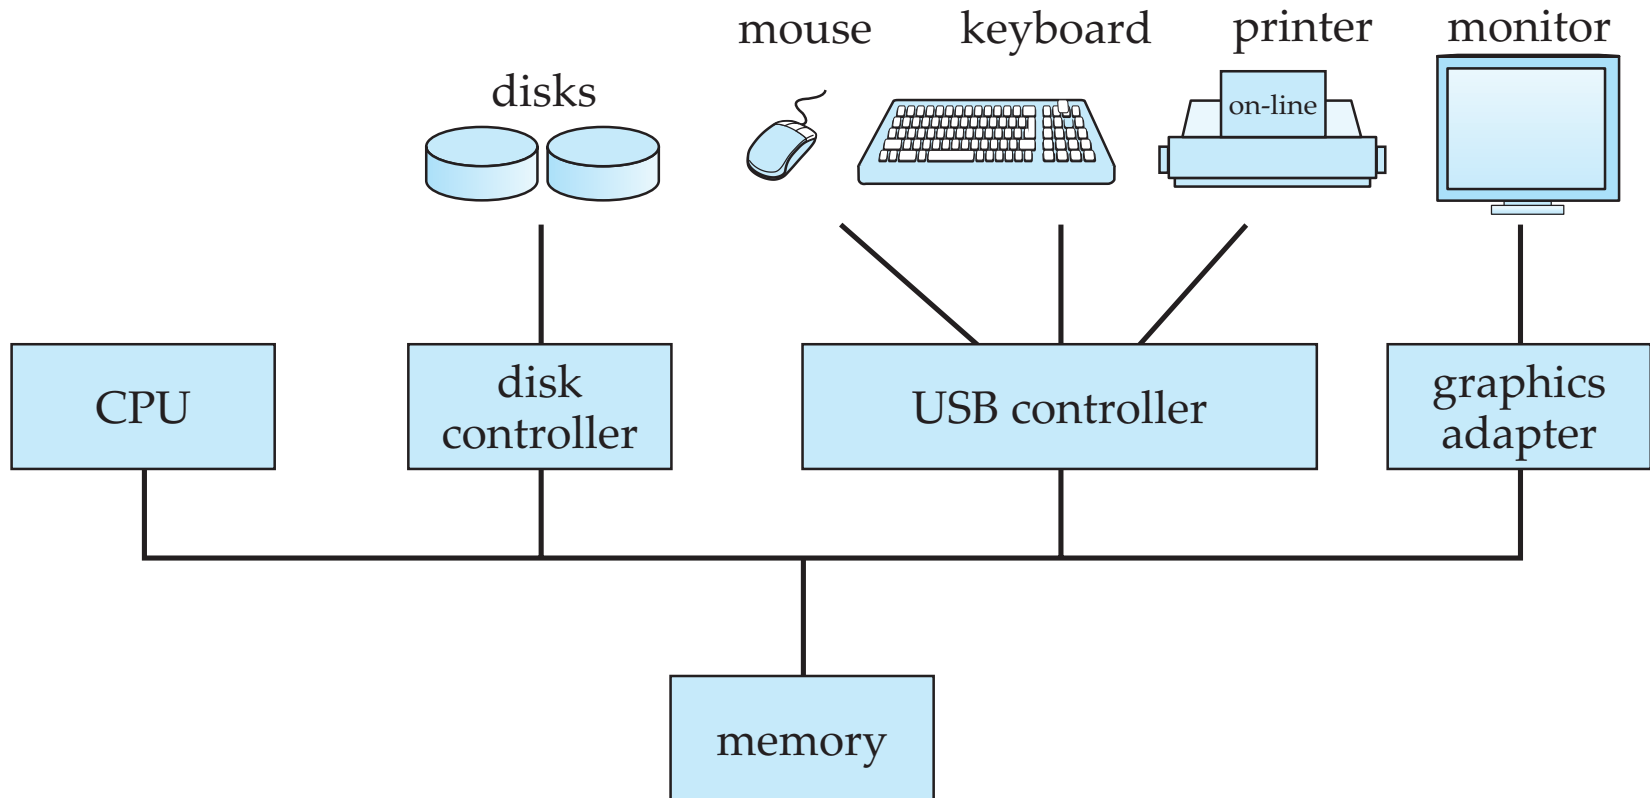
Learning Outcomes

- Centralised and Client-Server Systems
- Server System Architectures
- Database Connectivity
 - ODBC
 - JDBC
 - Java Programming with JDBC

Centralised Systems


- Run on a single computer system and do not interact with other computer systems.
 - **General-purpose computer system**: one to a few CPUs and a number of device controllers that are connected through a common bus that provides access to shared memory.
 - **Single-user system** (e.g., personal computer or workstation): desk-top unit, single user, usually has only one CPU and one or two hard disks; the OS may support only one user.
 - **Multi-user system**: more disks, more memory, multiple CPUs, and a multi-user OS. Serve a large number of users who are connected to the system via terminals. Often called *server systems*.

A Centralised Computer System



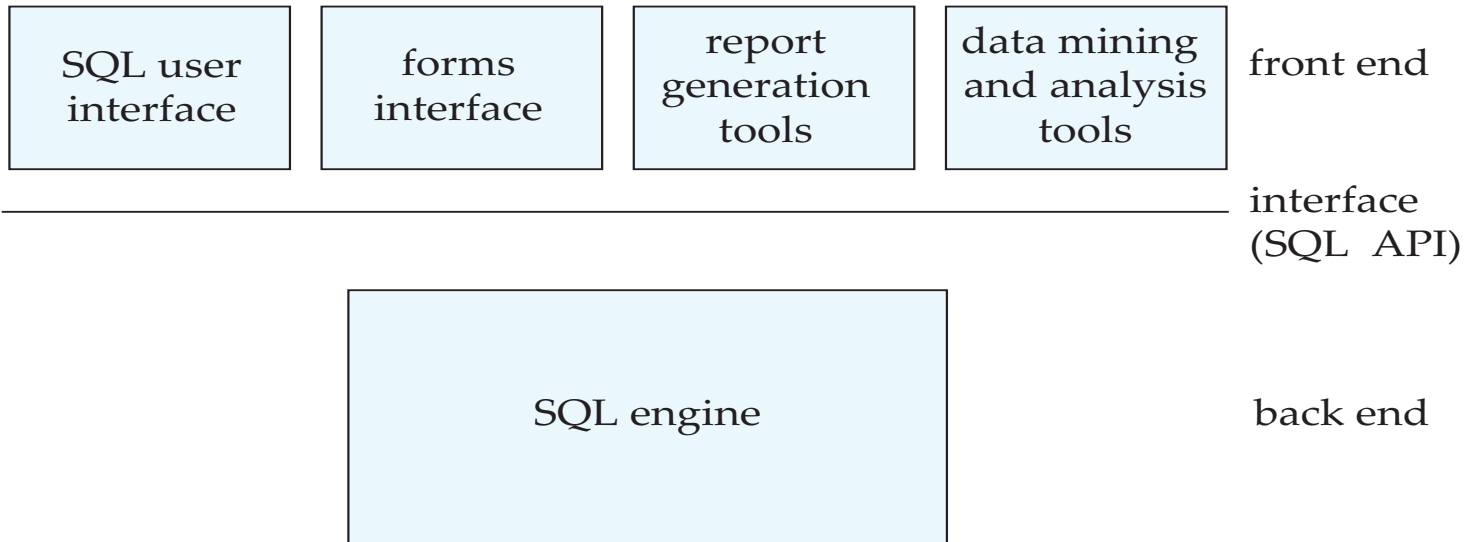
Client-Server Systems

- Server systems satisfy requests generated at client systems.



Client-Server Systems cont'd

- Database functionality can be divided into:
 - **Back-end**: manages access structures, query evaluation and optimisation, concurrency control and recovery, etc.
 - **Front-end**: consists of tools such as *forms*, *report-writers*, and graphical user interface facilities, etc.
- The interface between the front-end and the back-end is through SQL or through an application program interface.



Client-Server Systems cont'd

- Advantages of replacing mainframes with networks of workstations or personal computers connected to **back-end server** machines:
 - better functionality for the cost
 - flexibility in locating resources and expanding facilities
 - better user interfaces
 - easier maintenance

Server System Architecture

- Server systems can be broadly categorised into two kinds:
 - **transaction servers** widely used in relational database systems
 - **data servers** used in object-oriented database systems

Transaction Servers

- Also called **query server** systems or *SQL server systems*
 - Clients send requests to the server,
 - Transactions are executed at the server,
 - Results are shipped back to the client.
- Requests are specified in SQL, and communicated to the server through a **remote procedure call (RPC)** mechanism.
- Transactional RPC allows many RPC calls to form a transaction.
- **Open Database Connectivity (ODBC)** is a C language application program interface standard from Microsoft for connecting to a server, sending SQL requests, and receiving results.
- **JDBC** standard is similar to ODBC, but for Java.

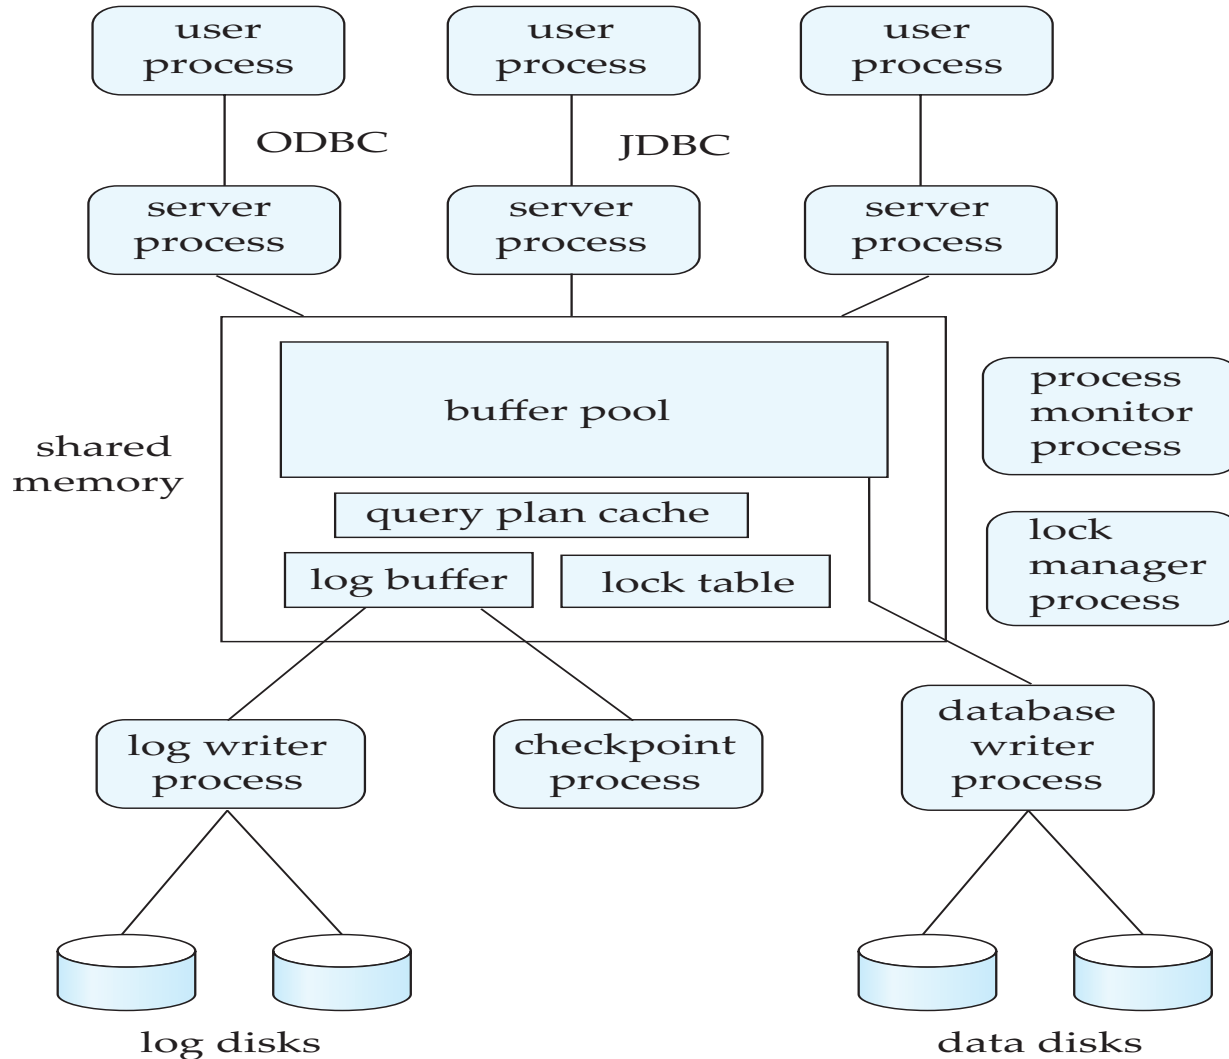
Transaction Server Process Structure

- A typical transaction server consists of multiple processes accessing data in shared memory.
- Server processes
 - receive user queries (transactions), execute them and send results back
 - may be **multithreaded** to support multiple user queries concurrently
 - may be multiple multithreaded server processes
- Lock manager process
- Database writer process
 - Output modified buffer blocks to disks continually

Transaction Server Processes cont'd

- Log writer process
 - Server processes simply add log records to log record buffer
 - Log writer process outputs log records to stable storage.
- Checkpoint process
 - Performs periodic checkpoints
- Process monitor process
 - Monitors other processes, and takes recovery actions if any of the other processes fail
 - e.g., aborting any transactions being executed by a server process and restarting it

Transaction System Processes (Cont.)



Transaction System Processes cont'd

- Shared memory contains shared data
 - Buffer pool
 - Lock table
 - Log buffer
 - Cached query plans (reused if same query submitted again)
- All database processes can access shared memory
- To ensure that no two processes are accessing the same data structure at the same time, databases systems implement **mutual exclusion** using either
 - Operating system semaphores
 - Atomic instructions such as test-and-set
- To avoid overhead of inter-process communication for lock request/grant, each database process operates directly on the lock table
 - instead of sending requests to lock manager process
- Lock manager process still used for deadlock detection

Data Servers

- Used in high-speed LANs, in cases where
 - The clients are comparable in processing power to the server
 - The tasks to be executed are computationally intensive.
- Data is shipped to clients where processing is performed, and then shipped results back to the server.
- This architecture requires full back-end functionality at the clients.
- Used in many object-oriented database systems
- Issues:
 - Page-Shipping versus Item-Shipping
 - Locking
 - Data Caching
 - Lock Caching

Data Servers cont'd

■ Page-shipping versus item-shipping

- Smaller unit of shipping VS more messages
- Worth **prefetching** related items along with requested item
- Page shipping can be thought of as a form of prefetching

■ Locking

- Overhead of requesting and getting locks from server is high due to message delays.
- Can grant locks on requested and prefetched items; with page shipping, transaction is granted lock on whole page.
- Locks on a prefetched item can be called back by the server, and returned by client transaction if the prefetched item has not been used.
- Locks on the page can be **deescalated** to locks on items in the page when there are lock conflicts. Locks on unused items can then be returned to server.

Data Servers cont'd

■ Data Caching

- Data can be cached at client even in between transactions
- But check that data is up-to-date before it is used (**cache coherency**)
- Check can be done when requesting lock on data item

■ Lock Caching

- Locks can be retained by client system even in between transactions
- Transactions can acquire cached locks locally, without contacting server
- Server **calls back** locks from clients when it receives conflicting lock request.
- Client returns lock once no local transaction is using it.

Accessing Database from Applications

- SQL commands can be called from within a host language (e.g., C++ or Java) program.
 - SQL statements can refer to **host variables** (including special variables used to return status).
 - Must include statement to **connect** to right database.
- Two main integration approaches:
 - Embed SQL in the host language (e.g., Pro*C, Embedded SQL, SQLJ)
 - Create special API (Call Level Interface) to call SQL commands (eg: JDBC, ODBC, PHP)

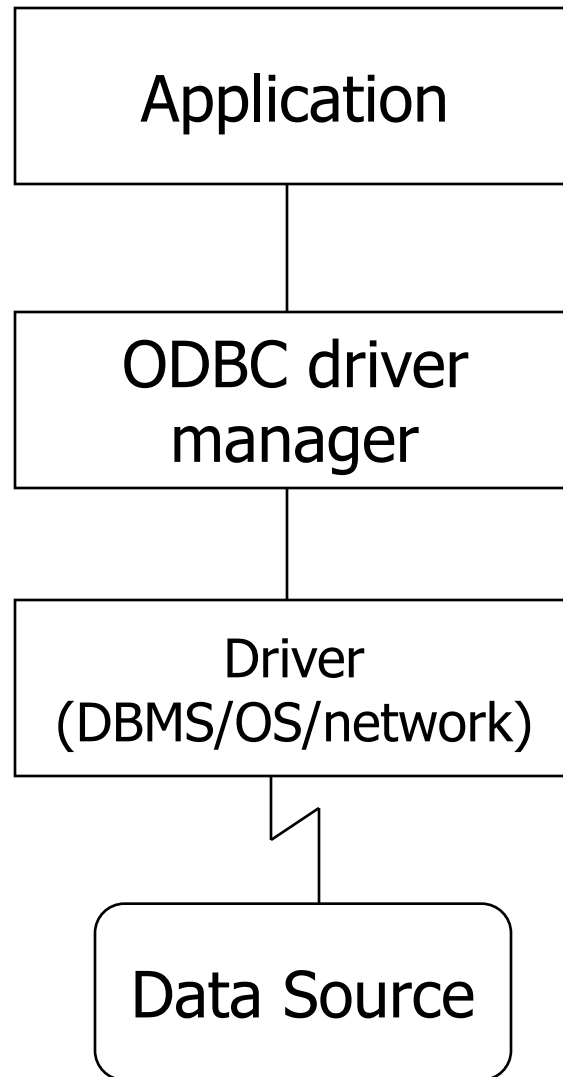
What is ODBC?

- ODBC is short for Open Database Connectivity
 - A standard open application programming interface (API) for accessing a database.
 - SQL Access Group, chiefly Microsoft, in 1992
 - By using ODBC statements in a program, you can access files in a number of different databases, including Access, dBase, DB2, Excel, and Text.
 - It allows programs to use SQL requests that will access databases without having to know the proprietary interfaces to the databases.
 - ODBC handles the SQL request and converts it into a request the individual database system understands.

More on ODBC

- You need:
 - the ODBC software, and
 - a separate module or *driver* for each database to be accessed (library that is dynamically connected to the application).
- Driver masks the heterogeneity of DBMS operating system and network protocol.
 - e.g. (Sybase, Windows/NT, Novell driver)

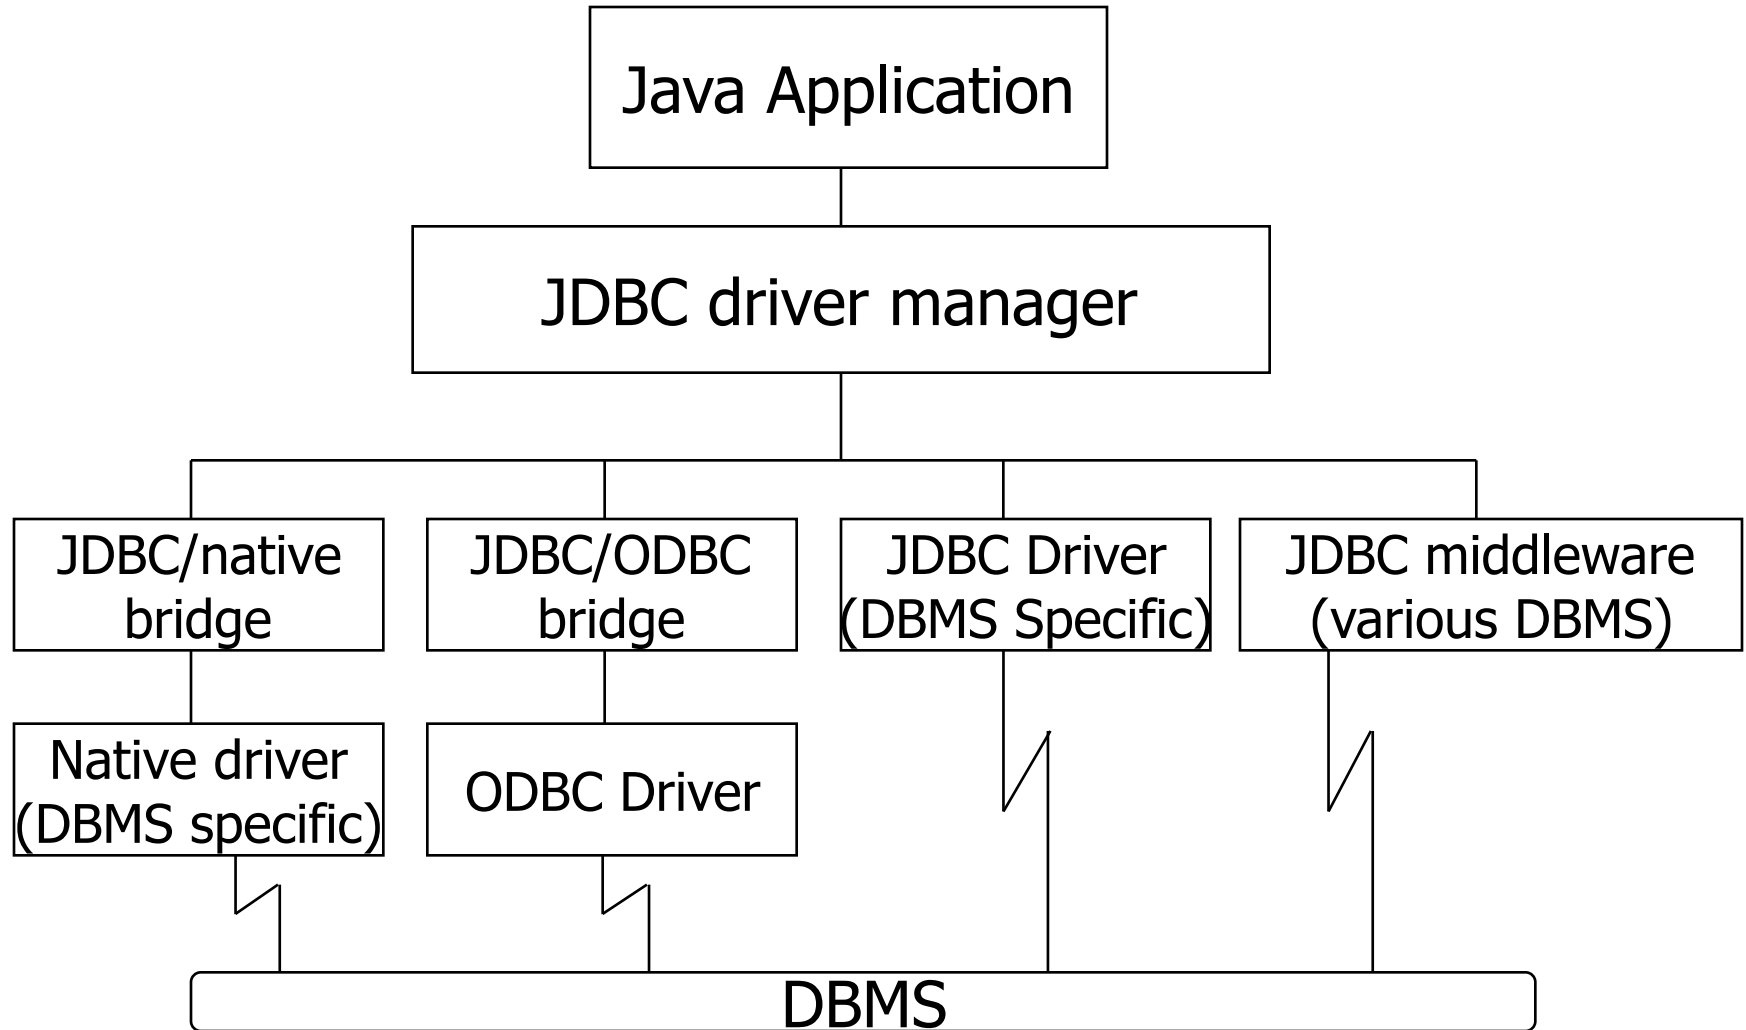
ODBC Architecture



What is JDBC?

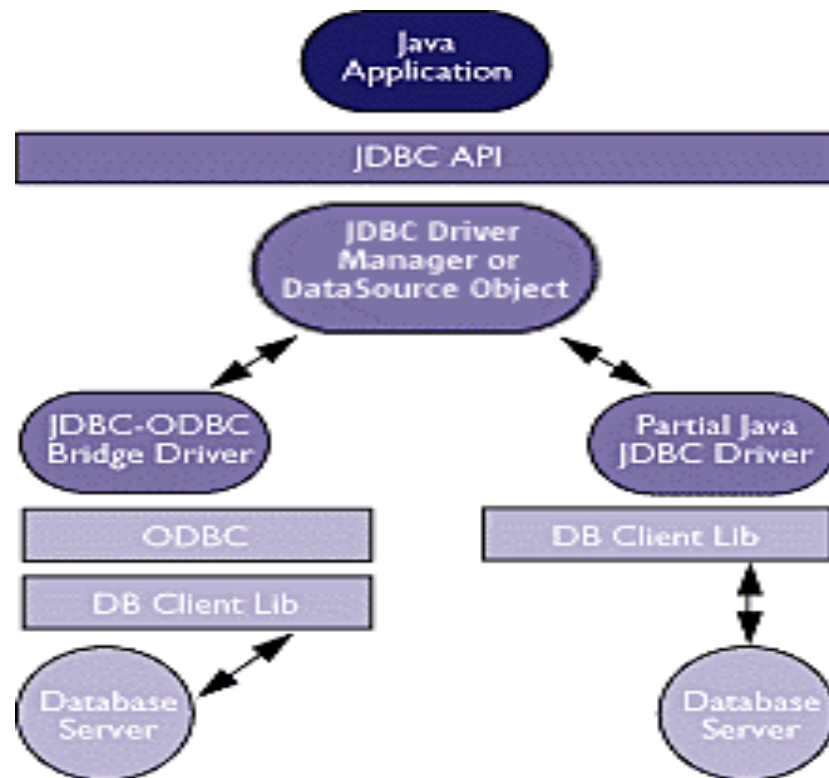
- JDBC is short for Java Database Connectivity
 - is a Java API for connecting programs written in Java to the data in relational databases.
 - consists of a set of classes and interfaces written in the Java programming language.
 - provides a standard API for tool/database developers and makes it possible to write database applications using a pure Java API.
 - The standard defined by Sun Microsystems, allowing individual providers to implement and extend the standard with their own JDBC drivers.
- JDBC:
 - establishes a connection with a database
 - sends SQL statements
 - processes the results

JDBC Architectures



JDBC connectivity using ODBC drivers

- With a small "bridge" program, you can use the JDBC interface to access ODBC-accessible databases.



JDBC vs ODBC

- ODBC is used between applications
- JDBC is used by Java programmers to connect to databases
- With a small "bridge" program, you can use the JDBC interface to access ODBC-accessible databases.
- JDBC allows SQL-based database access for EJB (Enterprise JavaBeans) persistence and for direct manipulation from CORBA and other server objects.



JDBC API

- The JDBC API supports both two-tier and three-tier models for database access.
- Two-tier model - a Java applet or application interacts directly with the database.
- Three-tier model - introduces a middle-level server for execution of business logic:
 - the middle tier to maintain control over data access.
 - user can employ an easy-to-use higher-level API which is translated by the middle tier into the appropriate low-level calls.

The JDBC Steps

- Importing Packages
- Registering the JDBC Drivers
- Opening a Connection to a Database
- Creating a Statement Object
- Executing a Query and Returning a Result Set Object
- Processing the Result Set
- Closing the Result Set, Statement and Connection Objects



1: Importing Packages

```
// Program name: JDBCExample.java
// Purpose: Basic selection using prepared statement
//
//Import packages

import java.sql.*; //JDBC packages
import java.math.*;
import java.io.*;
import oracle.jdbc.driver.*;
```

2: Registering JDBC Drivers

```
class JDBCExample {  
  
    public static void main (String args []) throws SQLException  
    {  
  
        //Load Oracle driver  
        DriverManager.registerDriver (new  
            oracle.jdbc.driver.OracleDriver());  
    }  
}
```

3: Opening connection to a Database

```
//Prompt user for username and password
```

```
String user;
```

```
String password;
```

```
user = readEntry("username: ");
```

```
password = readEntry("password: ");
```

```
//Connect to the local database
```

```
Connection conn = DriverManager.getConnection  
("jdbc:oracle:thin:@aardvark:1526:teach",  
    user,  
    password);
```

4. Creating a Statement Object

```
// Query the hotels table for resort = 'palma nova'
```

```
PreparedStatement pstmt = conn.prepareStatement  
    ("SELECT hotelname, rating FROM hotels WHERE  
    trim(resort) = ?");  
pstmt.setString(1, "palma nova");
```

5. Executing a Query, Returning a Result Set Object & 6. Processing the Result Set

```
ResultSet rset = pstmt.executeQuery();
```

```
//Print query results
```

```
while (rset.next ())
```

```
    System.out.println (rset.getString(1)+" "+  
                        rset.getString(2));
```

7. Closing the Result Set and Statement Objects & 8. Closing the Connection

```
// close the result set, statement, and the connection
    rset.close();
    pstmt.close();
    conn.close();
}
```


Mapping Data Types

- There are data types specified to SQL that need to be mapped to Java data types if the user expects Java to be able to handle them.
- Conversion falls into three categories:
 - SQL type to Java direct equivalents
 - SQL INTEGER direct equivalent of Java int data type
 - SQL type can be converted to a Java equivalent.
 - SQL CHAR, VARCHAR, and LONGVARCHAR can all be converted to the Java String data type
 - SQL data type is unique and requires a special Java data class object to be created specifically for their SQL equivalent
 - SQL DATE converted to the Java Date object that is defined in java.Date especially for this purpose

End of Lecture

■ Summary

- Centralised and Client-Server Systems
- Server System Architectures
- Database Connectivity: ODBC, JDBC, Java Programming with JDBC

■ Reading

- Textbook 6th edition, chapters 5.1, 9.1, 9.2, 17.1, and 17.2
- Textbook 7th edition, chapters 5.1, 9.1, 9.2, 20.2, and 20.3