

Database Development and Design (CPT201)

Lecture 10: Introduction to Object-Oriented Databases

Dr. Wei Wang
Department of Computing

Learning Outcomes

- Complex Data Types and Object Orientation
- Structured Data Types and Inheritance in SQL
- Table Inheritance
- Array and Multiset Types in SQL
- Object Identity and Reference Types in SQL
- Implementing O-R Features
- Persistent Programming Languages
- Comparison of Object-Oriented and Object-Relational Databases



Object-Relational Data Models

- Extend the relational data model by including object orientation and constructs to deal with added data types.
- Allow attributes of tuples to have complex types, including non-atomic values such as nested relations.
- Preserve relational foundations, in particular the declarative access to data, while extending modeling power.
- Upward compatibility with existing relational languages.



Complex Data Types

- Motivation:
 - Permit non-atomic domains (atomic \equiv indivisible)
 - Example of non-atomic domain: set of integers, or set of tuples
 - Allows more intuitive modeling for applications with complex data
- Intuitive definition:
 - allow relations whenever we allow atomic (scalar) values
 - relations within relations
 - Retains mathematical foundation of relational model
 - Violates first normal form (1NF)

Example of a Nested Relation

- Example: library information system
- Each book has
 - title,
 - a list (**array**) of authors,
 - Publisher, with subfields *name* and *branch*, and
 - a set of keywords (set)
- Non-1NF relation *books*

<i>title</i>	<i>author_array</i>	<i>publisher</i>	<i>keyword_set</i>
		(<i>name</i> , <i>branch</i>)	
Compilers	[Smith, Jones]	(McGraw-Hill, NewYork)	{parsing, analysis}
Networks	[Jones, Frick]	(Oxford, London)	{Internet, Web}

4NF Decomposition of Nested Relation

- Suppose for simplicity that title uniquely identifies a book
 - In real world ISBN is a unique identifier
- Decompose *books* into 4NF using the schemas:
 - $(title, author, position)$
 - $(title, keyword)$
 - $(title, pub_name, pub_branch)$
- 4NF design requires users to include joins in their queries.

<i>title</i>	<i>author</i>	<i>position</i>
Compilers	Smith	1
Compilers	Jones	2
Networks	Jones	1
Networks	Frick	2

authors

<i>title</i>	<i>keyword</i>
Compilers	parsing
Compilers	analysis
Networks	Internet
Networks	Web

keywords

<i>title</i>	<i>pub_name</i>	<i>pub_branch</i>
Compilers	McGraw-Hill	New York
Networks	Oxford	London

books4

Complex Types and SQL

- Extensions introduced in SQL:1999 to support complex types:
 - Collection and large object types
 - Nested relations are an example of collection types
 - Structured types
 - Nested record structures like composite attributes
 - Inheritance
 - Object orientation
 - Including object identifiers and references
- Not fully implemented in any database system currently
 - But some features are present in each of the major commercial database systems
 - Read the manual of your database system to see what it supports

Structured Types and Inheritance in SQL

- **Structured types** (a.k.a. **user-defined types**) can be declared and used in SQL

```
create type Name as  
    (firstname    varchar(20),  
    lastname     varchar(20))  
final
```

```
create type Address as  
    (street      varchar(20),  
    city         varchar(20),  
    zipcode     varchar(20))  
not final
```

- Note: **final** and **not final** indicate whether subtypes can be created
- Structured types can be used to create tables with composite attributes

```
create table person (  
    name      Name,  
    address   Address,  
    dateOfBirth date )
```

- Dot notation used to reference components: *name.firstname*

Structured Types

- User-defined row types

```
create type PersonType as (  
    name Name,  
    address Address,  
    dateOfBirth date)  
not final
```

- Can then create a table whose rows are a user-defined type

```
create table person of PersonType
```

- Alternative using unnamed row types

```
create table person_r(  
    name          row(firstname varchar(20),  
                      lastname varchar(20)),  
    address      row(street   varchar(20),  
                      city     varchar(20),  
                      zipcode varchar(20)),  
    dateOfBirth date)
```

Methods

- Can add a method declaration with a structured type.
method *ageOnDate* (*onDate* **date**)
 returns **interval year**
- Method body is given separately.
 create instance method *ageOnDate* (*onDate* **date**)
 returns **interval year**
 for *CustomerType*
 begin
 return *onDate* - **self.dateOfBirth**;
 end
- We can now find the age of each customer:
 select *name.lastname, ageOnDate* (**current_date**)
 from *customer*

Constructor Functions

- **Constructor functions** are used to create values of structured types

```
create function Name(firstname varchar(20), lastname varchar(20))  
returns Name  
begin  
    set self.firstname = firstname;  
    set self.lastname = lastname;  
end
```

- To create a value of type *Name*, we use
new *Name*('John', 'Smith')
- Normally used in insert statements
insert into *Person* values
 (new *Name*('John', 'Smith'),
 new *Address*('20 Main St', 'New York', '11001'),
 date '1960-8-22');

Type Inheritance

- Suppose that we have the following type definition for Person:
create type *Person*
 (*name* **varchar**(20),
 address **varchar**(20))
- Using inheritance to define the student and teacher types
create type *Student*
 under *Person*
 (*degree* **varchar**(20),
 department **varchar**(20))
create type *Teacher*
 under *Person*
 (*salary* **integer**,
 department **varchar**(20))
- Subtypes can redefine methods by using **overriding method** in place of **method** in the method declaration

Multiple Type Inheritance

- SQL:1999 and SQL:2003 do not support multiple inheritance
- **If** our type system supports multiple inheritance, we can define a type for teaching assistant as follows:

```
create type Teaching_Assistant  
under Student, Teacher
```

- To avoid a conflict between the two occurrences of *department* we can rename them

```
create type Teaching_Assistant  
under  
    Student with (department as student_dept),  
    Teacher with (department as teacher_dept)
```

- Each value must have a **most-specific type**

Table Inheritance

- Tables created from subtypes can further be specified as **subtables**
- E.g. `create table people of Person;`
 `create table students of Student under people;`
 `create table teachers of Teacher under people;`
- Tuples added to a subtable are automatically visible to queries on the supertable
 - E.g. query on *people* also sees *students* and *teachers*.
 - Similarly, updates/deletes on *people* also result in updates/deletes on subtables
 - To override this behaviour, use “**only people**” in query
- Conceptually, multiple inheritance is possible with tables
 - e.g. *teaching_assistants* under *students* and *teachers*
 - *But is not supported in SQL currently*
 - So we cannot create a person (tuple in *people*) who is both a student and a teacher

Consistency Requirements for Subtables

- Consistency requirements on subtables and supertables.
 - Each tuple of the supertable (e.g. *people*) can correspond to at most one tuple in each of the subtables (e.g. *students* and *teachers*)
 - Additional constraint in SQL:1999: All tuples corresponding to each other (that is, with the same values for inherited attributes) must be derived from one tuple (inserted into one table).
 - Each entity must have a most specific type
 - We cannot have a tuple in *people* corresponding to a tuple each in *students* and *teachers*

Array and Multiset Types in SQL

- Example of array and multiset declaration:

```
create type Publisher as  
  (name          varchar(20),  
   branch       varchar(20));  
create type Book as  
  (title         varchar(20),  
   author_array varchar(20) array [10],  
   pub_date      date,  
   publisher     Publisher,  
   keyword-set   varchar(20) multiset);  
create table books of Book;
```


Creation of Collection Values

- Array construction

```
array ['Silberschatz', `Korth', `Sudarshan']
```

- Multisets

```
multiset ['computer', 'database', 'SQL']
```

- To create a tuple of the type defined by the books relation:

```
('Compilers', array[ `Smith', `Jones'],  
new Publisher( `McGraw-Hill', `New York'),  
multiset [ `parsing', `analysis' ])
```

- To insert the preceding tuple into the relation books

```
insert into books  
values
```

```
('Compilers', array[ `Smith', `Jones'],  
new Publisher( `McGraw-Hill', `New York'),  
multiset [ `parsing', `analysis' ]);
```

Querying Collection-Valued Attributes

- To find all books that have the word "database" as a keyword,

```
select title
from books
where 'database' in (unnest(keyword-set))
```
- We can access individual elements of an array by using indices
 - E.g.: If we knew that a particular book has three authors, we could write:

```
select author_array[1], author_array[2], author_array[3]
from books
where title = 'Database System Concepts'
```
- To get a relation containing pairs of the form "title, author_name" for each book and each author of the book

```
select B.title, A.author
from books as B, unnest (B.author_array) as A (author)
```
- To retain ordering information we add a **with ordinality** clause

```
select B.title, A.author, A.position
from books as B, unnest (B.author_array) with ordinality as
A (author, position)
```

Unnesting

- The transformation of a nested relation into a form with fewer (or no) relation-valued attributes is called **unnesting**.
- E.g.

```
select B.title, A.author, B.publisher.name as pub_name,  
       B.publisher.branch as pub_branch, K.keyword  
from books as B, unnest(B.author_array) as A(author),  
       unnest(B.keyword_set) as K(keyword)
```
- Result relation: *flat_books*

<i>title</i>	<i>author</i>	<i>pub_name</i>	<i>pub_branch</i>	<i>keyword</i>
Compilers	Smith	McGraw-Hill	New York	parsing
Compilers	Jones	McGraw-Hill	New York	parsing
Compilers	Smith	McGraw-Hill	New York	analysis
Compilers	Jones	McGraw-Hill	New York	analysis
Networks	Jones	Oxford	London	Internet
Networks	Frick	Oxford	London	Internet
Networks	Jones	Oxford	London	Web
Networks	Frick	Oxford	London	Web

Nesting

- **Nesting** is the opposite of unnesting, creating a collection-valued attribute
- Nesting can be done in a manner similar to aggregation, but using the function **collect()** in place of an aggregation operation, to create a multiset or array.
- To nest the *flat_books* relation on the attribute *keyword*:

```
select title, author, Publisher(pub_name, pub_branch) as publisher,  
       collect(keyword) as keyword_set  
from flat_books  
group by title, author, publisher
```
- To nest on both authors and keywords:

```
select title, collect(author) as author_set,  
       Publisher(pub_name, pub_branch) as publisher,  
       collect(keyword) as keyword_set  
from flat_books  
group by title, publisher
```

Nesting cont'd

- Another approach to creating nested relations is to use subqueries in the **select** clause, starting from the 4NF relation *books4* (see *Textbook*)

```
select B.title,  
       array (select author  
              from authors as A  
              where A.title = B.title  
              order by A.position) as  
          author_array,  
       B.Publisher (pub-name, pub-branch) as publisher,  
       multiset (select keyword  
                 from keywords as K  
                 where K.title = B.title) as keyword_set  
from books4 as B
```

Object-Identity and Reference Types

- Define a type *Department* with a field *name* and a field *head* which is a reference to the type *Person*, with table *people* as scope:

```
create type Department (  
    name varchar (20),  
    head ref (Person) scope people)
```

- We can then create a table *departments* as follows

```
create table departments of Department
```

- We can omit the declaration **scope** *people* from the type declaration and instead make an addition to the **create table** statement:

```
create table departments of Department  
    (head with options scope people)
```

- Referenced table must have an attribute that stores the identifier, called the **self-referential attribute**

```
create table people of Person  
ref is person_id system generated;
```



Initialising Reference-Typed Values

- To create a tuple with a reference value, we can first create the tuple with a null reference and then set the reference separately:

```
insert into departments  
values ('CS', null)
```

```
update departments  
set head = (select p.person_id  
            from people as p  
            where name = 'John')  
where name = 'CS'
```

User Generated Identifiers

- The type of the object-identifier must be specified as part of the type definition of the referenced table, and
- The table definition must specify that the reference is user generated

```
create type Person
(name varchar(20)
 address varchar(20))
ref using varchar(20)
create table people of Person
ref is person_id user generated
```

- When creating a tuple, we must provide a unique value for the identifier:

```
insert into people (person_id, name, address) values
('01284567', 'John', '23 Coyote Run')
```

- We can then use the identifier value when inserting a tuple into *departments*
 - Avoids need for a separate query to retrieve the identifier:

```
insert into departments
values('CS', '02184567')
```


User Generated Identifiers

- Can use an existing primary key value as the identifier:

```
create type Person
  (name varchar (20) primary key,
   address varchar(20))
  ref from (name)
create table people of Person
  ref is person_id derived
```

- When inserting a tuple for *departments*, we can then use

```
insert into departments
  values(`CS`,`John`)
```

Path Expressions

- Find the names and addresses of the heads of all departments:

```
select head->name, head->address  
from departments
```
- An expression such as "head->name" is called a **path expression**
- Path expressions help avoid explicit joins
 - If department head were not a reference, a join of *departments* with *people* would be required to get at the address
 - Makes expressing the query much easier for the user

Implementing O-R Features

- Similar to how E-R features are mapped onto relation schemas
 - Object-relational database systems are basically extensions of existing relational database systems.
 - To minimise changes to the storage-system code (relation storage, indices, etc.), the complex data types can be translated to the simpler type system of relational databases.
- Subtable implementation
 - Each table stores primary key and those attributes defined in that table, or
 - Each table stores both locally defined and inherited attributes

Why Persistent Programming Language?

- Persistent programming language is a programming language extended with constructs to handle **persistent** data.
- Why need persistent programming language ?
 - Access to a database is only one component of any real-world application.
 - Data-manipulation language like SQL is effective for accessing data.
 - But programming language is required for implementing other components such as user interfaces or communication with other computers.
- The traditional way of interfacing database languages to programming languages is by embedding SQL within the programming language.

Difference between two Languages

- Persistent programming languages can be distinguished from languages with embedded SQL in at least two ways:
 - With an embedded language, the **type system** of the host language usually differs from the type system of the data-manipulation language. The programmer is responsible for any type conversions between the host language and SQL. The programmer using an embedded query language is responsible for **writing explicit code** to fetch/store data from/to databases.
 - In contrast, in a persistent programming language, the programmer can manipulate persistent data without writing code explicitly.

Persistent Programming Languages

- Languages extended with constructs to handle persistent data
- Programmer can manipulate persistent data directly
 - no need to fetch it into memory and store it back to disk (unlike embedded SQL)
- Approaches to make persistent objects:
 - **Persistence by class** - explicit declaration of persistence
 - **Persistence by creation** - special syntax to create persistent objects
 - **Persistence by marking** - make objects persistent after creation
 - **Persistence by reachability** - object is persistent if it is declared explicitly to be so or is reachable from a persistent object

Object Identity and Pointers

- Degrees of permanence of object identity
 - **Intraprocedure**: only during execution of a single procedure
 - **Intraprogram**: only during execution of a single program or query
 - **Interprogram**: across program executions, but not if data-storage format on disk changes
 - **Persistent**: interprogram, plus persistent across data reorganisations
- Persistent versions of C++ and Java have been implemented
 - C++
 - ODMG C++
 - ObjectStore
 - Java
 - Java Database Objects (JDO)

Object-Relational Mapping

- **Object-Relational Mapping (ORM)** systems built on top of traditional relational databases
- Implementer provides a mapping from objects to relations
 - Objects are purely transient, no permanent object identity
- Objects can be retrieved from database
 - System uses mapping to fetch relevant data from relations and construct objects
 - Updated objects are stored back in database by generating corresponding update/insert/delete statements
- The **Hibernate** ORM system is widely used
 - An implementation of the Java Persistence API.
 - Provides API to start/end transactions, fetch objects, etc
 - Provides query language operating directly on object model
 - queries translated to SQL
- Limitations: overheads, especially for bulk updates

Comparison of Databases

- **Relational systems**
 - simple data types, powerful query languages, high protection.
- **Persistent-programming-language-based OODBs**
 - complex data types, integration with programming language, high performance.
- **Object-relational systems**
 - complex data types, powerful query languages, high protection.
- **Object-relational mapping systems**
 - complex data types integrated with programming language, but built as a layer on top of a relational database system
- **Note: Many real systems blur these boundaries, e.g.**
 - persistent programming language built as a wrapper on a relational database offers first two benefits, but may have poor performance.

End of Lecture

■ Summary

- Complex Data Types and Object Orientation
- Structured Data Types and Inheritance in SQL, Table Inheritance
- Array and Multiset Types in SQL
- Object Identity and Reference Types in SQL
- Implementing O-R Features
- Persistent Programming Languages
- Comparison of Object-Oriented and Object-Relational Databases

■ Reading

- Textbook 6th edition, chapter 22
- Textbook 7th edition, shortened. PDF will be provided.