

Xi'an Jiaotong-Liverpool University

PAPER CODE	EXAMINER	DEPARTMENT	TEL
CSE201		Computer Science & Software Engineering	

FIRST SEMESTER 2017/2018 FINAL EXAMINATIONS

BACHELOR DEGREE – Year 3

DATABASE DEVELOPMENT AND DESIGN

TIME ALLOWED: 2 Hours

INSTRUCTIONS TO CANDIDATES

- 1、 Total marks available are 100. This will count for 80% in the final assessment.
- 2、 Answer four questions.
- 3、 The number in the column on the right indicates the marks for each section.
- 4、 Answer should be written in the answer booklet(s) provided.
- 5、 The university approved calculator - Casio FS82ES/83ES can be used.
- 6、 All the answers must be in English.

THIS PAPER MUST NOT BE REMOVED FROM THE EXAMINATION ROOM

Xi'an Jiaotong-Liverpool University

Question 2. Consider the following three relations:

- i) *author* (*authorID*, *name*, *affiliation*, *email*)
- ii) *publishes* (*articleID*, *authorID*)
- iii) *article* (*articleID*, *title*, *journal*, *year*, *publisher*)

“*authorID*” and “*articleID*” are the primary keys for relations *author* and *article*, respectively. *authorID* is the join attribute for *author* and *publishes*. *articleID* is the join attribute for *publishes* and *article*. Relation *author* has 5,000 tuples stored on 500 blocks. Relation *publishes* has 100,000 tuples stored on 100 blocks. Answer the following questions.

[25 marks]

- a) Describe how the selection $\delta_{year < 2014}$ on relation *article* can be evaluated in the most cost effective way by using a primary index on year and a secondary index on year, respectively.
[4/25]
- b) Assume that the memory can only hold one block for each relation. Using the blocked nested loop join algorithm and *publishes* as the outer relation, how many block transfers and seeks are needed to compute “*author* \bowtie *publishes*”, respectively?
[4/25]
- c) Assume that a B+ tree index with the number of pointers in a node, $N=10$, is built for *author* on the attribute *authorID*. How many block transfers and seeks are needed to compute “*author* \bowtie *publishes*” using the indexed nested loop join, respectively?
[6/25]
- d) If both relations have been sorted and the merge join algorithm is used to evaluate “*author* \bowtie *publishes*”. Assume that the buffer for input and output, $b_b=1$; also, it takes $t_T=0.1ms$ for a block transfer and $t_S=4ms$ for a seek. Compared to the results from Question 2.b) and 2.c), which algorithm is the most cost effective in terms of total time for block transfers and seeks? Justify your answer.
[7/25]
- e) Briefly outline the procedure of performing a hash join.
[4/25]

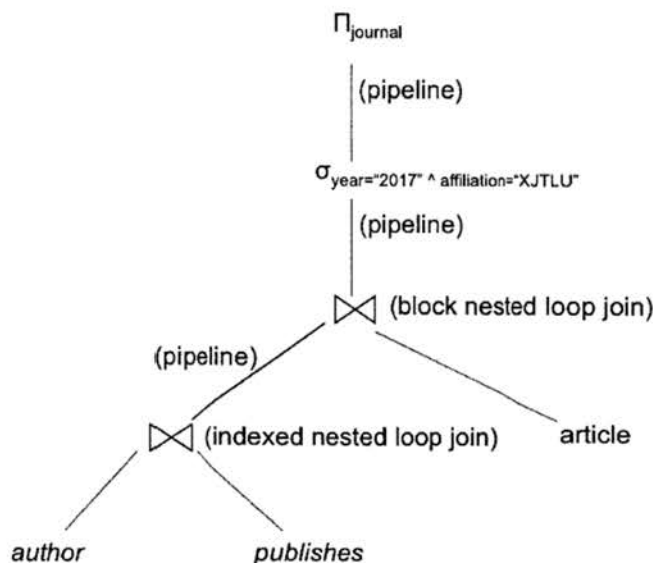
Xi'an Jiaotong-Liverpool University

Question 3. Consider the three relations in Question 2 and their catalog information. “*authorID*” and “*articleID*” are the primary keys for relations *author* and *article*, respectively. The two attributes in *publishes* are the foreign keys referencing *author* and *article*, respectively.

- i) *author* (*authorID*, *name*, *affiliation*, *email*)
- ii) *publishes* (*articleID*, *authorID*)
- iii) *article* (*articleID*, *title*, *journal*, *year*, *publisher*)

- number of records in *author*, $n_{author} = 5,000$;
- number of blocks in *author*, $b_{author} = 500$;
- number of records in *publishes*, $n_{publishes} = 100,000$;
- number of blocks in *publishes*, $b_{publishes} = 100$;
- number of records in *article*, $n_{article} = 20,000$;
- number of blocks in *article*, $b_{article} = 2,000$;
- index: a primary B⁺-tree index of height 3 on the *authorID* attribute of *publishes* relation;
- number of distinct values for the attribute *affiliation* in the *author* relation, $V(author, affiliation) = 500$;
- the number of distinct values for the attribute *year* in the *article* relation, $V(article, year) = 10$;

A query evaluation plan is shown below.



Xi'an Jiaotong-Liverpool University

Question 5. Answer the following questions.

[25 Marks]

- a) The two-phase commit protocol is one of the most widely used protocols in distributed database systems to ensure transaction atomicity. Describe how the protocol works in distributed database systems.

[4/25]

- b) Assume that the transaction coordinator fails during execution of the two-phase commit protocol in distributed database systems, how can the status of the transaction be determined based on the logs?

[4/25]

- c) Consider the three relations in Question 2. Assume that they are stored in a distributed database at two sites:

XJTLU site:

author

publishes

IEEE site:

article

To answer the query “*find the names of XJTLU staff and titles of articles that were published by them*”, one needs to compute the join of the three relations. It is known that only a small number of articles stored at the IEEE site were published by authors stored at the XJTLU site. Assume that a query is made to the XJTLU site; describe how semi-join can be used to optimise the query (e.g., reduce cost).

[6/25]

- d) Can the centralised deadlock detection approach be used to detect deadlock in distributed database system? If yes, describe how it works. If no, explain why.

[5/25]

- e) Consider the following transaction table. Columns represent the items and rows represent transactions. Each cell represents whether an item is purchased in a transaction, “1” means yes and “0” means no. Compute the following:

(1) support (A)

(2) support (A->>C)

(3) confidence (A->>C)

Xi'an Jiaotong-Liverpool University

ID\Item	A	B	C	D	E	F
T1	1	1	1	0	0	0
T2	1	0	1	0	0	0
T3	1	0	0	1	0	0
T4	0	1	0	0	1	1

[6/25]

END OF EXAM PAPER