


Question 1. Answer the following questions.

[25 marks]

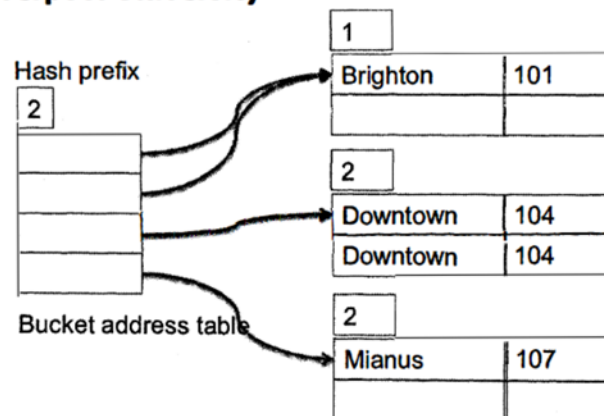
A relation called *employee* holds 30,000 tuples. Each tuple is stored as fixed length and fixed format record. Each tuple has the length of 250 bytes and the key attribute *ID* has the length of 20 bytes. The tuples are stored sequentially in a number of blocks, ordered by the *ID* attribute. Each block has the size of 4,096 bytes (i.e., 4 Kilobytes). Consider creating a primary index on the *ID* attribute, and a secondary index on the *city* attribute (stores the city where an employee is from). Each index entry consists of a search key and a 10-byte long pointer to data. Assume that a record or an index entry can only be stored in one block.

- Compute the number of disk blocks needed to store the relation *employee*. **[4/25]**
- If the primary index on *ID* is sparse (i.e., one index entry for one block), compute the number of blocks needed to store the index. **[4/25]**
- Is a sparse secondary index on the *city* attribute useful? Why? **[3/25]**
- Suppose that the number of pointers in a B+ tree node is 20 and there are one million distinct search key values to be indexed. What is the maximum height of the B+ tree? **[2/25]**
- The initial B+ tree for indexing *names* in the *employee* is shown below. Draw the B+ trees after the following operations: (1) insert Amy; (2) insert Linda; (3) delete Salman. Each subsequent operation should be performed based on the result of the previous operation. **[9/25]**



- An extendable index (see diagram below) is created on the *city* attribute. Describe in detail how the index will be updated if another two records with the search key value of "Mianus" are inserted.

Xi'an Jiaotong-Liverpool University



[3/25]

Question 2. Consider the following two relations and their catalogue information:

- i) *account*(*account_Number*, *customer_Name*, *balance*, *branch_ID*)
- ii) *branch*(*branch_ID*, *branch_Name*, *branch_City*, *postcode*)

The “*account_Number*” is the key for the *account* relation, and the “*branch_Name*” is the key for the *branch* relation. The *account* relation contains 300,000 records stored in 60,000 blocks, and the *branch* relation contains 500 records stored in 50 blocks. Assume that both relations are sequentially stored by the key attributes. Answer the following questions.

[25 marks]

- a) Suppose that the linear search algorithm is used to evaluate the selection $\delta_{balance > 5,000}$ in the *account* relation, how many block transfers are needed? How many seeks are needed?
[4/25]
- b) Suppose that a sparse B+ tree index (with the number of pointers in a node, $N=7$) has been created for the *account* relation on the attribute “*account_Number*”, how many block transfers are needed for the selection $\delta_{account_Number = 11737}$? How many seeks are needed? (Hint: in a sparse index, every index entry contains a pointer to one block)
[4/25]
- c) Suppose that none of the relations can fit in memory, and the nested loop join algorithm is used to evaluate “*account* \bowtie *branch*”. Which relation should be used as outer relation? How many block transfers are needed? How many seeks are needed?
[5/25]
- d) Suppose that the external sort merge algorithm is used to sort the *account* relation on the *account_Number* attribute. Assume that the memory size $M=30$ and the buffer for reading and writing $b_b=2$. How many block transfers are needed? How many seeks are needed?
[4/25]
- e) Suppose that the hash join algorithm is used to evaluate “*account* \bowtie *branch*”, the number of partitions, $n_h=70$, and the size of the buffer for reading and writing, $b_b=2$. How many block transfers are needed? How many seeks are needed?
[4/25]
- f) With regard to the results obtained from Questions 2.c) and 2.e), discuss which algorithm is more efficient in evaluating the join.
[4/25]


Question 3. Consider the following two relations and their catalog information. Answer the questions below.

wechatGroup(group_ID, group_Name, size, yearOfCreation, owner)

member(member_ID, member_Name, city, group_ID)

- *group_ID* and *member_ID* are the keys for the two relations, respectively.
- number of tuples in relation *wechatGroup*, $n_r = 1,000$
- number of blocks in *wechatGroup*, $b_r = 200$
- number of distinct values on attribute *yearOfCreation* in *wechatGroup*, $V(\text{yearOfCreation}, \text{wechatGroup}) = 100$
- number of tuples in relation *member*, $n_s = 10,000$
- number of blocks in relation *member*, $b_s = 1,000$
- number of distinct values on attribute *city* in *member*, $V(\text{city}, \text{member}) = 100$

A query evaluation plan is shown below.



[25 marks]

a) What is the relational algebra expression for the given evaluation plan?

[3/25]

b) One of the heuristic rules for query optimisation is to perform selection operations as early as possible. Write the equivalent algebra expression for the answer from Question 3.a).

[4/25]

Xi'an Jiaotong-Liverpool University

- c) What are the differences between materialisation and pipelining in evaluating an expression?

[4/25]

- d) Suppose that all selections are evaluated by using linear scan and pipelining is used for projection. Draw an annotated evaluation tree for the relational algebra expression obtained from Question 3.b).

[4/25]

- e) Based on the catalog information, what is the estimated size of the selection

$\sigma_{yearOfCreation="2017"}(wechatGroup)?$

[4/25]

- f) Assume that for each tuple in $\sigma_{yearOfCreation="2017"}(wechatGroup)$, the average cost of performing the hash join is 3 block transfers. What is the total number of block transfers for the whole evaluation plan in Question 3.d)?

[6/25]

Question 4. Answer the following questions.

[25 marks]

- a) Draw a precedence diagram for the following schedule. Is it conflict serialisable?

T1	T2	T3	T4
read(Y)			
read(X)			
write(X)			
		write(Y)	
	read(X)		
	write(Y)		
			read(W)
			write(W)
			read(Y)
	read(W)		

[4/25]

- b) Is the following schedule recoverable? Justify your answer.

Schedule: T1:write(X); T1:write(Y); T2:read(X); T2:write(Y); T2:read(Z); T1:write(Z); T1:commit; T3:read(Y); T2: commit; T3:write(Z); T4:read(Z); T3:commit; T4:abort.

[4/25]

- c) Consider the following schedule. Assume that the database failure happens at time=18, answer the following questions: (1) at the checkpoint, what transactions are in the list *L*? (2) Which transactions need to be redone? (3) Which need to be undone? Justify your answer.

Time	T7	T8	T9
0			start
1			read(B)
2	-----Checkpoint {L}-----		
3			B=B+10
4	start		
5	read(A)		
6	A=A+1		
7			write(B)
8			commit
9		start	
10		read(A)	
11		read(B)	
12		B=A+2B	

Xi'an Jiaotong-Liverpool University

13		write(B)	
14		commit	
15	read(B)		
16	$B=B+3A$		
17	write(A)		
18	-----System Failure-----		

[5/25]

- d) Describe how the centralised deadlock detection method is used to detect deadlocks in a distributed database system.

[6/25]

- e) In the context of distributed database, briefly describe how the “primary site” and “peer-to-peer” methods work for asynchronous data replication, respectively.

[6/25]**END OF EXAM PAPER**