Database Development and Design (CSE210)

Tutorial 3

Dr. Wei Wang

Department of Computing

Q1

branch_name	h(*branch_name*)
Brighton	0010 1101 1111 1011 0010 1100 0011 0000
Downtown	1010 0011 1010 0000 1100 0110 1001 1111
Mianus	1100 0111 1110 1101 1011 1111 0011 1010
Perryridge	1111 0001 0010 0100 1001 0011 0110 1101
Redwood	0011 0101 1010 0110 1100 1001 1110 1011
Round Hill	1101 1000 0011 1111 1001 1100 0000 0001

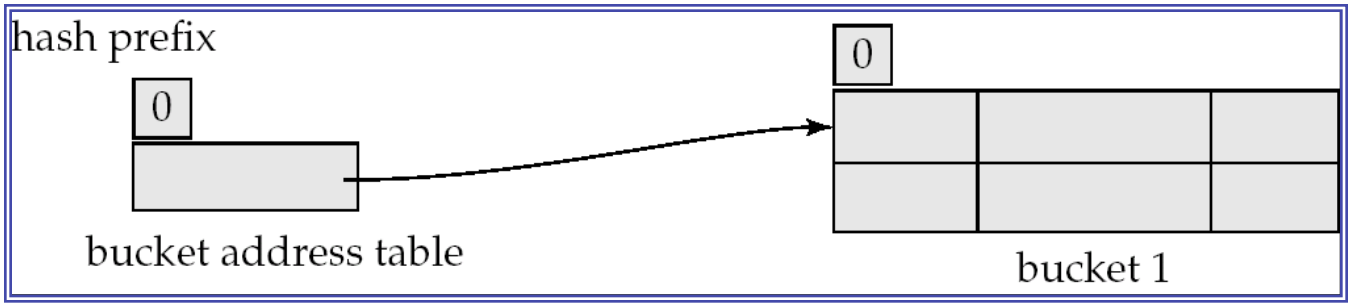

- Given the above hash values for the branch_name and the initial hash index, draw the hash indices after inserting the following 8 values. Assume that one block can hold two tuples.
  - Brighton, Downtown, Downtown, Mianus, Perryridge, Perryridge, Perryridge, Redwood

Q1 Solution

- Insert "Brighton"

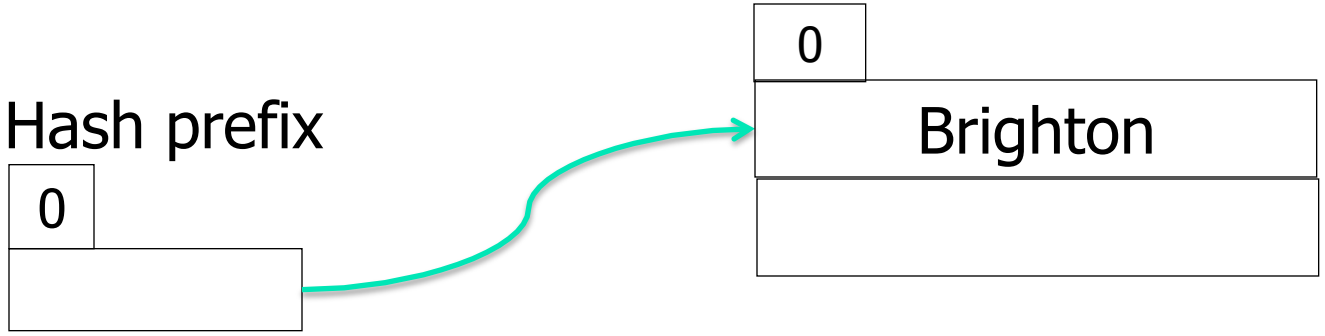

Bucket address table

Q1 Solution cont'd

- Insert "Downtown"

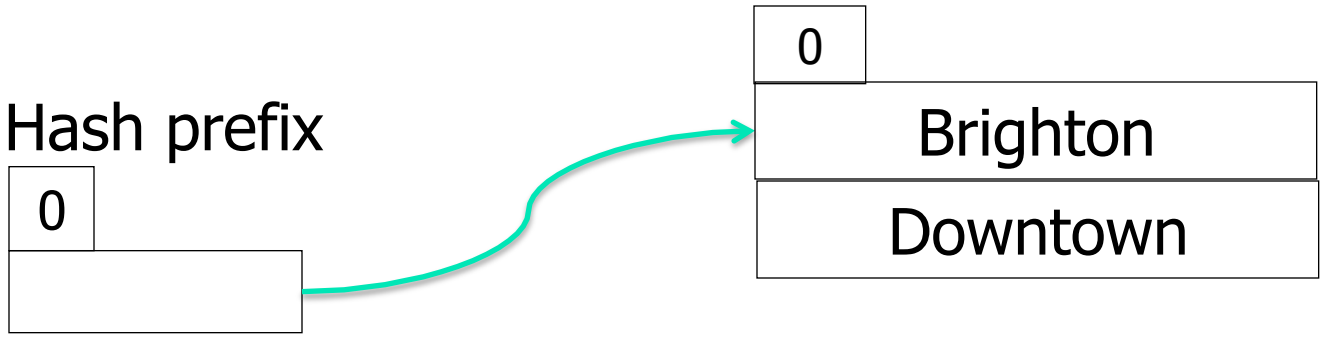

Bucket address table

Q1 Solution cont'd

- Insert "Downtown"

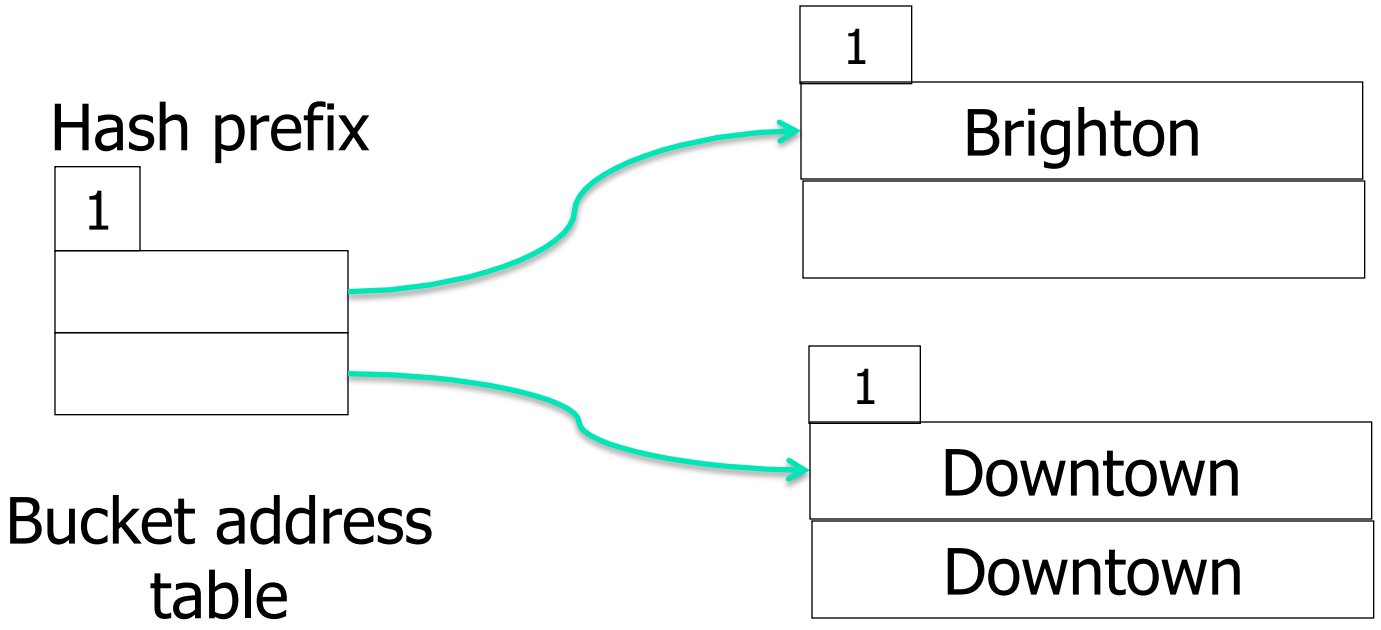

Q1 Solution cont'd

- Insert "Mianus"

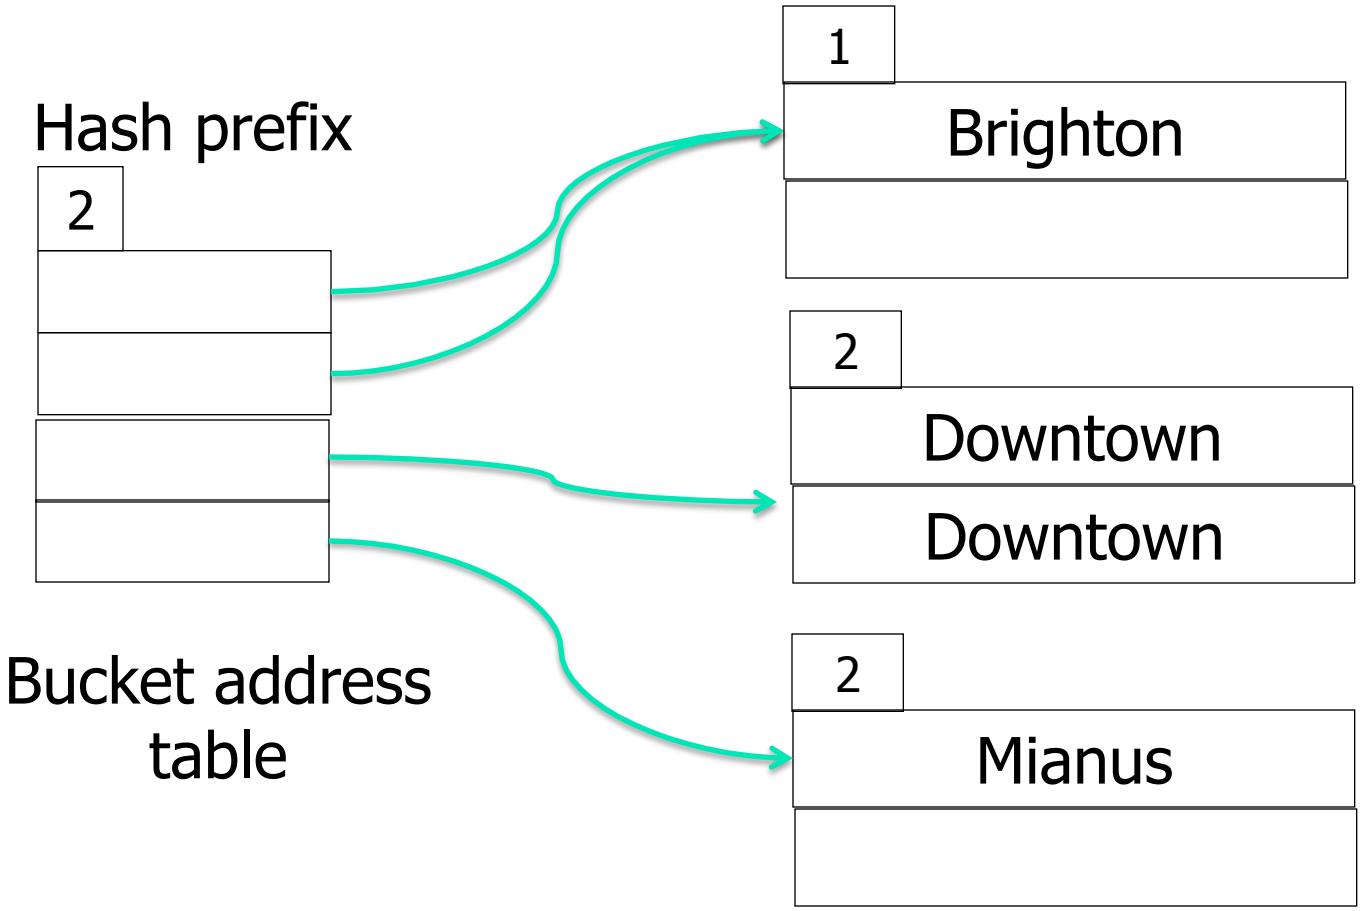

Q1 Solution cont'd

- Insert "Perryridge"

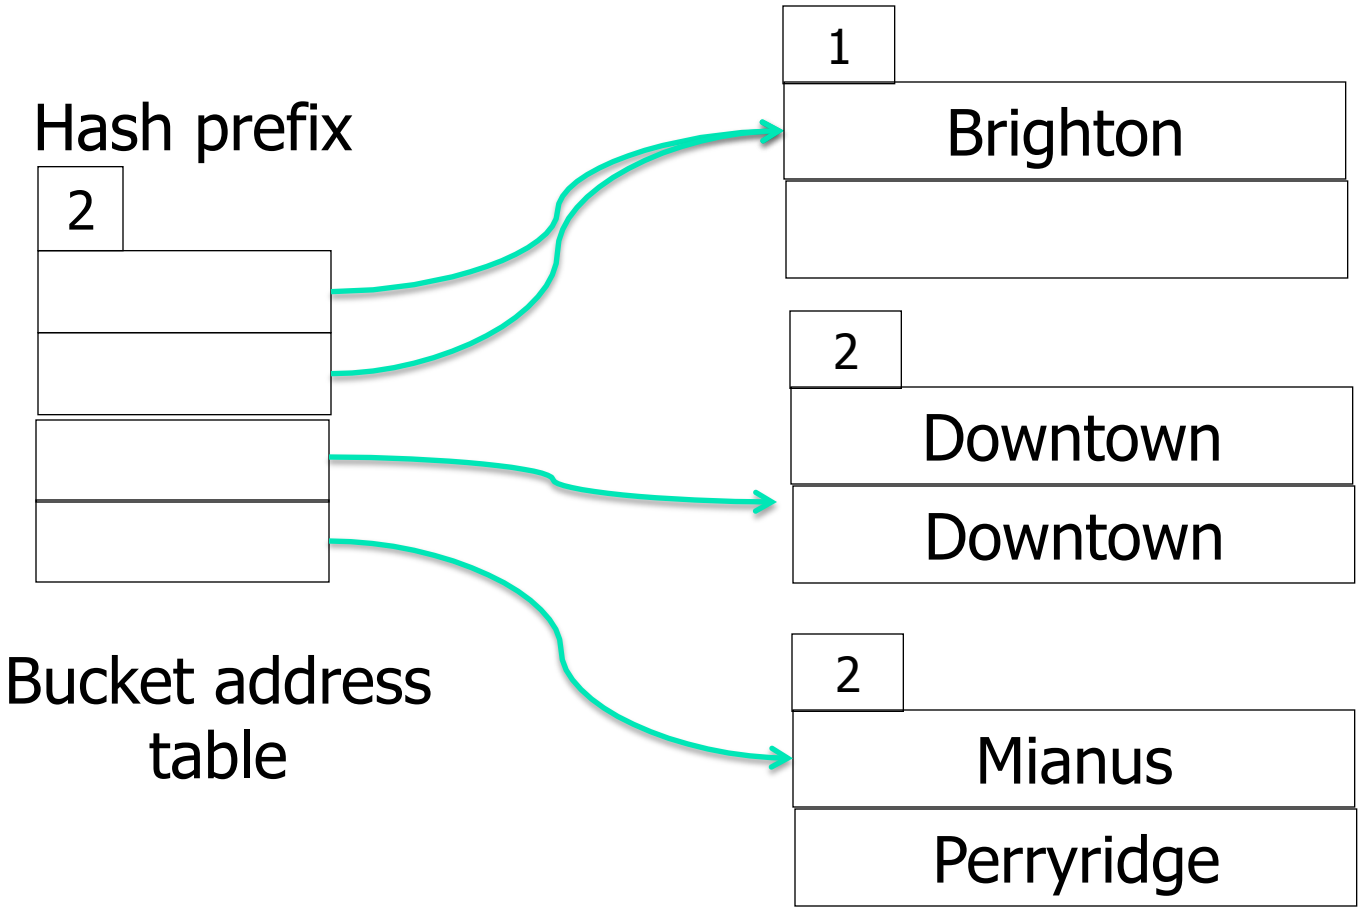

Q1 Solution cont'd

- Insert "Perryridge"

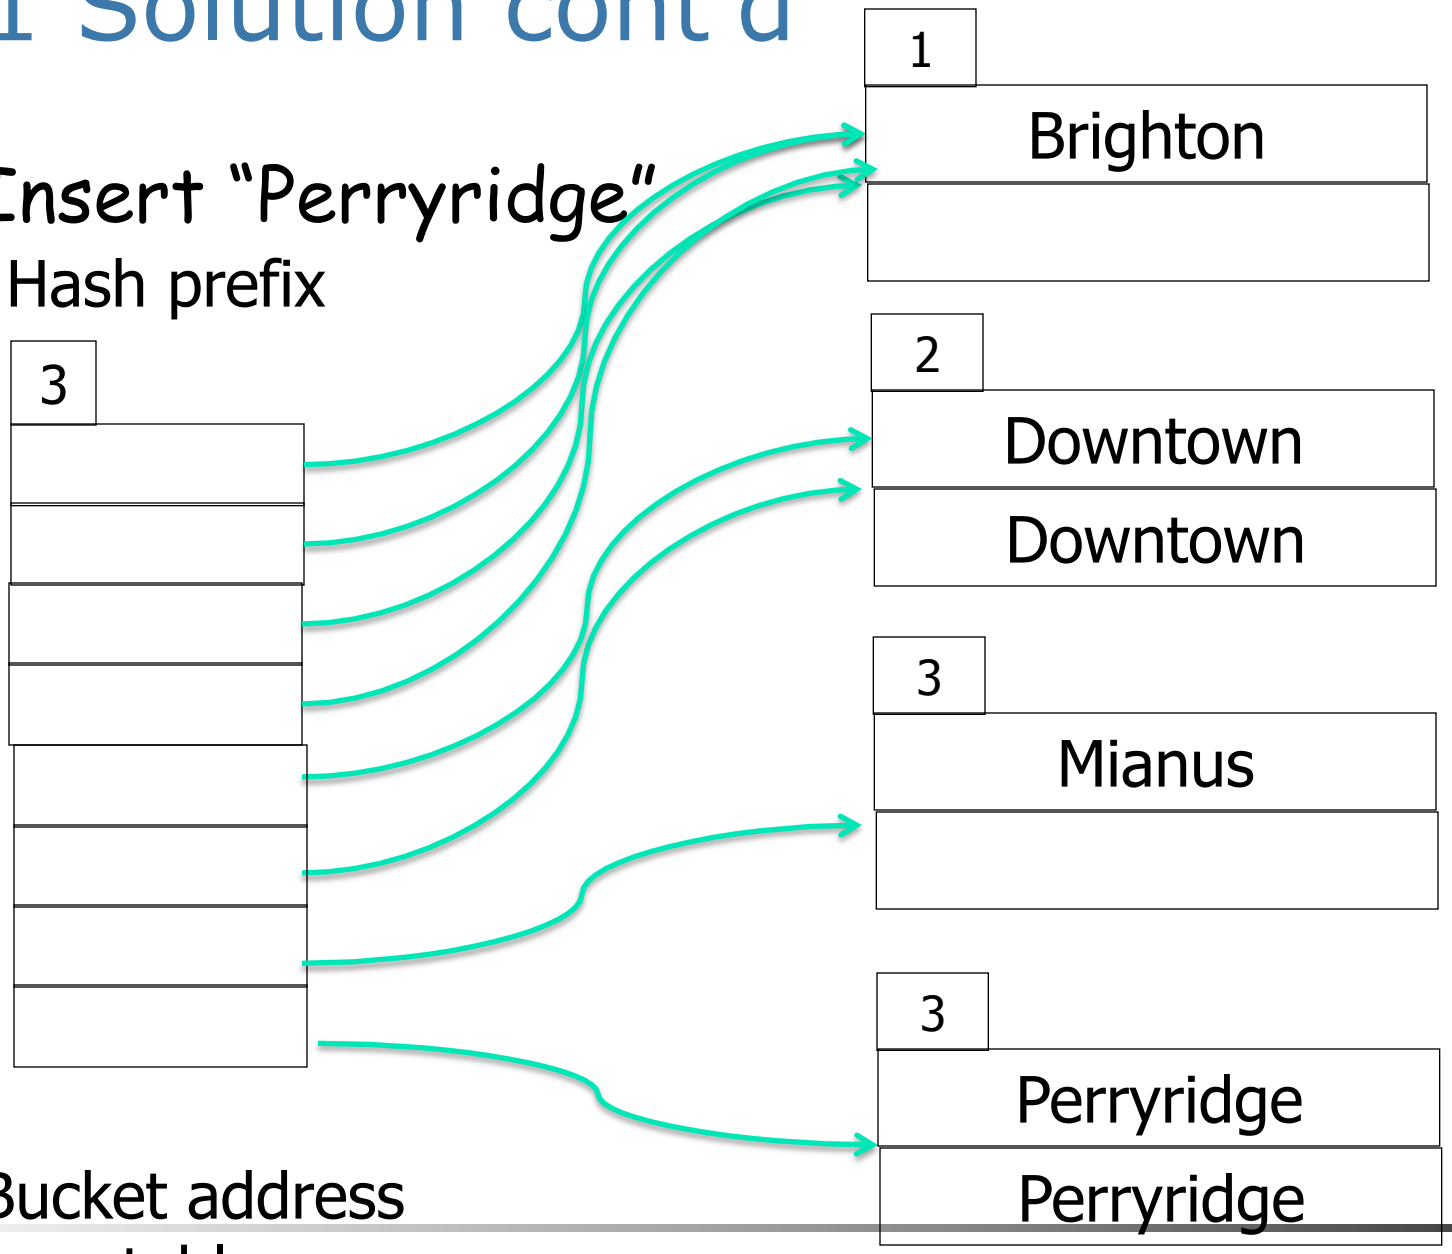

Q1 Solution cont'd

- Insert "Perryridge"

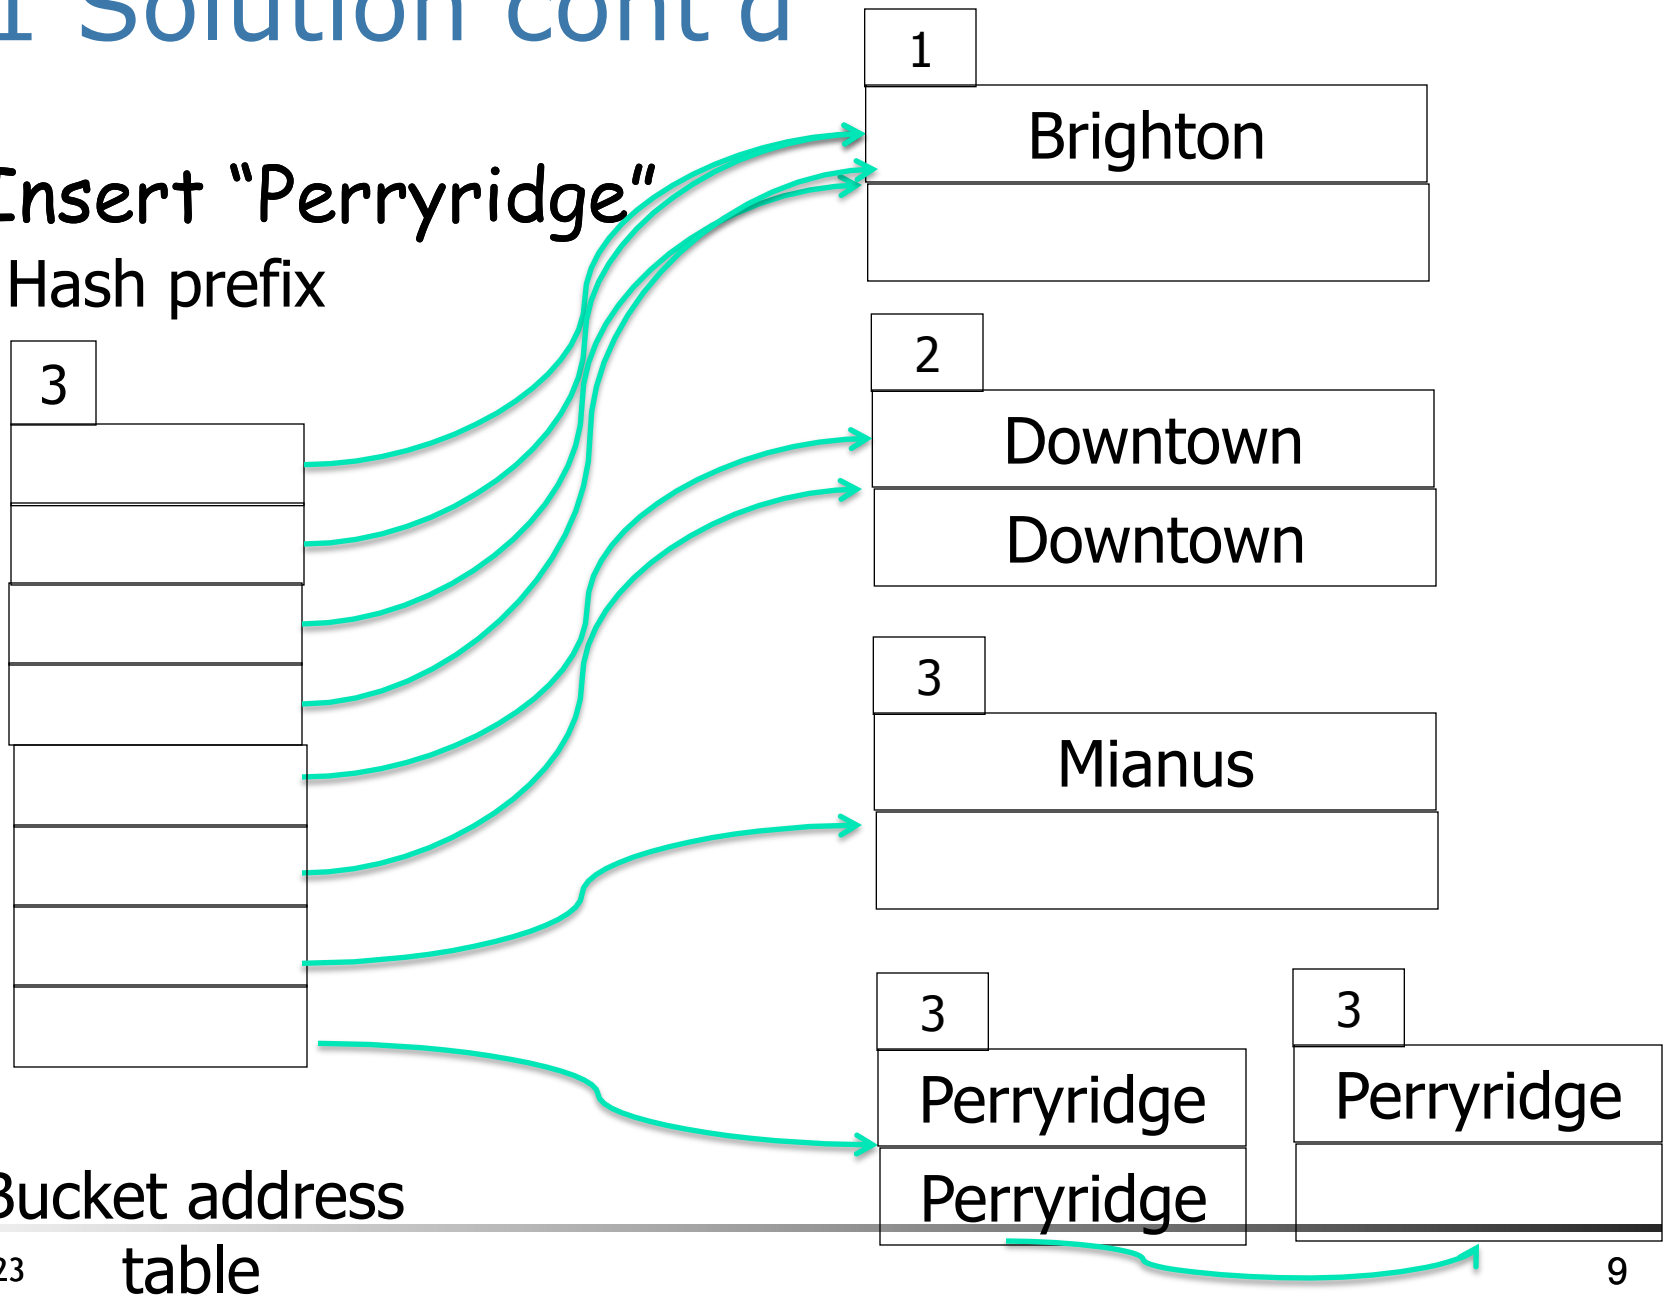

Q1 Solution cont'd

- Insert "Redwood"

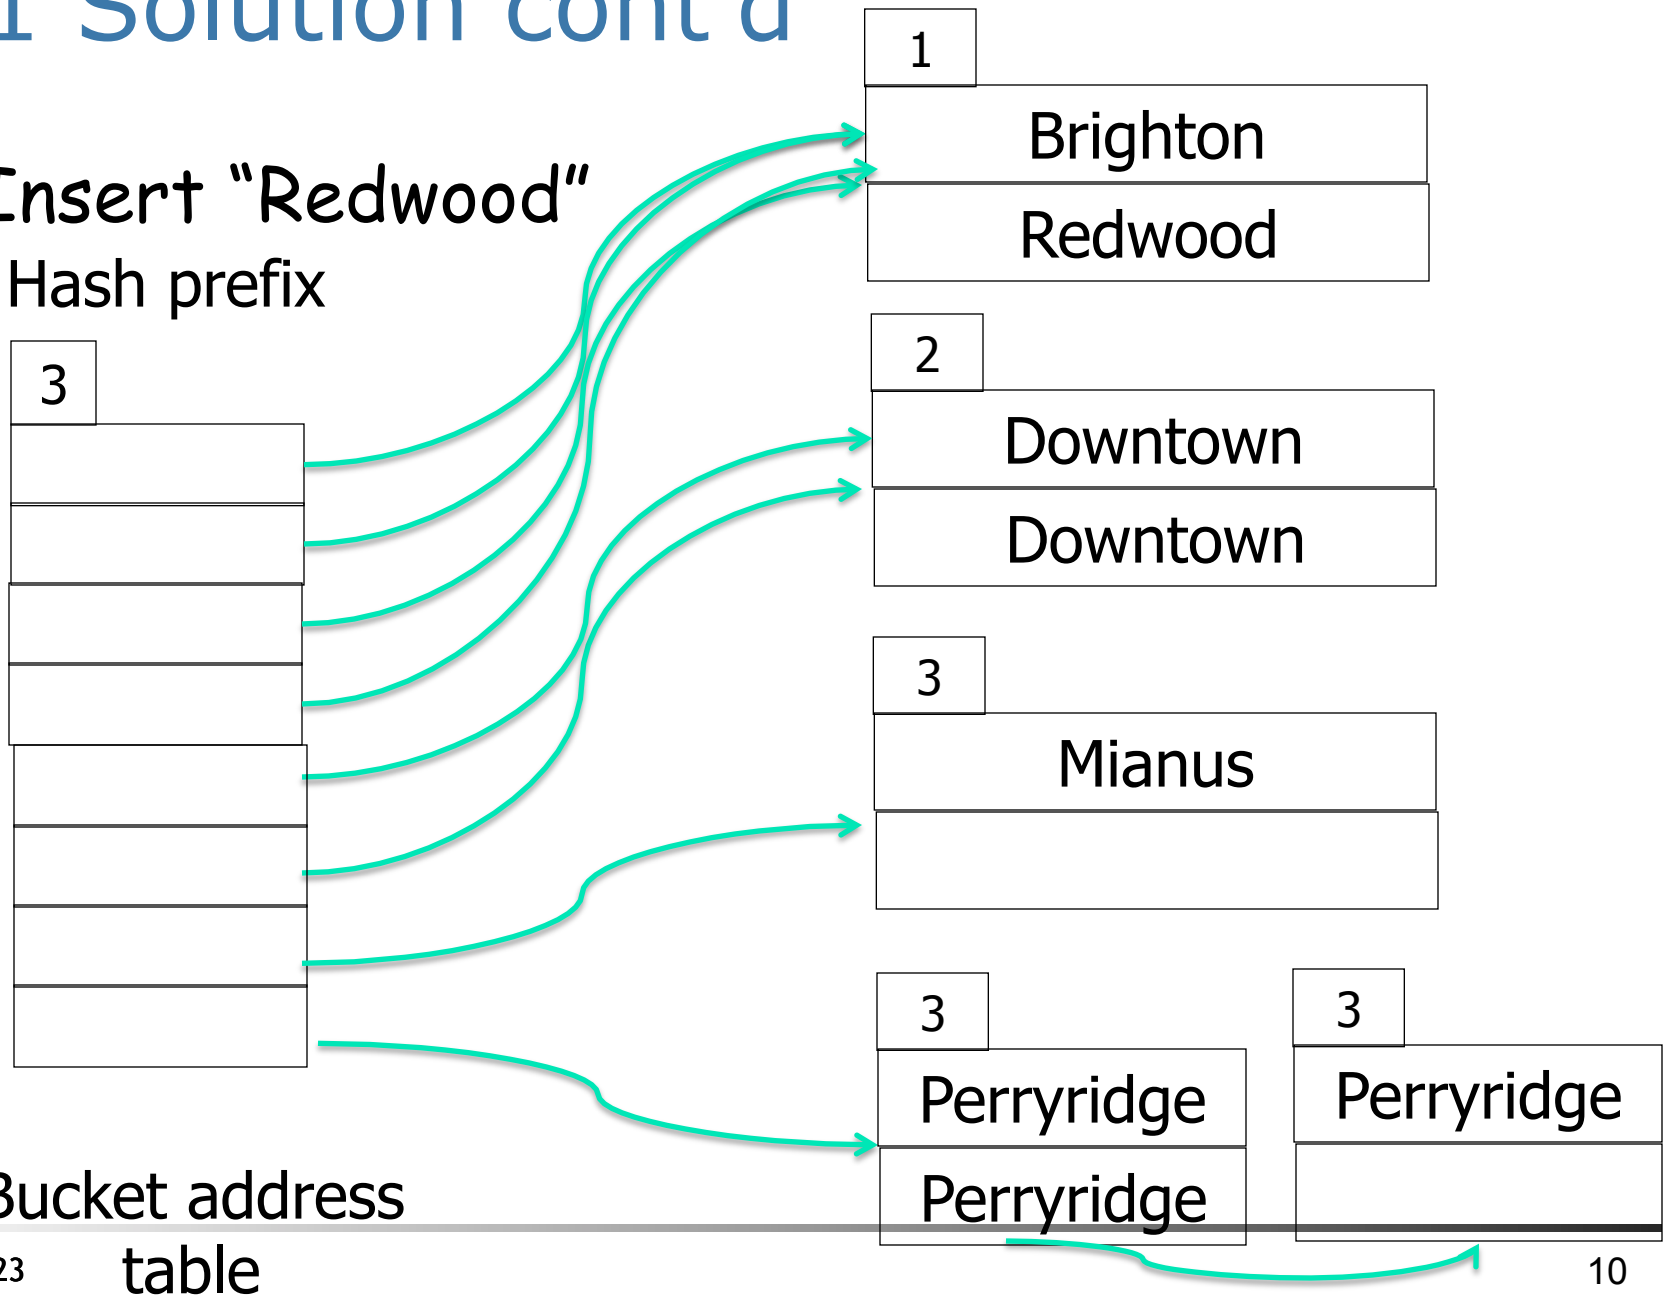

Q2

branch_name	h(*branch_name*)
Brighton	0010 1101 1111 1011 0010 1100 0011 0000
Downtown	1010 0011 1010 0000 1100 0110 1001 1111
Mianus	1100 0111 1110 1101 1011 1111 0011 1010
Perryridge	1111 0001 0010 0100 1001 0011 0110 1101
Redwood	0011 0101 1010 0110 1100 1001 1110 1011
Round Hill	1101 1000 0011 1111 1001 1100 0000 0001

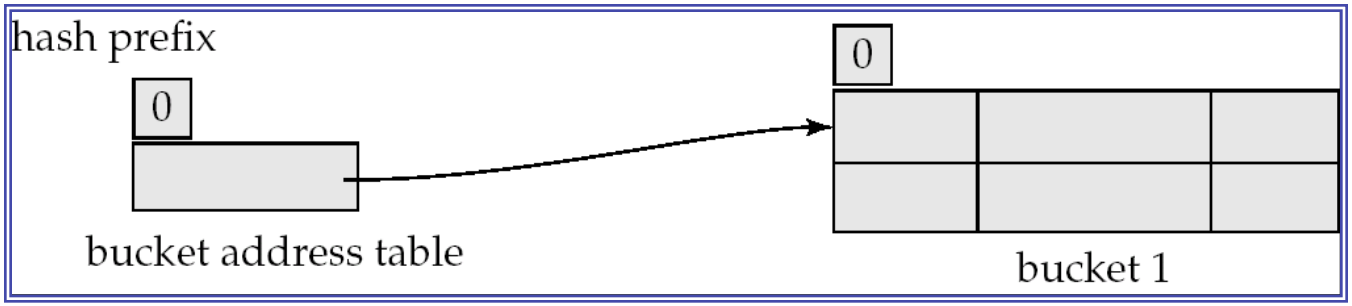

- Given the above hash values for the branch_name and the initial hash index, draw the hash indices after inserting the following 8 values. Assume that one block can hold two tuples.
  - Redwood, Perryridge, Perryridge, Perryridge, Mianus, Downtown, Downtown, Brighton

Q2 Solution

- Insert "Redwood"

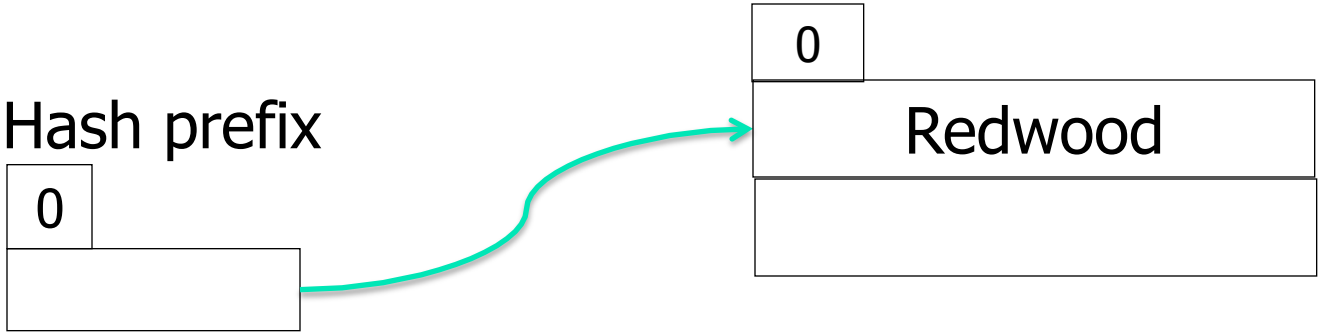

Bucket address table

Q2 Solution cont'd

- Insert "Perridge"

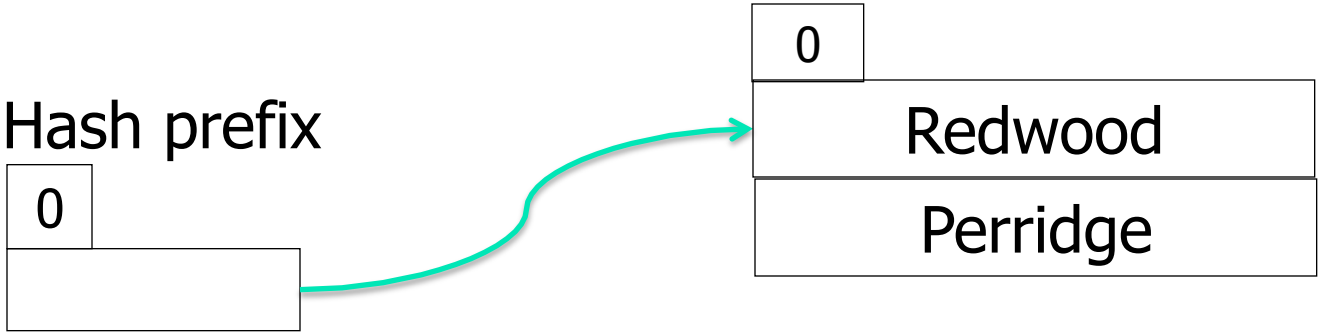

Bucket address table

Q2 Solution cont'd

- Insert "Perridge"

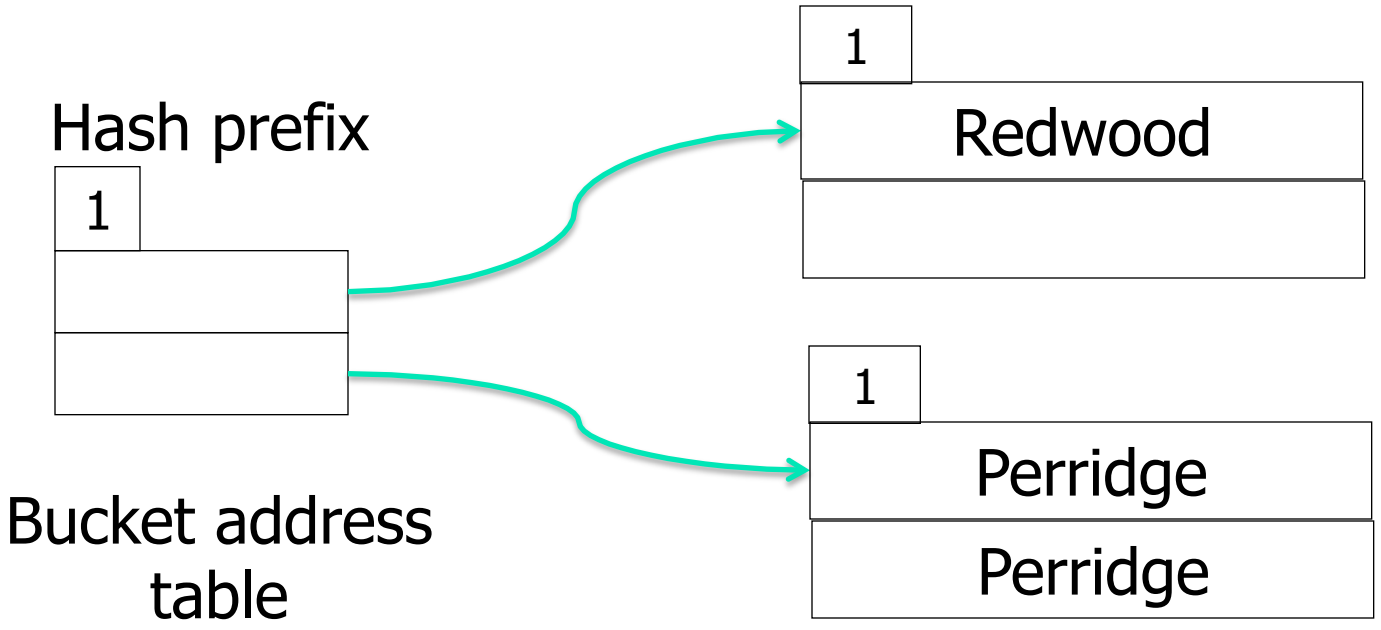

Q2 Solution cont'd

- Insert "Perridge"

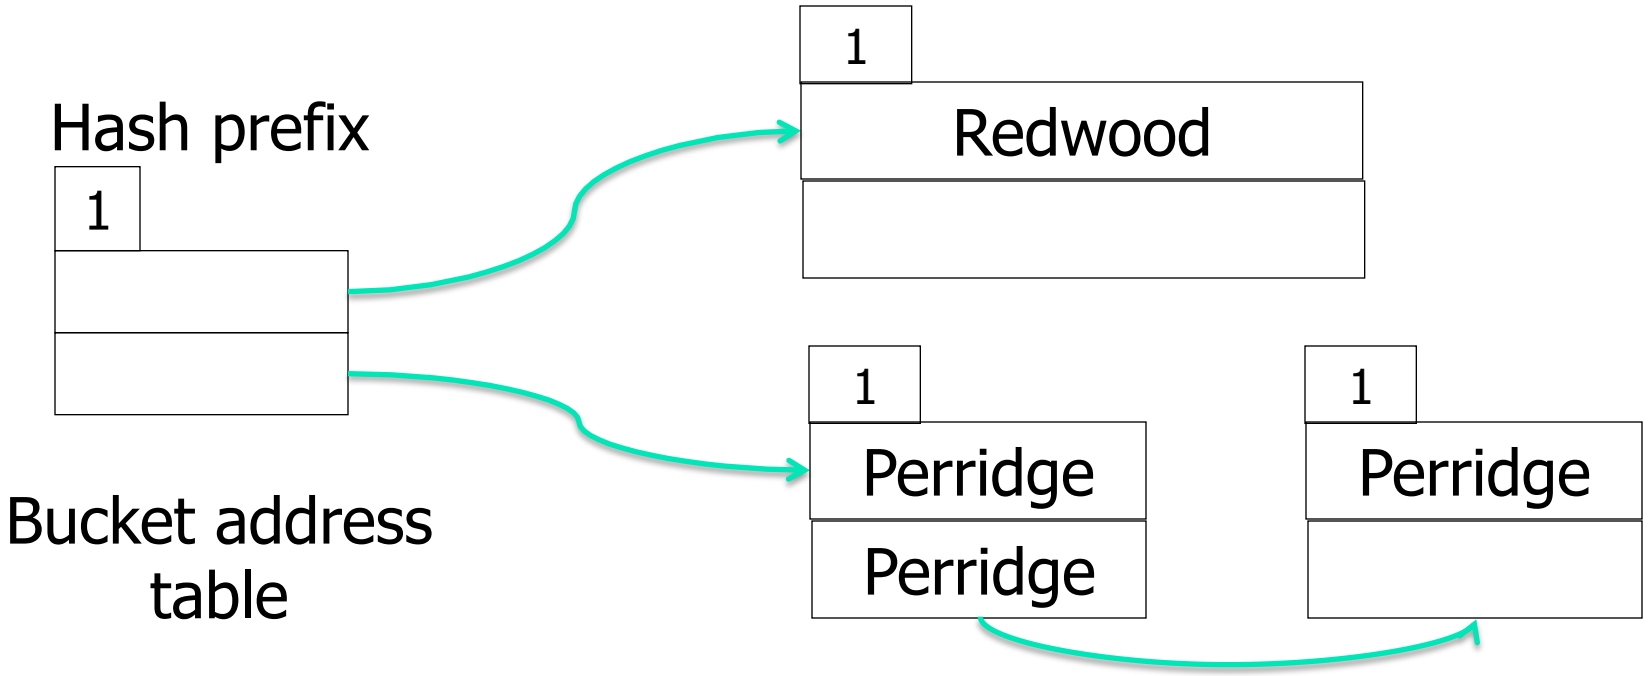

Q2 Solution cont'd

- Insert "Mianus"

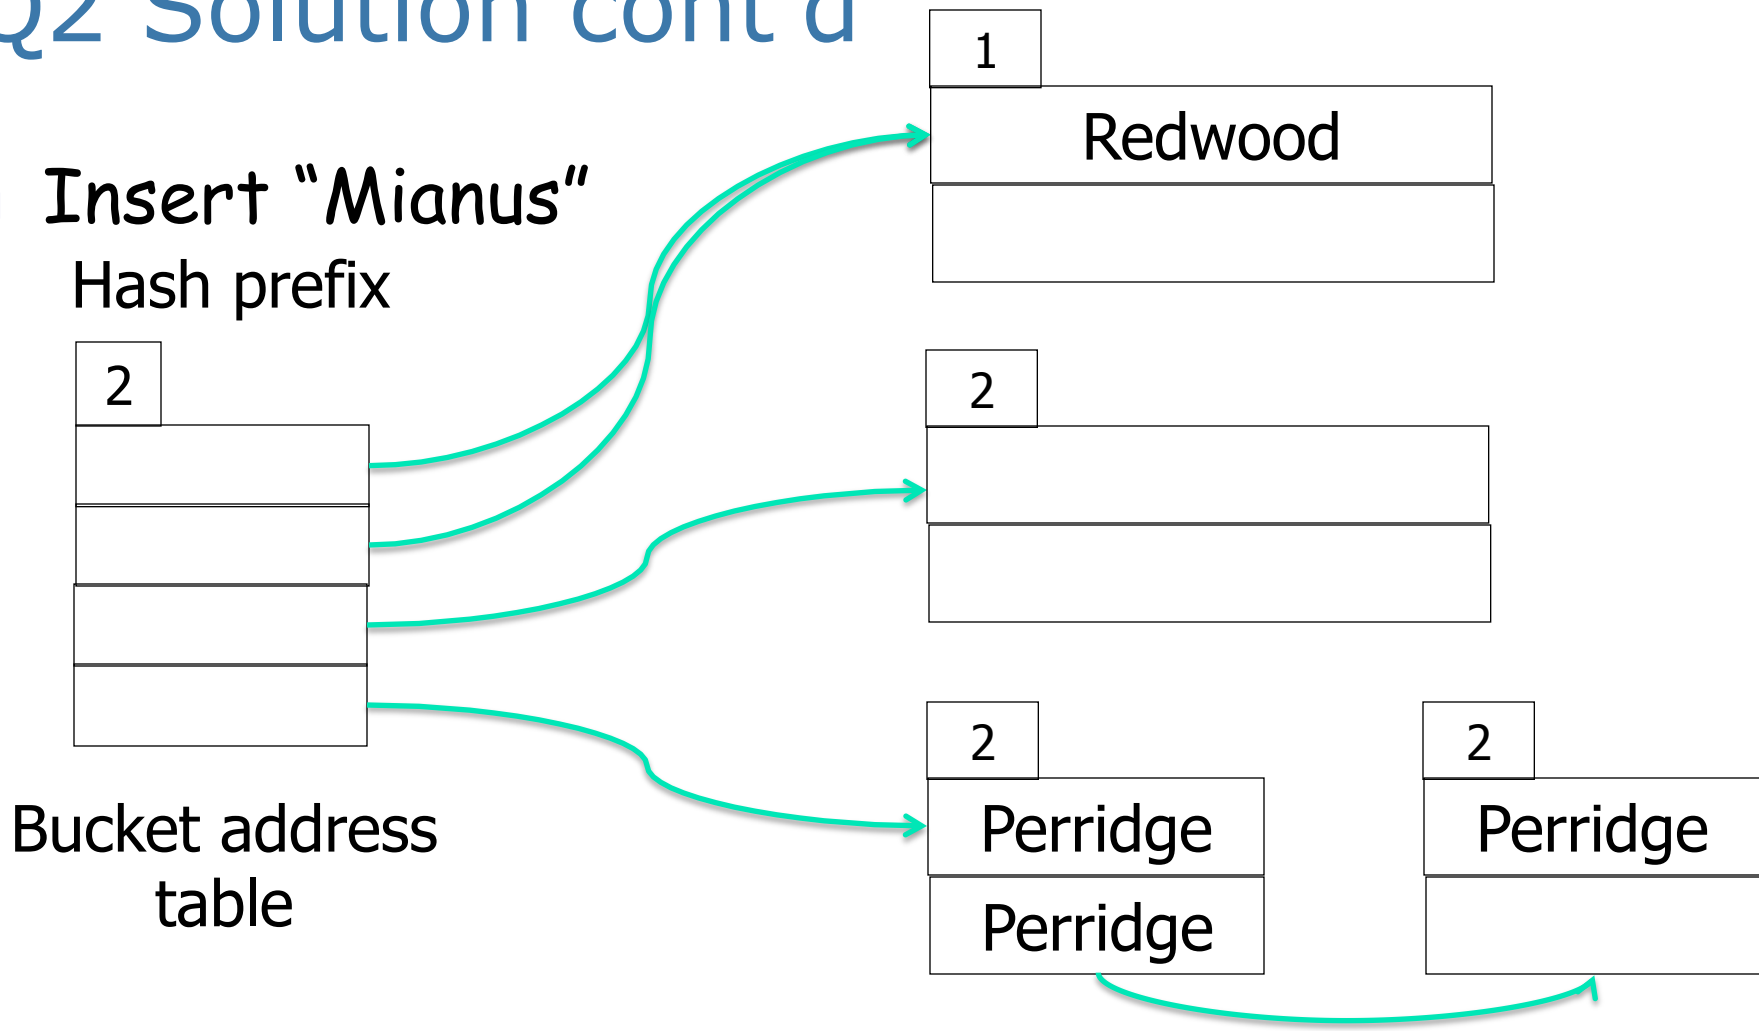

NOT COMPLETE. STILL CANNOT INSERT MIANUS.
SEE NEXT SLIDE.

Q2 Solution cont'd

- Insert "Mianus"

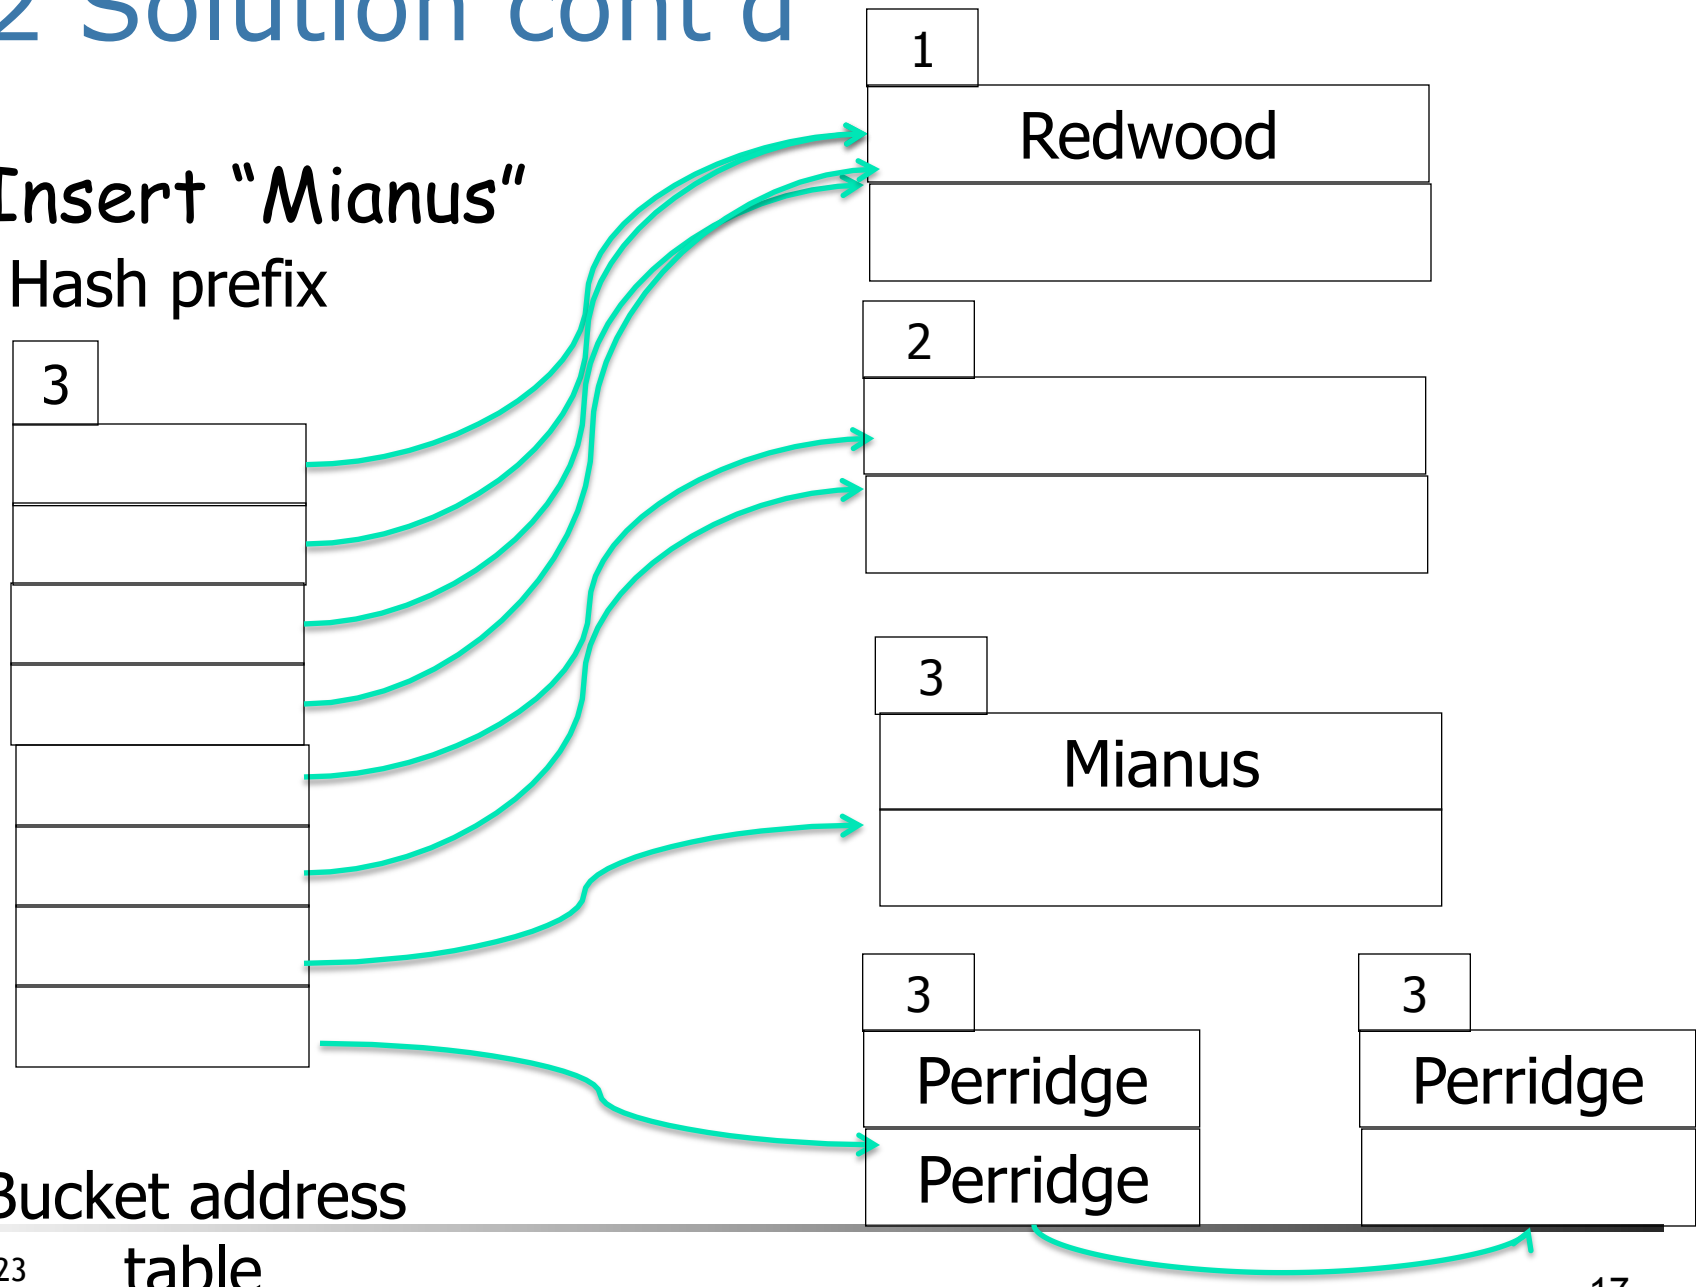

Q2 Solution cont'd

- Insert "Downtown"

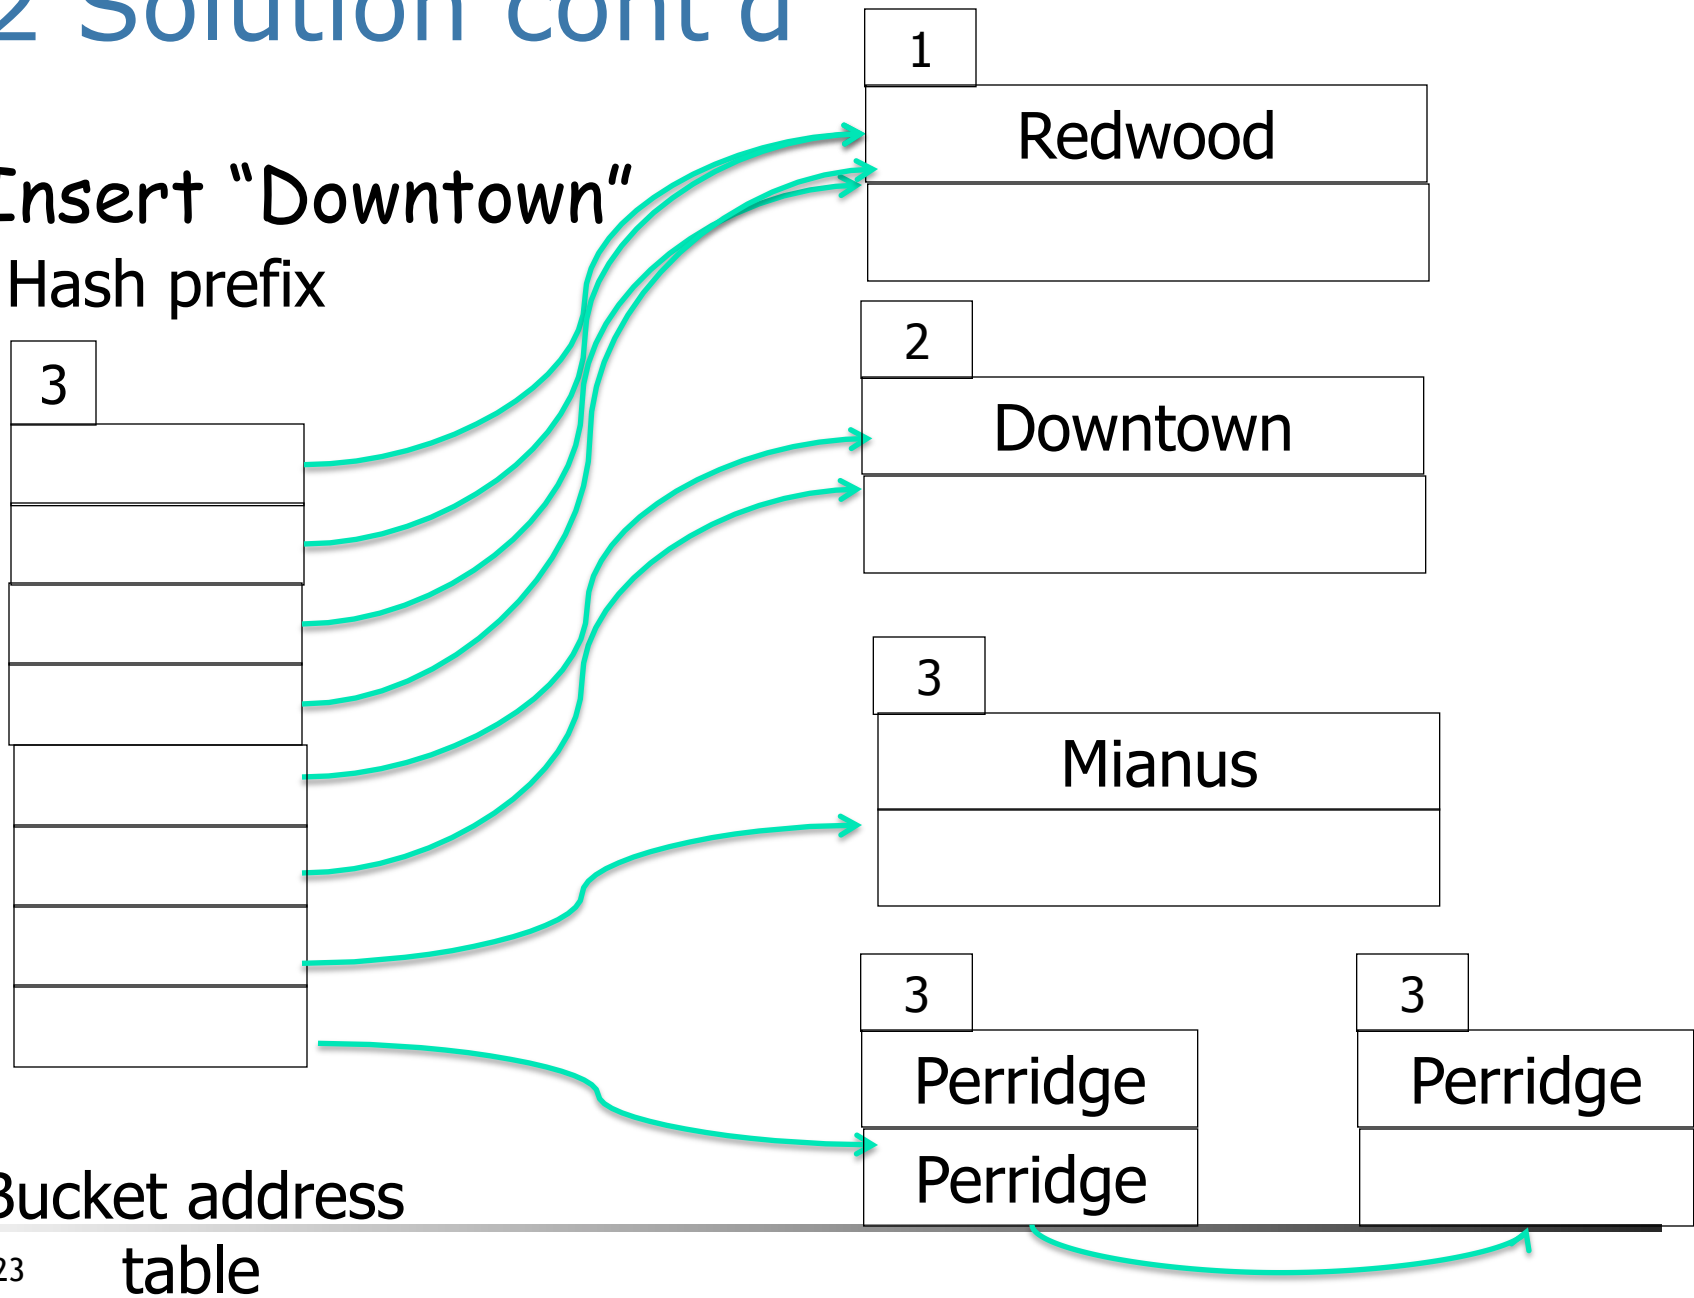

Q2 Solution cont'd

- Insert "Downtown"

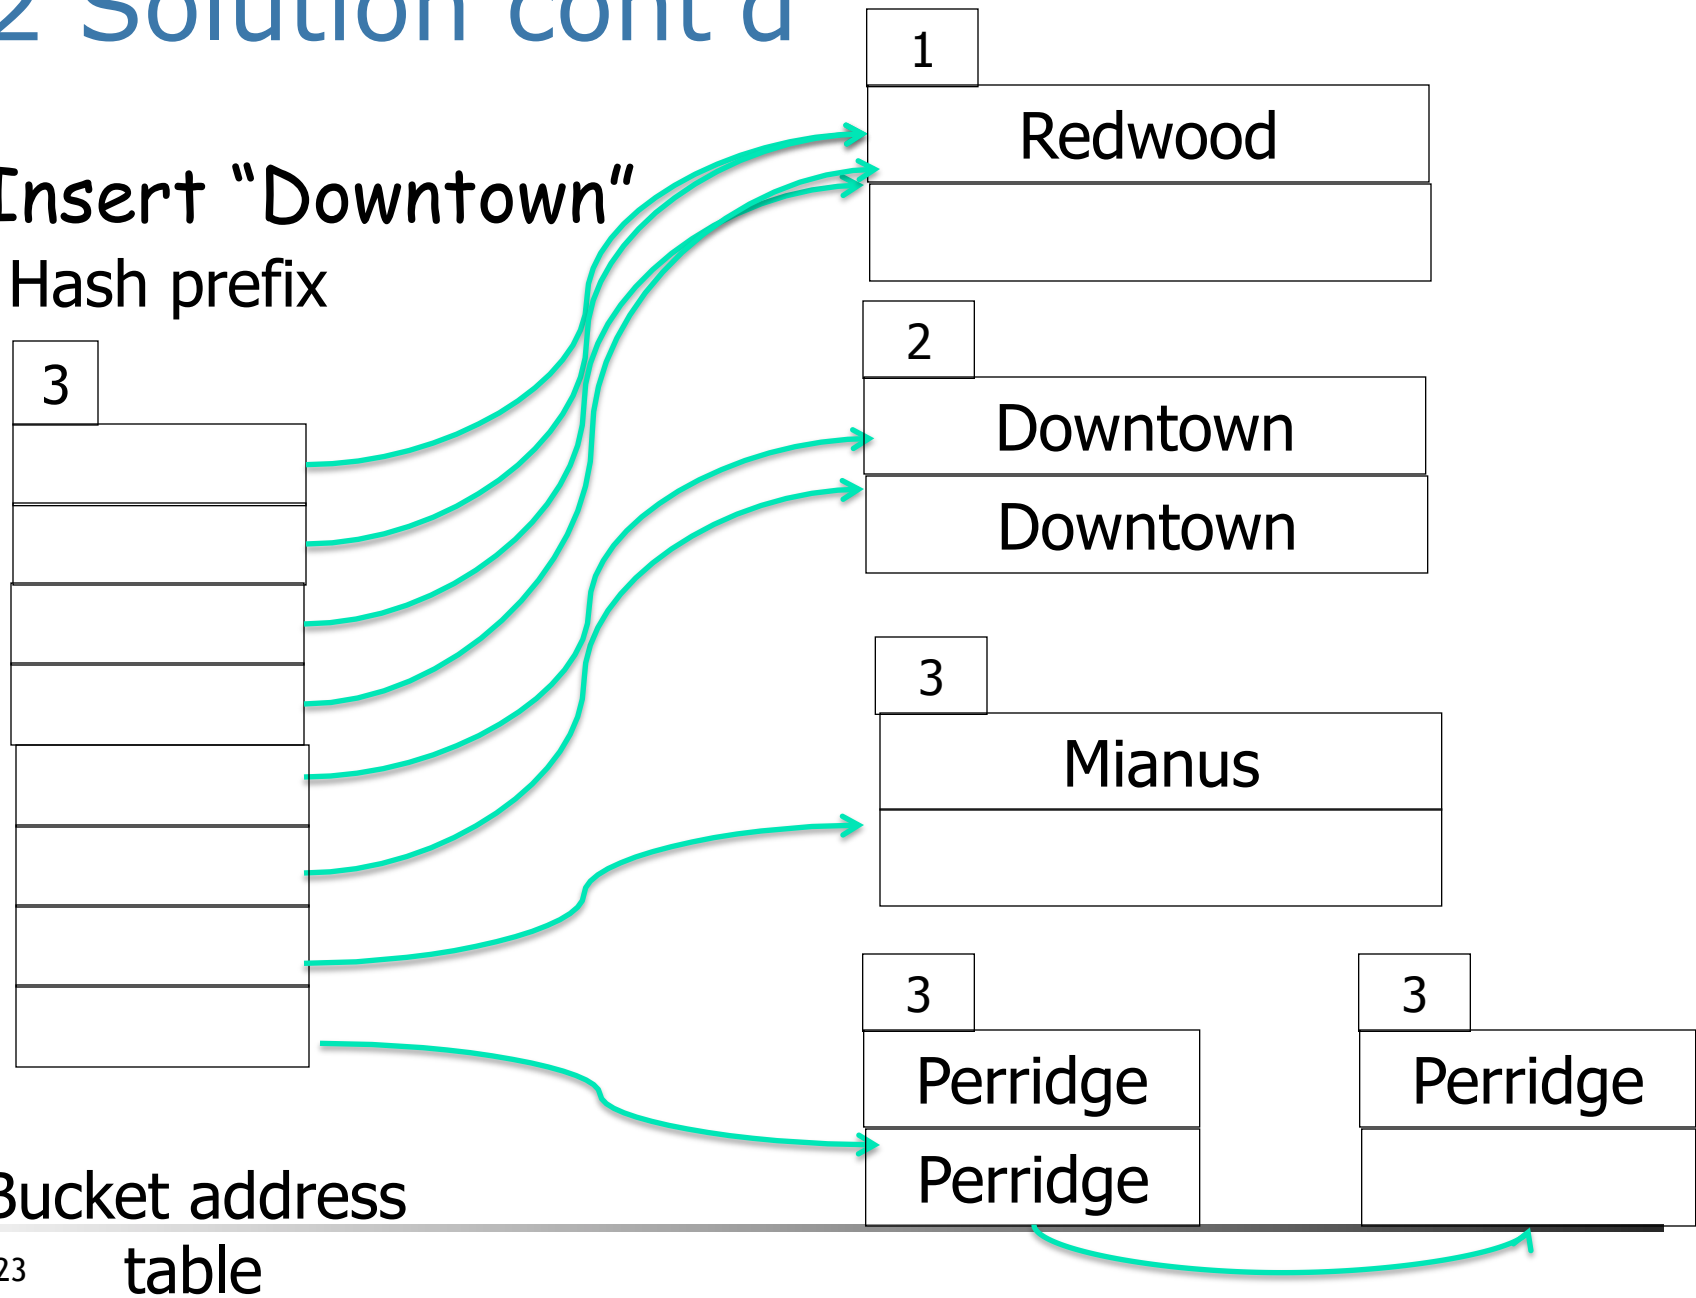

Q2 Solution cont'd

- Insert "Brighton"

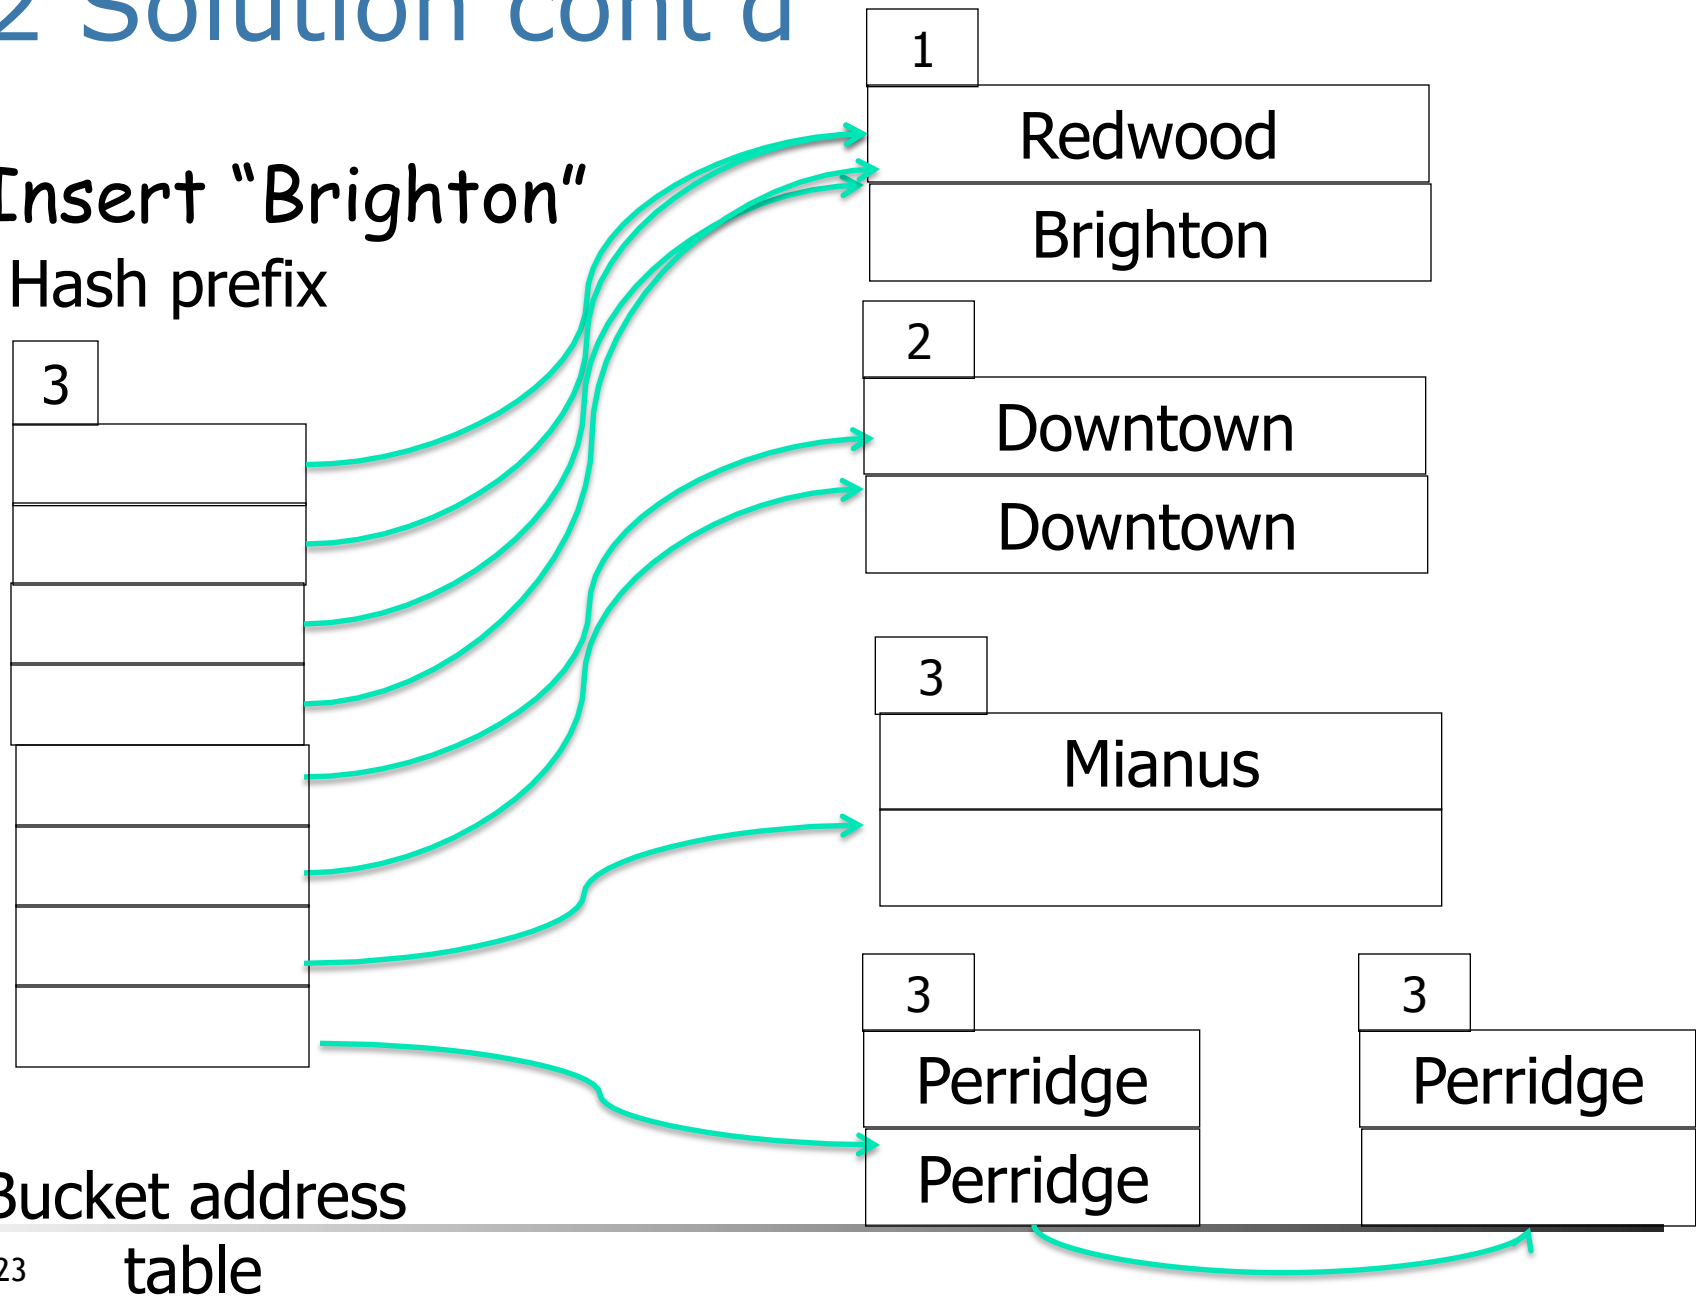

Q3

- Consider the extendable hashing with hash function h(x) = x mod 8 and a bucket can hold three records. Draw the hash index after inserting 13.

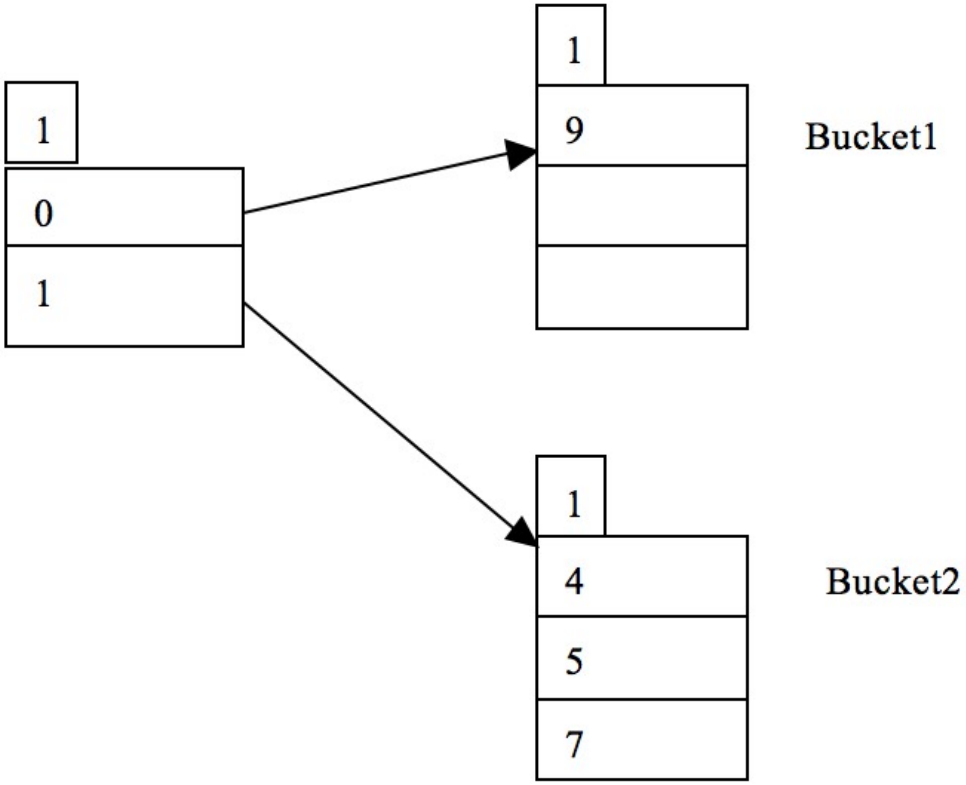

Q3 Solution

- Insertion of 13
  - h(13) = 13 mod 8 = 101, 13 should be in bucket 2, but there is no space and the bucket 2 also can't be split so we have to extend the structure by
  - Doubling the size of address table
  - Splitting the bucket 2 and redistributing the entries in it
  - Insert 13 in bucket 2

Q3 cont'd