

Database Development and Design (CPT201)

Lecture 5a: Introduction to Query Optimisation 1

Dr. Wei Wang

Department of Computing

Learning Outcomes

- Introduction to Query Optimisation
 - Transformation of Relational Expressions




Evaluation of Expressions

- So far we have seen algorithms for individual operations
 - These have then to be combined to evaluate complex expressions, with multiple operations.
- Alternatives for evaluating an entire expression tree
 - **Materialisation:** generate results of an expression whose inputs are relations or relations that are already computed. Temporary relations must be **materialised** (stored) on disk.
 - **Pipelining:** pass on tuples to parent operations even as the operation is being executed.

Materialisation

- **Materialised evaluation:** evaluate one operation at a time, starting at the lowest-level. Use intermediate results materialised into temporary relations to evaluate next-level operations.
 - e.g., in figure below, compute and store the selection, then compute its join with *customer* and store the result, and finally compute the projections on *customer-name*.



Materialisation cont'd

- Materialised evaluation is always applicable
- It may require considerable storage space.
Moreover, cost of writing results to disk and reading them back can be quite high
 - Our cost formulas for operations ignore cost of writing final results to disk, so:
 - Overall cost = Sum of costs of individual operations +
cost of writing intermediate results to disk
- Double buffering: use two output buffers for each operation, when one is full, write it to disk while the other is getting filled
 - Allows overlap of disk writes with computation and reduces execution time

Pipelining

- **Pipelined evaluation:** evaluate several operations simultaneously, passing the results of one operation on to the next.
 - e.g., in previous expression tree, don't store result of the selection
 - instead, pass tuples directly to the join.
 - Similarly, don't store result of join, pass tuples directly to projection.
- It is much cheaper than materialisation: there is no need to store a temporary relation to disk.
- Pipelining may not always be possible - e.g., sort and hash-join where a preliminary phase is required over the whole relations.
- Pipelines can be executed in two ways: **demand driven** and **producer driven**.

Producer-Driven Pipelining



- In producer-driven (or eager or push) pipelining
 - Operators produce tuples eagerly and pass them up to their parents
 - buffer maintained between operators, child puts tuples in buffer, parent removes tuples from buffer
 - if buffer is full, child waits till there is space in the buffer, and then generates more tuples
 - System schedules operations that have space in output buffer and can process more input tuples.

Demand-Driven Pipelining

- In demand driven (or lazy, or pull) evaluation
 - system repeatedly requests next tuple from top level operation
 - Each operation requests next tuple from child operations as required, in order to output its next tuple
 - In between calls, operation has to maintain "state" so it knows what to return next.

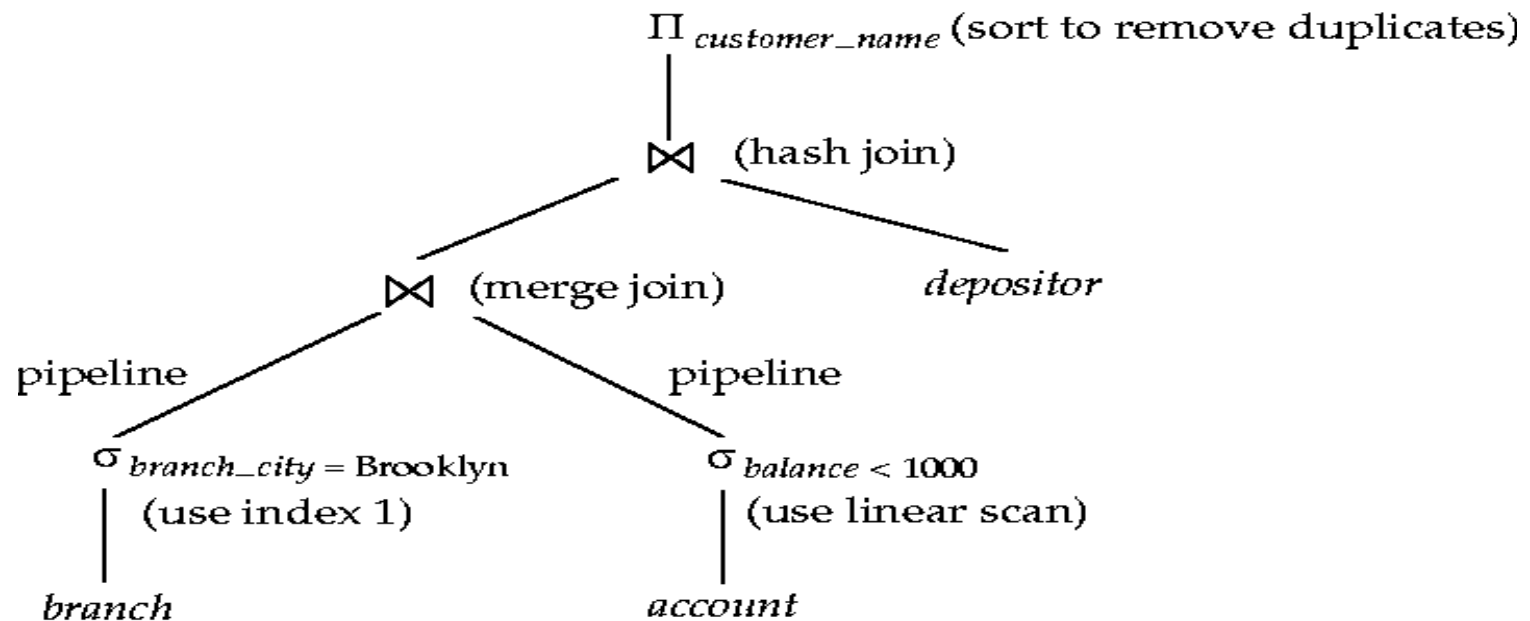
Equivalent expressions

- Alternative ways of evaluating a given query
 - Equivalent expressions
 - Different algorithms for each operation



Evaluation Plan

- An **evaluation plan** defines exactly what algorithm is used for each operation, and how the execution of the operations is coordinated.



Cost-based query Optimisation

- Cost difference between evaluation plans for a query can be enormous
 - E.g. seconds vs. days in some cases
- **Cost-based query optimisation**
 - Find logically equivalent expressions of the given expression (but more efficient to execute)
 - Select a detailed strategy for processing the query, such as choosing the algorithm to use for executing an operation or choosing the specific indices to use
- Estimation of plan cost based on:
 - **Statistical information** about relations, e.g., number of tuples, number of distinct values for an attribute
 - **Statistical estimation for intermediate results** to compute cost of complex expressions
 - **Cost formulae** for algorithms, computed using statistics
- It should be noted that since the cost is an estimate, the selected plan is **not** necessarily the least-costly plan; however, as long as the estimates are good, the plan will not be much more costly than it.

Transformation of Relational Expressions

- Two relational algebra expressions are said to be **equivalent** if the two expressions generate the same set of tuples on every legal database instance
 - Note: order of tuples is irrelevant
 - In SQL, inputs and outputs are multisets of tuples
 - Two expressions in the multiset version of the relational algebra are said to be equivalent if the two expressions generate the same multiset of tuples on every legal database instance.
- An **equivalence rule** says that expressions of two forms are equivalent if
 - Can replace expression of first form by second, or vice versa

Equivalence Rules

- Rule 1: Conjunctive selection operations can be deconstructed into a sequence of individual selections.

$$\sigma_{\theta_1 \wedge \theta_2}(E) = \sigma_{\theta_1}(\sigma_{\theta_2}(E))$$

- Rule 2: Selection operations are commutative.

$$\sigma_{\theta_1}(\sigma_{\theta_2}(E)) = \sigma_{\theta_2}(\sigma_{\theta_1}(E))$$

- Rule 3: Only the last one in a sequence of projection operations is needed, the others can be omitted.

$$\Pi_{L_1}(\Pi_{L_2}(\dots(\Pi_{L_n}(E))\dots)) = \Pi_{L_1}(E)$$

- Rule 4: Selections can be combined with Cartesian products and theta joins.

- (a). $\sigma_{\theta}(E_1 \times E_2) = E_1 \bowtie_{\theta} E_2$
- (b). $\sigma_{\theta_1}(E_1 \bowtie_{\theta_2} E_2) = E_1 \bowtie_{\theta_1 \wedge \theta_2} E_2$

Equivalence Rules cont'd

- Rule 5: Theta-join operations (and natural joins) are commutative.

$$E_1 \bowtie_{\theta} E_2 = E_2 \bowtie_{\theta} E_1$$

- Rule 6.

- (a) Natural join operations are associative:

$$(E_1 \bowtie E_2) \bowtie E_3 = E_1 \bowtie (E_2 \bowtie E_3)$$

- (b) Theta joins are associative in the following manner:

$$(E_1 \bowtie_{\theta_1} E_2) \bowtie_{\theta_2 \wedge \theta_3} E_3 = E_1 \bowtie_{\theta_1 \wedge \theta_3} (E_2 \bowtie_{\theta_2} E_3)$$

where θ_2 involves attributes from only E_2 and E_3 .

Equivalence Rules cont'd

- Rule 7. The selection operation distributes over the theta join operation under the following two conditions:

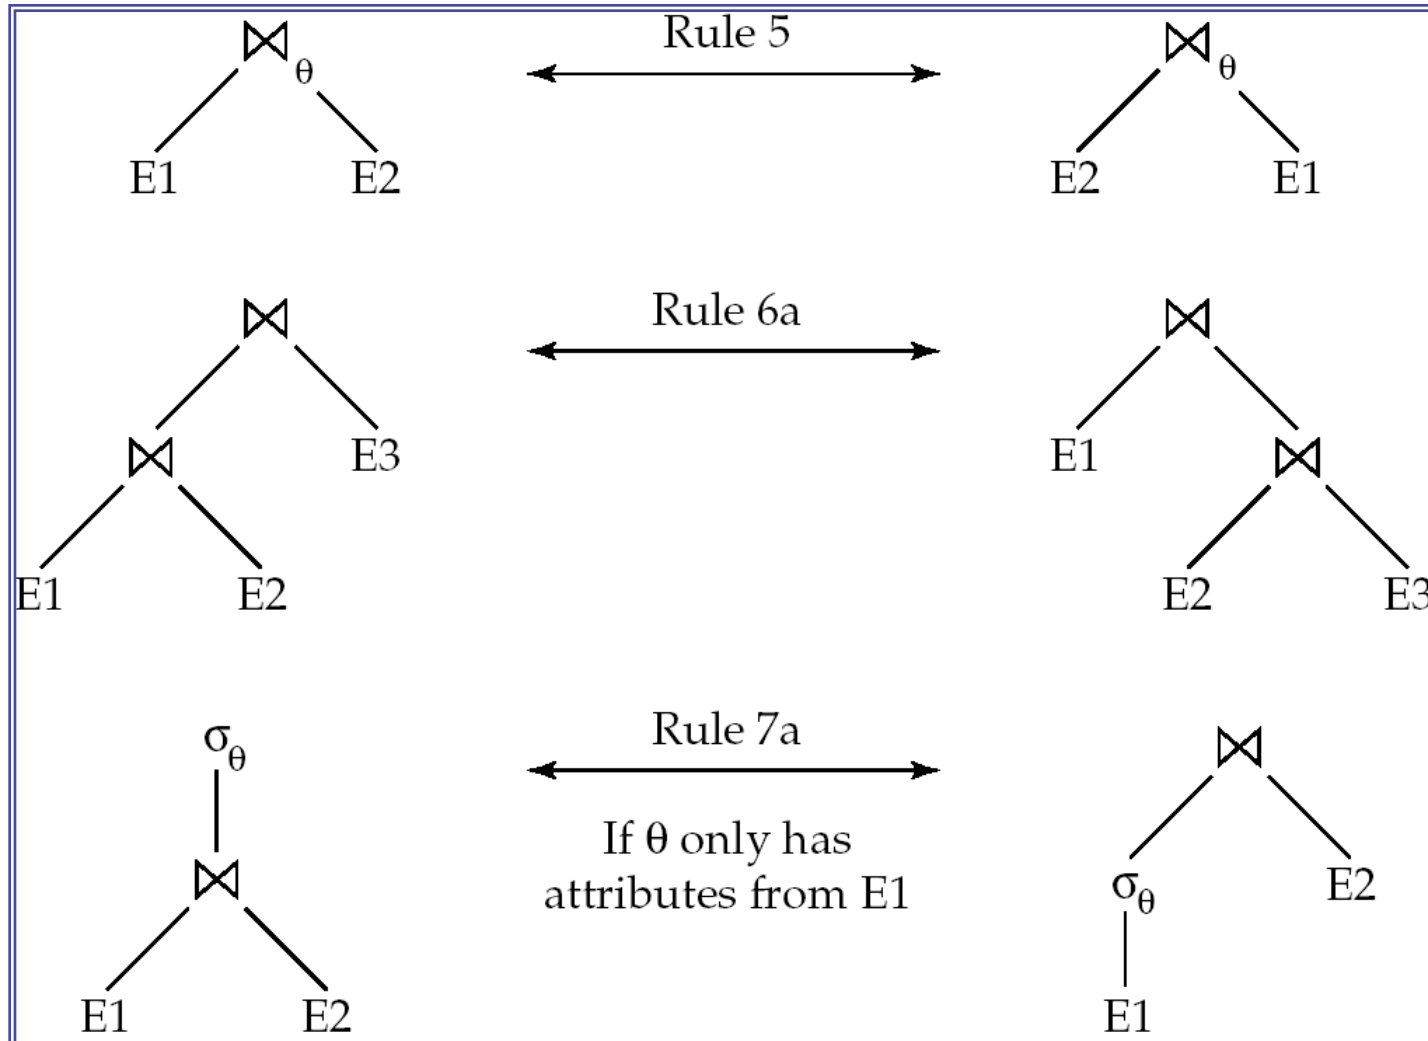
- (a) When θ_0 involves only the attributes of one of the expressions (E_1) being joined.

$$\sigma_{\theta_0}(E_1 \bowtie_{\theta} E_2) = (\sigma_{\theta_0}(E_1)) \bowtie_{\theta} E_2$$

- (b) When θ_1 involves only the attributes of E_1 and θ_2 involves only the attributes of E_2 .

$$\sigma_{\theta_1 \wedge \theta_2}(E_1 \bowtie_{\theta} E_2) = (\sigma_{\theta_1}(E_1)) \bowtie_{\theta} (\sigma_{\theta_2}(E_2))$$

Pictorial Depiction of Equivalence Rules



Equivalence Rules cont'd

■ Rule 8. The projection operation distributes over the theta join operation as follows:

- (a) Let L_1 and L_2 be attributes from E_1 and E_2 , if θ involves only attributes from $L_1 \cup L_2$:

$$\Pi_{L_1 \cup L_2} (E_1 \bowtie_{\theta} E_2) = (\Pi_{L_1} (E_1)) \bowtie_{\theta} (\Pi_{L_2} (E_2))$$

- (b) Consider a join $E_1 \bowtie_{\theta} E_2$.
 - let L_1 and L_2 be sets of attributes from E_1 and E_2 , respectively.
 - let L_3 be attributes of E_1 that are involved in join condition θ , but are not in $L_1 \cup L_2$, and
 - let L_4 be attributes of E_2 that are involved in join condition θ , but are not in $L_1 \cup L_2$.

$$\Pi_{L_1 \cup L_2} (E_1 \bowtie_{\theta} E_2) = \Pi_{L_1 \cup L_2} ((\Pi_{L_1 \cup L_3} (E_1)) \bowtie_{\theta} (\Pi_{L_2 \cup L_4} (E_2)))$$

Equivalence Rules cont'd

- Rule 9. The set operations union and intersection are commutative (set difference is not commutative)

$$\begin{aligned}E_1 \cup E_2 &= E_2 \cup E_1 \\E_1 \cap E_2 &= E_2 \cap E_1\end{aligned}$$

- Rule 10. Set union and intersection are associative.

$$\begin{aligned}(E_1 \cup E_2) \cup E_3 &= E_1 \cup (E_2 \cup E_3) \\(E_1 \cap E_2) \cap E_3 &= E_1 \cap (E_2 \cap E_3)\end{aligned}$$

- Rule 11. The selection operation distributes over \cup , \cap and $-$.

$$\sigma_{\theta}(E_1 - E_2) = \sigma_{\theta}(E_1) - \sigma_{\theta}(E_2)$$

Also:
$$\sigma_{\theta}(E_1 - E_2) = \sigma_{\theta}(E_1) - E_2$$

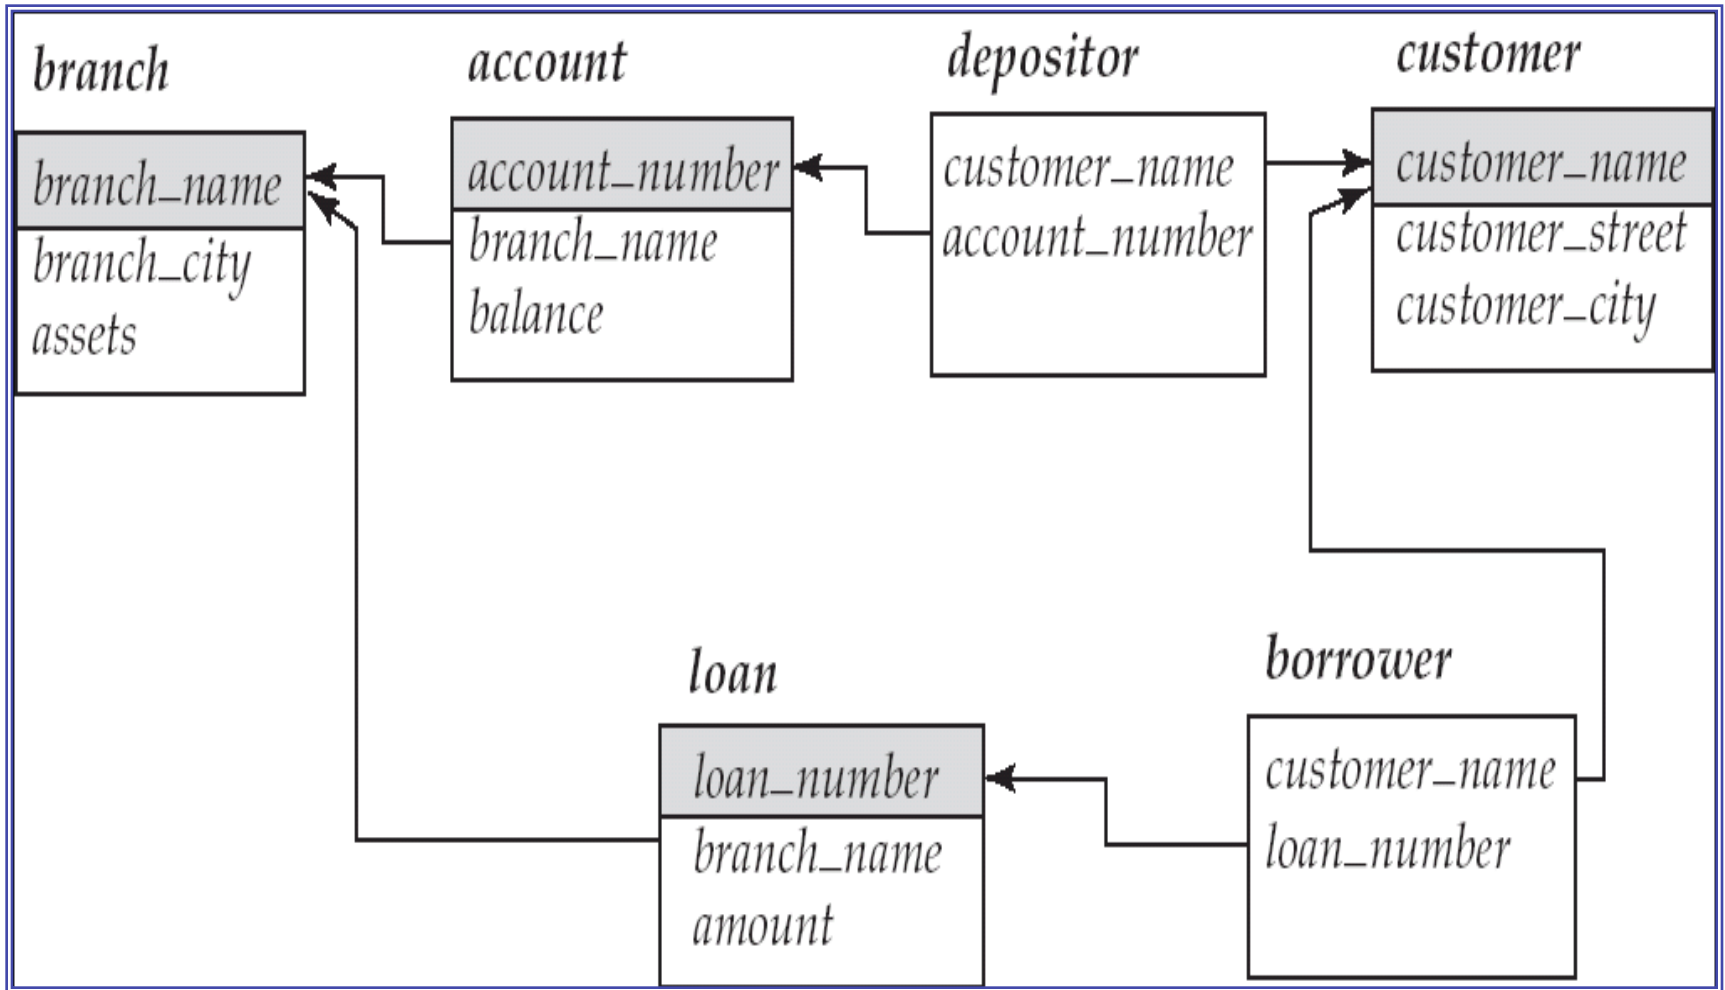
and similarly for \cap in place of $-$, but not for \cup

- Rule 12. The projection operation distributes over union

$$\Pi_L(E_1 \cup E_2) = (\Pi_L(E_1)) \cup (\Pi_L(E_2))$$



Banking Example



Example: Pushing Selections

- Query: Find the names of all customers who have an account at some branch located in Brooklyn.

$\Pi_{customer_name}(\sigma_{branch_city = "Brooklyn"}(branch \bowtie (account \bowtie depositor)))$

- Transformation using rule 7a (distribute the selection).

$\Pi_{customer_name}((\sigma_{branch_city = "Brooklyn"}(branch)) \bowtie (account \bowtie depositor))$

- Performing the selection as early as possible reduces the size of the relation to be joined.

Example: Multiple Transformations

- Query: Find the names of all customers with an account at a Brooklyn branch whose account balance is over \$1000.

$\Pi_{customer_name}(\sigma_{branch_city = "Brooklyn" \wedge balance > 1000} (branch \bowtie (account \bowtie depositor)))$

- Transformation using join associatively (Rule 6a and 7a):


$\Pi_{customer_name}((\sigma_{branch_city = "Brooklyn" \wedge balance > 1000} (branch \bowtie account)) \bowtie depositor)$

- Second form provides an opportunity to apply the “**perform selections early**” rule, resulting in the subexpression 7b

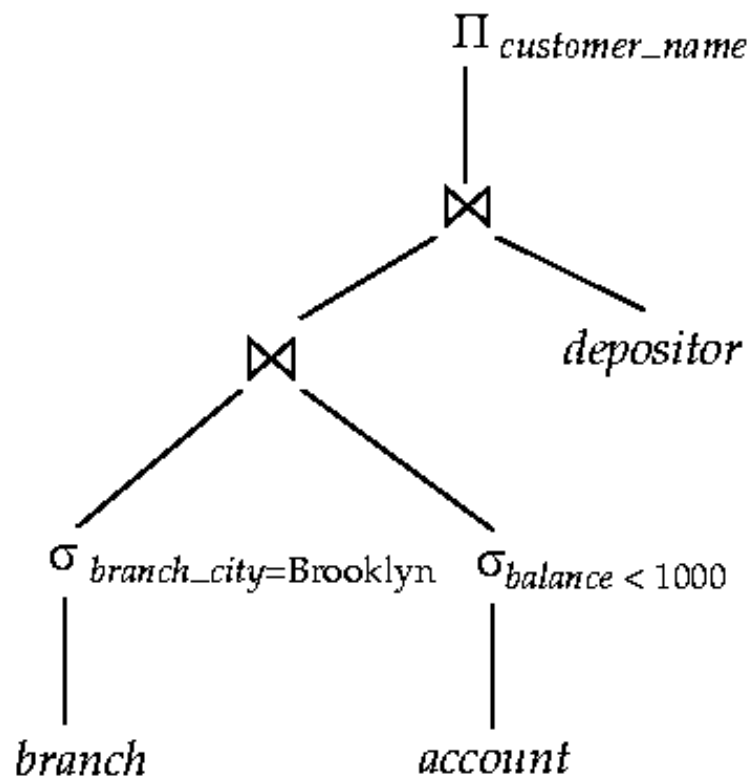
$\sigma_{branch_city = "Brooklyn"} (branch) \bowtie \sigma_{balance > 1000} (account)$

- Thus a sequence of transformations can be useful

Multiple Transformations cont'd



(a) Initial expression tree



(b) Tree after multiple transformations

Transformation Example: Pushing Projections

$\Pi_{customer_name}((\sigma_{branch_city = "Brooklyn"} (branch) \bowtie account) \bowtie depositor)$

- When we compute

$(\sigma_{branch_city = "Brooklyn"} (branch) \bowtie account)$

we obtain a relation whose schema is:

$(branch_name, branch_city, assets, \textcolor{red}{account_number}, balance)$

- Push projections using equivalence rules 8b; eliminate unneeded attributes from intermediate results to get:

$\Pi_{customer_name}((\Pi_{account_number}(\sigma_{branch_city = "Brooklyn"} (branch) \bowtie account)) \bowtie depositor))$

(HINT: L1 is null, L2 is customer_name; L3=L4=account_number)

- Performing projection as early as possible reduces the size of the tuples to be joined.

Join Ordering Example

- For all relations r_1, r_2 , and r_3 ,

$$(r_1 \bowtie r_2) \bowtie r_3 = r_1 \bowtie (r_2 \bowtie r_3)$$

(Join Associativity)

- If $r_2 \bowtie r_3$ is quite large and $r_1 \bowtie r_2$ is small, we choose

$$(r_1 \bowtie r_2) \bowtie r_3$$

so that we compute and store a smaller temporary relation.

Join Ordering Example cont'd

- Consider the expression

$\Pi_{customer_name} ((\sigma_{branch_city = \text{"Brooklyn"}}(branch)) \bowtie (account \bowtie depositor))$

- Could compute $account \bowtie depositor$ first, and join result with

$\sigma_{branch_city = \text{"Brooklyn"}}(branch)$
but $account \bowtie depositor$ is likely to be a large relation.

- Only a small fraction of the bank's customers are likely to have accounts in branches located in Brooklyn
 - it is better to compute first

$\sigma_{branch_city = \text{"Brooklyn"}}(branch) \bowtie account$

Enumeration of Equivalent Expressions

- Query optimisers use equivalence rules to systematically generate expressions equivalent to the given expression
- The approach is very expensive in space and time

```
procedure genAllEquivalent( $E$ )
```

```
 $EQ = \{E\}$ 
```

```
repeat
```

```
    Match each expression  $E_i$  in  $EQ$  with each equivalence rule  $R_j$ 
```

```
    if any subexpression  $e_i$  of  $E_i$  matches one side of  $R_j$ 
```

```
        Create a new expression  $E'$  which is identical to  $E_i$ , except that  
         $e_i$  is transformed to match the other side of  $R_j$ 
```

```
        Add  $E'$  to  $EQ$  if it is not already present in  $EQ$ 
```

```
until no new expression can be added to  $EQ$ 
```

End of Lecture

- Summary
 - Transformation of Relational Expressions
- Reading
 - Textbook chapter 13.1, 13.2, 13.3, and 13.4

Database Development and Design (CPT201)

Lecture 5b: Introduction to Query Optimisation 2

Dr. Wei Wang

Department of Computing

Learning Outcomes

- Introduction to Query Optimisation
 - Catalog Information for Cost Estimation
 - Cost-based optimisation




Catalog Information for Cost Estimation

- n_r : number of tuples in a relation r .
- b_r : number of blocks containing tuples of r .
- l_r : size of a tuple of r .
- f_r : blocking factor of r . i.e., the number of tuples of r that fit into one block.
- $V(A, r)$: number of **distinct** values that appear in r for attribute A ; same as the size of $\Pi_A(r)$.
- If tuples of r are stored together physically in a file, then:
$$b_r = \left\lceil \frac{n_r}{f_r} \right\rceil$$

Histograms

- Histogram on attribute *age* of relation *person*
- Equi-width histograms
- Equi-depth histograms



Estimation of the Size of Selection

- $\sigma_{A=v}(r)$
 - $n_r / V(A,r)$: number of records that will satisfy the selection
 - Equality condition on a **key** attribute (primary key): *size estimate* = 1
- $\sigma_{A \leq v}(r)$ (case of $\sigma_{A \geq v}(r)$ is symmetric)
 - Let c denote the estimated number of tuples satisfying the condition. Let $\min(A,r)$ and $\max(A,r)$ denote the lowest and highest values for attribute A .
 - If $\min(A,r)$ and $\max(A,r)$ are available in catalog
 - $c = 0$ if $v < \min(A,r)$
 - $c = n_r \cdot \frac{v - \min(A,r)}{\max(A,r) - \min(A,r)}$
 - If histograms available, can refine above estimate
 - In absence of statistical information c is assumed to be $n_r / 2$.

Estimation of the Size of Joins

- The Cartesian product $r \times s$ contains $n_r \cdot n_s$ tuples; each tuple occupies $s_r + s_s$ bytes.
- If $R \cap S = \emptyset$, then $r \bowtie s$ is the same as $r \times s$.
- If $R \cap S$ is a key for R , then a tuple of s will join with at most one tuple from r
 - therefore, the number of tuples in $r \bowtie s$ is no greater than the number of tuples in s .
- If $R \cap S$ is a foreign key in S referencing R , then the number of tuples in $r \bowtie s$ is exactly the same as the number of tuples in s .
 - The case for $R \cap S$ being a foreign key referencing S is symmetric.
- In the example query *depositor* \bowtie *customer*, *customer_name* in *depositor* is a foreign key (of *customer*)
 - hence, the result has exactly $n_{\text{depositor}}$ tuples, which is 5000



Estimation of the Size of Joins

cont'd

- If $R \cap S = \{A\}$ is not a key for R or S .

If we assume that every tuple t in R produces tuples in $R \bowtie S$, the number of tuples in $R \bowtie S$ is estimated to be:

$$\frac{n_r * n_s}{V(A, s)}$$

If the reverse is true, the estimate obtained will be:

$$\frac{n_r * n_s}{V(A, r)}$$

The **lower of these two estimates** is probably the more accurate one.

- Can improve on above if histograms are available
 - Use formula similar to above, for each cell of histograms on the two relations

Join Operation: Running Example

- Running example: *depositor* ⋈ *customer*
- Catalog information for join examples:
 - $n_{customer} = 10,000$.
 - $f_{customer} = 25$, which implies that $b_{customer} = 10,000/25 = 400$.
 - $n_{depositor} = 5000$.
 - $f_{depositor} = 50$, which implies that $b_{depositor} = 5,000/50 = 100$.
 - $V(customer_name, depositor) = 2,500$, which implies that, on average, each customer has two accounts.
 - Also assume that *customer_name* in *depositor* is a foreign key on *customer*.
 - $V(customer_name, customer) = 10,000$ (primary key)

Join Operation: Running Example cont'd

- Compute the size estimates for *depositor* ⋈ *customer* without using information about foreign keys:
 - $V(\text{customer_name}, \text{depositor}) = 2,500$, and
 $V(\text{customer_name}, \text{customer}) = 10,000$
 - The two estimates are $5,000 * 10,000 / 2,500 = 20,000$
and $5,000 * 10,000 / 10,000 = 5,000$
 - We choose the lower estimate, which in this case, is the same as our earlier computation using foreign keys.

Size Estimation for Other Operations

- Projection: estimated size of $\Pi_A(r) = V(A,r)$
- Set operations
 - For unions/intersections of selections on the **same** relation: rewrite and use size estimate for selections
 - e.g., $\sigma_{\theta_1}(r) \cup \sigma_{\theta_2}(r)$ can be rewritten as $\sigma_{\theta_1 \vee \theta_2}(r)$
 - For operations on **different** relations:
 - estimated size of $r \cup s$ = size of r + size of s .
 - estimated size of $r \cap s$ = minimum size of r and size of s .
 - estimated size of $r - s$ = r .
 - All the three estimates may be quite **inaccurate**, but provide **upper bounds** on the sizes.

Estimation of Number of Distinct Values in Selection

- If θ forces A to take a specified value: $V(A, \sigma_{\theta}(r)) = 1$.
 - e.g., $A = 3$
- If θ forces A to take on one of a specified set of values:
$$V(A, \sigma_{\theta}(r)) = \text{number of specified values.}$$
 - (e.g., $(A = 1 \vee A = 3 \vee A = 4)$),
- If the selection condition θ is of the form $A \text{ op } v$ (op is $>$, $<$, etc),
$$V(A, \sigma_{\theta}(r)) = V(A, r) * s$$
 - where s is the selectivity of the selection.
- In all the other cases: use approximate estimate of $\min(V(A, r), n_{\sigma_{\theta}(r)})$

Estimation of Distinct Values cont'd

Joins: $r \bowtie s$

- If all attributes in A are from r ,
estimated $V(A, r \bowtie s) = \min(V(A, r), n_{r \bowtie s})$
- If A contains attributes $A1$ from r and $A2$ from s , then
estimated
 $V(A, r \bowtie s) =$
 $\min(V(A1, r) * V(A2 - A1, s), V(A1 - A2, r) * V(A2, s), n_{r \bowtie s})$
 - More accurate estimate can be got using probability theory, but this one works fine generally
- Projections: Estimation of distinct values are straightforward for projections.
 - They are the same in $\Pi_A(r)$ as in r .

Choice of Evaluation Plans

- Must consider the **interaction** of evaluation techniques when choosing evaluation plans
 - choosing the cheapest algorithm for each operation independently may not yield best overall algorithm, e.g.
 - merge-join may be costlier than hash-join, but may provide a sorted output which reduces the cost for an outer level aggregation.
 - nested-loop join may provide opportunity for **pipelining**
- Practical query optimisers incorporate elements of the following two broad approaches:
 - Search all the plans and choose the best plan in a cost-based fashion.
 - Uses heuristics to choose a plan.

Cost-Based Join Order Optimisation

- Consider finding the best join-order for

$$r_1 \bowtie r_2 \bowtie \dots \bowtie R_n.$$

- There are $(2(n-1))!/(n-1)!$ different join orders for above expression. With $n = 7$, the number is 665280, with $n = 10$, the number is greater than 176 billion!
- No need to generate all the join orders. Using **dynamic programming**, the least-cost join order for any subset of $\{r_1, r_2, \dots, r_n\}$ is computed only once and stored for future use.

Dynamic Programming in Optimisation

- To find best plan (join tree) for a set of n relations:
 - Consider all possible plans of the form: $S_1 \bowtie (S - S_1)$, where S_1 is any non-empty subset of S .
 - Recursively compute cost for joining subsets of S to find the cost of each plan. Choose the cheapest of the alternatives.
 - Base case for recursion: single relation access plan
 - Find the best selection strategy for a particular relation R_i
 - When plan for any subset is computed, store it and reuse it when it is required again, instead of re-computing it.

Join Order Optimisation Algorithm

```
// initialise bestplan[S].cost to  $\infty$ 
procedure findbestplan(S)
  if (bestplan[S].cost  $\neq \infty$ )
    return bestplan[S]
  // else bestplan[S] has not been computed earlier, compute it now
  if (S contains only 1 relation)
    set bestplan[S].plan and bestplan[S].cost based on the best way
    of accessing S /* Using selections on S and indices on S */
  else for each non-empty subset S1 of S such that S1  $\neq$  S
    P1= findbestplan(S1)
    P2= findbestplan(S - S1)
    A = best algorithm for joining results of P1 and P2
    cost = P1.cost + P2.cost + cost of A
    if cost < bestplan[S].cost
      bestplan[S].cost = cost
      bestplan[S].plan = "execute P1.plan; execute P2.plan;
                          join results of P1 and P2 using A"
  return bestplan[S]
```

Cost of Join Order Optimisation

- With dynamic programming time complexity of optimisation with bushy trees is $O(3^n)$.
 - With $n = 10$, this number is 59000 instead of 176 billion!
- Space complexity is $O(2^n)$ as the number of subsets of the S is 2^n .
- Although both numbers still increase rapidly with n , commonly occurring joins usually have less than 10 relations, and can be handled easily.

Cost-Based Optimisation with Equivalence Rules

- Many optimisers follow an approach based on
 - Using heuristic transformations to handle constructs other than joins
 - applying the cost-based join order selection algorithm to subexpressions involving only joins and selections
- General-purpose cost-based optimiser based on equivalence rules
 - easy to extend the optimiser with new rules to handle different query constructs
 - but the procedure to enumerate all equivalent expressions is very expensive

Cost-Based Optimisation with Equivalence Rules cont'd


- To make the approach work efficiently requires the following:
 - A space-efficient representation of expressions
 - Efficient techniques for detecting duplicate derivations of the same expression
 - dynamic programming based on memoisation
 - avoid generating all possible equivalent plans

Heuristic Optimisation

- Cost-based optimisation is expensive, even with dynamic programming.
- Systems may use *heuristics* to reduce the number of choices that must be made in a cost-based fashion.
- Heuristic optimisation transforms the query-tree by using a set of rules that typically (but not in all cases) improve execution performance:
 - Perform selection early (reduces the number of tuples)
 - Perform projection early (reduces the number of attributes)
 - Perform the most restrictive selection and join operations (i.e., with smallest result size) before other similar operations.
- Some systems use only heuristics, others combine heuristics with partial cost-based optimisation.

Other heuristics: Left Deep Join Trees

- In **left-deep join trees**, the right-hand-side input for each join is a relation, not the result of an intermediate join.



Cost of left-deep join Optimisation

- To find best left-deep join tree for a set of n relations:
 - Consider n alternatives with one relation as right-hand side input and the other relations as left-hand side input.
 - Modify optimisation algorithm:
 - Replace "for each non-empty subset S_1 of S such that $S_1 \neq S$ "
 - By: for each relation r in S , let $S_1 = S - r$.
- If only left-deep trees are considered, time complexity of finding best join order is $O(n!)$, with dynamic programming this can be reduced to $O(n 2^n)$
 - Space complexity remains at $O(2^n)$
- Cost-based optimisation is expensive, but worthwhile for queries on large datasets (typical queries have small n , generally < 10)

Structure of Query Optimisers

- Many optimisers considers only left-deep join orders.
 - Plus heuristics to push selections and projections down the query tree
 - Reduces optimisation complexity and generates plans amenable to pipelined evaluation.
- Heuristic optimisation used in some versions of Oracle:
 - Repeatedly pick “best” relation to join next
 - Starting from each of n starting points. Pick best among these

Structure of Query Optimisers

cont'd

- Some query optimisers integrate heuristic selection and the generation of alternative access plans.
 - Frequently used approach
 - heuristic rewriting of nested block structure and aggregation
 - followed by cost-based join-order optimisation for each block
 - Some optimisers (e.g. SQL Server) apply transformations to entire query and do not depend on block structure
- Even with the use of heuristics, cost-based query optimisation imposes a substantial overhead.
 - But is worth for expensive queries
 - Optimisers often use simple heuristics for very cheap queries, and perform exhaustive enumeration for more expensive queries

End of Lecture

■ Summary

- Transformation of Relational Expressions
- Catalog Information for Cost Estimation
- Cost-based optimisation
- Dynamic Programming for Choosing Evaluation Plans

■ Reading

- Textbook chapter 13.1, 13.2, 13.3, and 13.4